

VOLUME II | RESEARCH DERIVATIONS

One Geometry, Two Research Questions

Superposition, robustness
and recursive-training stability

A derivation companion to the monosemanticity paper

Exact toy results, proofs, counterexamples, numerical checks,
and a route to testing the same mechanism in diffusion models.

Project 1: the superposition-to-robustness phase diagram.

Project 2: model collapse as monosemanticity decay.

Structure: mathematics $\rightarrow$ controlled toy models $\rightarrow$ real-model validation.

Horizon: eight-week research core, with a two-week extension if available.

Prepared: 20 September 2026; Project 1 derivations updated 6 October 2026. Conditional derivations are distinguished from empirical hypotheses and from claims of research novelty.

What this volume delivers

This volume keeps the two agreed projects and their shared monosemanticity motivation. It does not replace them with a benchmark, an agent system or a different application area. It supplies a mathematically explicit foundation that can support their development.

The exact formulas delivered here include: operational perturbation radii; reconstruction risk; a finite-state feature-detection risk; a small, nonlinear superposition phase boundary with feature frequency and importance; a general interference-based error bound; exact recursive toy dynamics; a bound connecting one-generation feature errors to cumulative distribution drift; a local derivative for a genuinely retrained model; and a Gaussian denoising calculation clarifying the diffusion connection.

These are *finished conditional calculations*, not a claim that the proposed conference contribution is already established. Some are standard facts presented with proofs; others are worked combinations specialized to this program. Their novelty has not been certified. The general learned-geometry phase diagram and a validated monosemanticity-based predictor in real diffusion models remain the research targets.

That distinction is essential to using these notes productively. A theorem can be correct and still be too elementary for a new paper. A promising experiment can still contradict a preferred mechanism. The purpose of the derivations is to make the remaining research time test precise claims, not to promise an outcome that has not been observed.

Project 1 update: exact scope of the completed mathematics

The added clean-training and noise chapters derive the finite bias profiles, dense-case optima, a certified low-frequency geometry, the global clean frequency transition, and the at-least-two noise-crossing result. They preserve the distinction between reconstruction-selected representations and downstream binary detection. All new mathematics remains in this volume rather than only in working Markdown files. The arbitrary-load/importance target and real-model transfer are not complete; the original Project 2 material is preserved as background, with no new recursive-training experiment introduced.

The scientific question in plain language

When a model shares representational directions between concepts, it can store more information in limited space. That sharing can also create confusion. We want to understand the conditions under which the gain in coverage outweighs the cost in errors, and whether the same errors accumulate when generated outputs become future training data.

This is a general learning-and-representation question. Safety is a possible application, but no adversary is needed for the main study. Recursive training supplies the connection to the team’s self-improvement interests; a self-consuming loop does not count as improvement unless quality actually improves. Vision diffusion provides a real generative setting in which to test the mechanism.

Read these chapters in sequence

Chapters 1–3 establish definitions and operational interpretations. Chapters 4–6 derive fixed-geometry Project 1 results and a general interference bound. Chapters 7–8 derive clean-selected geometry, the global frequency transition and the certified nonmonotone noise comparison. Chapters 9–12 derive Project 2 and the mathematical connection. Chapters 13–15 explain diffusion, controlled validation and the original eight-week execution outline. The appendices give a formula ledger, exercises, numerical results and primary references.

No team member is assigned ownership in this document. The work has a continuous theoretical and experimental thread, with room for complementary interests in interpretation, robustness, evaluation, post-training and generative modeling.

Contents

What this volume delivers	i
1 Fix the question before manipulating the symbols	1
1.1 The commitments that remain fixed	1
1.2 The variables	1
1.3 Three noise locations	1
1.4 Three outcomes	2
2 What the original margin formula really means	3
2.1 The general linear decision radius	3
2.2 A normalized matched-readout example	3
2.3 It is a matched readout, not always the optimal readout	4
2.4 Interpreting isolation, rarity and signed overlap	4
2.5 Other threat models	4
3 Linear reconstruction: the useful calculation and its limit	5
3.1 The exact reconstruction risk	5
3.2 An isotropic example with no distributional advantage	5
3.3 A controlled way forward	6
4 Project 1: an exact risk for arbitrary fixed geometry	7
4.1 A matched feature detector	7
4.2 What the formula buys us	7
4.3 Rarity also affects the best decision threshold	8
4.4 Defining the phase boundary	8
4.5 Distinguishing geometric overlap from statistical interference	8
5 Project 1: a closed tradeoff for a superposed pair	9
5.1 Two features, one dimension	9
5.2 Enumerate the four states	9
5.3 The clean boundary in closed form	10
5.4 Noise moves the boundary	10
5.5 A reconstruction version, to show task dependence	10
5.6 What is solved and what is still research	11
6 Project 1: load, interference and a general bound	12
6.1 Centering interference	12
6.2 A robust region with an operational threshold	13

6.3	How load enters	13
6.4	A capacity lower bound	13
6.5	Importance, correlations and learned geometry	13
7	Project 1: clean training selects the feature geometry	15
7.1	The training-to-geometry question	15
7.2	Global decoder-bias minimization at fixed geometry	15
7.3	Monosemantic retention under the same budget	16
7.4	The equal antipodal pair with trained biases	17
7.5	Importance changes the selected amplitudes	18
7.6	Exact global optima at activation probability one half	18
7.6.1	Bias profiles in both sign sectors	18
7.6.2	Equal importance: unequal norms are globally selected	19
7.6.3	Importance ratio two: mono retention is global	19
7.7	A computer-assisted global certificate at probability one fifth	20
7.8	The global frequency transition at importance ratio two	21
7.8.1	Excluding all same-sign sharing	22
7.8.2	Opposite-sign strong-feature profile	22
7.8.3	Every competing weak-feature bias branch	22
7.8.4	Global comparison above the threshold	23
7.8.5	Strict sharing improvement below the threshold	24
7.9	What is established and what remains a research question	24
8	Project 1: robustness of the clean-selected representation	25
8.1	A fixed evaluation rule after clean training	25
8.2	Deriving the exact four-state Gaussian risk	25
8.3	The clean-selected mono control	26
8.4	Why reconstruction and detection can disagree	26
8.5	The infinite-noise comparison	26
8.6	A rigorous example with at least two crossings	27
8.7	Deriving the finite-noise rational certificate	28
8.7.1	From an algebraic root to interval probabilities	28
8.7.2	An exact enclosure of the normalizing constant	29
8.7.3	Rational upper and lower Gaussian CDF bounds	29
8.7.4	The certified sign and what is independently checked	30
8.8	The fixed numerical checks and their interpretation	30
8.9	Verification failures and stopping scope	31
8.10	The remaining paper-critical decisions	31
8.11	Does the reconstruction reversal survive decoder calibration?	32
8.11.1	A fixed, globally solved comparison	32
8.11.2	Gaussian moments and the scalar calibration objective	33
8.11.3	Derivatives and a uniform curvature bound	33

8.11.4	Explicitly excluding the two infinite tails	34
8.11.5	The declared numerical branch-and-bound calculation	34
8.11.6	Why both clean-optimal relative signs have the same profile	34
8.11.7	Prespecified outcome	35
8.11.8	A linear reference separates generic overlap from rectification	36
8.12	A joint storage and robustness boundary near the clean transition	38
8.12.1	Deriving the clean critical law	38
8.12.2	Frozen-bias corrupted risk and its boundary	39
8.12.3	A second frozen-bias ordering reversal	39
8.13	How symmetric oracle bias calibration changes the critical law	40
8.14	An arbitrary-load bound for fixed geometries under bias calibration	40
8.15	Verification status and contribution scope	41
8.16	An exact global certificate for the prescribed algebraic frequencies	42
8.16.1	Exact root brackets without rationalizing the probability	44
8.17	Prescribed finite reconstruction-risk maps	44
8.18	Numerical calibration bounds at critical precision	45
8.19	Finite calibrated boundary comparisons and numerical resolution	46
8.20	The prescribed eight-feature clean test and its limitations	47
8.21	A bounded repair of the larger-toy optimizer	48
8.21.1	Population gradients and the energy constraint	48
8.21.2	Retraction and actual-loss backtracking	49
8.21.3	Why merely decaying projected Adam is inconclusive	49
8.21.4	Recorded repair outcomes and the stopping decision	50
8.22	The critical law persists over an interval of feature importances	50
8.22.1	Uniform clean critical expansion	53
8.22.2	Gaussian population risks and the frozen-bias law	53
8.22.3	Global bias localization under calibration	54
8.22.4	Mono calibration and the uniform calibrated law	54
8.22.5	What boundary is proved	55
8.23	Fixed finite validation of the importance-family law	55
8.24	A singular endpoint separates frozen and calibrated robustness	56
8.24.1	Global clean selection, including competing sign sectors	58
8.24.2	The two clean critical branches	59
8.24.3	Frozen risk at the singular endpoint	59
8.24.4	Symmetric global bias calibration	60
8.24.5	Prospective finite validation and stopping scope	60
8.25	The two prescribed endpoint-policy comparisons	61
8.26	Continuous-target profiling for the learned-vision pilot	61
8.26.1	Complete clean bias profiling	63
8.26.2	Gaussian conditional moments for continuous targets	63
8.26.3	Finite interval and exterior bounds	64

8.26.4	Numerical tolerances, matched policies and scope	64
8.27	A fixed learned-vision check of decoder-policy dependence	65
8.27.1	Inputs, selected geometry and evaluation settings	65
8.27.2	All prescribed held-out comparisons	66
8.27.3	Audit, relevance and stopping decision	66
8.27.4	Saved-record gate and contribution audit	68
8.28	Class-mapped evidence transfer: a resolved sharing advantage	69
8.28.1	Prospective feature definition and controls	69
8.28.2	Held-out risks and the unfavorable transfer limit	69
8.28.3	Calibration contrast, independent audit and stop	70
9	Project 2: the correct meaning of the covariance recursion	72
9.1	A fixed linear map	72
9.2	Why rank deficiency does not imply gradual decay	72
9.3	The same margin can mean shrinkage, stability or growth	72
9.4	Adaptive maps and injected variation	73
9.5	Finite-sample shrinkage is estimator-specific	73
10	Project 2: an exact nonlinear generation operator	74
10.1	Define the transition before discussing collapse	74
10.2	The pair kernel in closed form	74
10.3	An explicit order parameter for a defined loss	75
10.4	What the eigenvalues mean here	75
10.5	A monosemantic comparator	76
11	The mathematical link between the two projects	77
11.1	Use the generation mechanism as a coupling	77
11.2	Importance-weighted version	77
11.3	A replenishment bound	78
11.4	When the bound is informative	78
11.5	Where monosemanticity enters the actual research claim	78
12	What changes when the model is genuinely retrained	79
12.1	Training is part of the dynamics	79
12.2	Deriving the response of the learned parameters	79
12.3	The local generation Jacobian	80
12.4	A tractable staged experiment	80
12.5	Order parameters worth testing	80
13	Diffusion: the bridge that is valid, and the one to test	81
13.1	A diffusion model in five equations	81
13.2	Two clocks, two different processes	81
13.3	A Gaussian denoiser that can be solved completely	82

13.4	Posterior sampling changes the answer completely	82
13.5	Where a representation-level contribution can enter	83
13.6	Existing work the new claim must distinguish itself from	83
14	Experiments that test the derivations	84
14.1	Stage A: exact identities and finite-state calculations	84
14.2	Stage B: learn the geometry in controlled toys	84
14.3	Stage C: recursive training on the same toys	84
14.4	Stage D: one controlled image generator	85
14.5	Stage E: a small natural-image diffusion demonstration	85
14.6	How to measure geometry consistently	85
14.7	Interventions that support a mechanism	86
14.8	Evaluation and uncertainty	86
14.9	A compact evidence table	86
15	An eight-week program that preserves both projects	87
15.1	Status of this execution outline	87
15.2	One integrated paper structure	87
15.3	Weekly deliverables	87
15.4	What to reduce if the compute budget is tight	88
15.5	What would make the paper substantive	88
A	Formula ledger: retain, qualify or replace	90
B	Further derivations and worked exercises	91
B.1	Gaussian ReLU moments for continuous reconstruction	91
B.2	Mixture covariance includes a mean term	91
B.3	Exercise 1: an exact adversarial perturbation	91
B.4	Exercise 2: rare feature versus common nuisance	91
B.5	Exercise 3: why a largest overlap and a sum differ	92
B.6	Exercise 4: solve the mixed fixed kernel	92
B.7	Exercise 5: covariance does not identify concepts	92
B.8	Exercise 6: a sharply stated falsification	92
C	Numerical verification record	93
D	Primary references and claim boundaries	94

Chapter 1

Fix the question before manipulating the symbols

Learning objective. Specify the prediction task, corruption, representation and training loop so a formula has a definite meaning.

1.1 The commitments that remain fixed

Project 1 seeks a robust-performance tradeoff in sparsity, feature-to-dimension load and importance decay. Project 2 seeks a representation-level account of recursive-training degradation and a quantity that predicts it. Their intended combination uses shared feature geometry, shared toy infrastructure and a real-model demonstration. The source anchor is Zhang et al.'s monosemanticity-and-robustness paper, taught in Volume I.

The mathematical choices below instantiate this direction; they do not exhaust it. We begin with a task whose ground truth is unambiguous: recovering binary feature presence. This allows an actual error probability to replace an informal “cleanness” score. Continuous-amplitude reconstruction is also derived as a companion task.

1.2 The variables

Let $b = (b_1, \dots, b_n) \in \{0, 1\}^n$ denote genuine toy concepts. Initially $b_i \sim \text{Bernoulli}(p_i)$ independently, unless otherwise stated. Let $s_i = 1 - p_i$ and let $I_i > 0$ weight errors. The encoder is $W \in \mathbb{R}^{m \times n}$, with columns w_i . Define

$$h = Wb, \quad G = W^\top W. \quad (1.1)$$

The load is n/m . Later, a training generation is indexed by g , and the geometry may become W_g, G_g . A diffusion denoising step uses a different index τ .

1.3 Three noise locations

1. **Feature-space corruption:** b or a continuous feature vector is perturbed before encoding. A feature-space adversary is an idealized mathematical model.
2. **Code-space corruption:** $h = Wb + \xi$ with $\xi \sim \mathcal{N}(0, \sigma^2 I_m)$. This makes each matched readout's variance depend on G_{ii} .
3. **Observed-input corruption:** pixels or waveforms change before a nonlinear encoder. The encoder's Jacobian and nonlinear activation changes matter.

These yield different formulas. A proof for the first cannot be described as a guarantee for the third without an additional argument.

1.4 Three outcomes

Feature detection error asks whether a concept is correctly declared present. **Reconstruction error** asks how accurately its amplitude is recovered. **Distributional fidelity** asks whether a generator produces the correct frequencies and combinations of concepts. A classifier can retain accuracy while a generator loses diversity, so these outcomes should be linked mathematically rather than treated as interchangeable.

For binary detection, the main risk is

$$R(W, \theta; P) = \sum_i I_i \mathbb{P}_{b \sim P, \xi} \{\hat{b}_i \neq b_i\}. \quad (1.2)$$

For a phase diagram, compare the risks of two admissible representation families at the same stated resource budget. A boundary is the set where their risks are equal. Calling it a “phase boundary” is descriptive here; a finite toy model does not automatically have a thermodynamic singularity.

Before moving on. For every proposed formula, write one sentence specifying what is random, what is fixed and what outcome it predicts.

What the original margin formula really means

Learning objective. Retain the valid geometric identity while making its adversarial interpretation exact.

2.1 The general linear decision radius

Consider a binary decision with label $y \in \{-1, 1\}$ and score $f(x) = \beta^\top x - c$, where $\beta \neq 0$. If $yf(x) > 0$, the example is correctly classified. For a perturbation with $\|\delta\|_q \leq \varepsilon$, Hölder's inequality gives

$$\min_{\|\delta\|_q \leq \varepsilon} yf(x + \delta) = yf(x) - \varepsilon \|\beta\|_{q^*}, \quad \frac{1}{q} + \frac{1}{q^*} = 1. \quad (2.1)$$

Equality is attained by a dual-norm maximizing direction. Therefore the exact distance to the decision boundary is

$$r_q(x, y) = \frac{yf(x)}{\|\beta\|_{q^*}}. \quad (2.2)$$

This assumes unrestricted perturbations in the chosen vector space. If valid inputs must remain nonnegative, binary or on a manifold, the feasible attack set changes and this may become a lower bound on the constrained radius rather than the exact answer.

2.2 A normalized matched-readout example

For the tied linear readout of feature i , take $\beta = Ge_i$, $x = e_i$, $c = 0$, $y = 1$ and $q = 2$. Then

$$M_i = \frac{G_{ii}}{\sqrt{\sum_j G_{ij}^2}} = \left(1 + \sum_{j \neq i} \frac{G_{ij}^2}{G_{ii}^2}\right)^{-1/2}. \quad (2.3)$$

This requires $G_{ii} > 0$. If the feature is dropped, the entire row is zero and the ratio is 0/0, not a meaningful robustness score.

Equation (2.3) is also the cosine between the row of G and the target feature axis. It is a legitimate geometric quantity. It equals an attack radius only in the normalized input and decision problem just specified. A different input x , threshold c or feature amplitude changes the radius through Equation (2.2).

The squared alignment score also equals the established per-feature capacity in Scherlis et al., *Polysemanticity and Capacity in Neural Networks*, Equation (3):

$$C_i = \frac{(w_i^\top w_i)^2}{\sum_j (w_i^\top w_j)^2} = M_i^2. \quad (2.4)$$

This identity prevents treating the score itself as a new contribution. A dropped feature requires a separate convention because the ratio is undefined.

2.3 It is a matched readout, not always the optimal readout

The choice $w_i^\top h$ corresponds to the tied decoder. A general linear estimator of x_i from h need not use w_i . For centered data with covariance C and additive code-noise covariance N , the least-squares coefficient is

$$a_i^* = (WCW^\top + N)^\dagger WCe_i, \quad (2.5)$$

with the appropriate identifiability conditions or pseudoinverse convention. Its effective feature-space direction is $W^\top a_i^*$. Calling Ge_i “the optimal readout” without these qualifications was unjustified.

2.4 Interpreting isolation, rarity and signed overlap

Can an attacker move a perfectly isolated feature? Yes. If $G = I$, then $M_i = 1$, but changing x_i changes its reading. An ability to move the score does not imply $M_i < 1$.

Does a negative overlap count? In this unrestricted worst-case ℓ_2 setting, G_{ij} enters squared, because an adversary can choose the harmful sign. For ordinary nonnegative data and a threshold decoder, the sign can be beneficial or harmful depending on co-occurrence. The square loses that distinction.

Does rarity determine M_i ? No. At fixed G , Equation (2.3) has no activation probability in it. Rarity may affect what G learns, or statistical collision probabilities, but it cannot be inserted into this fixed-geometry identity by multiplication.

Is M_i scale invariant? It is invariant under a common rescaling $W \mapsto aW$, since all entries of G scale by a^2 . Rescaling only one column changes its overlaps relative to its diagonal and generally changes M_i . The diagonal has not disappeared physically merely because it is hidden in normalized overlaps.

2.5 Other threat models

For a code-space perturbation $h \mapsto h + \delta_h$, the tied readout changes by $w_i^\top \delta_h$, so its worst-case swing is $\varepsilon \|w_i\|_2$. For random feature noise with covariance N_x , the output variance is $(Ge_i)^\top N_x (Ge_i)$. These differ from each other and from the feature-space worst-case swing.

For a nonlinear observed-input map $h = f(u)$, a local first-order calculation uses $\nabla_u[w_i^\top f(u)]$. To obtain a finite-radius guarantee, bound the remainder or a Jacobian norm over the entire perturbation region. A Jacobian at one point is a local sensitivity measurement, not automatically a certificate.

Status of this result. Equations in this chapter are standard linear estimation and norm-duality facts. They repair the operational meaning of the earlier margin; they are not, by themselves, a novel conference contribution.

Chapter 3

Linear reconstruction: the useful calculation and its limit

Learning objective. Put sparsity and importance into a risk correctly, and see why a purely linear model can miss the desired tradeoff.

3.1 The exact reconstruction risk

Let x have mean μ and covariance C . Observe $x + \epsilon$, where ϵ has mean zero, covariance N and is independent of x . Reconstruct $\hat{x} = G(x + \epsilon)$. Set $D = \text{diag}(I_1, \dots, I_n)$. Then

$$L = \mathbb{E}[(x - \hat{x})^\top D(x - \hat{x})] \quad (3.1)$$

$$= \text{tr}[D\{(I - G)C(I - G)^\top + GNG^\top\}] + \|D^{1/2}(I - G)\mu\|_2^2. \quad (3.2)$$

To prove this, write the error as $(I - G)(x - \mu) + (I - G)\mu - G\epsilon$. Cross terms involving centered independent variables vanish. Use $\mathbb{E}[z^\top Dz] = \text{tr}(D \text{Cov } z) + (\mathbb{E}z)^\top D(\mathbb{E}z)$.

For independent spike-and-uniform features, $\mu_i = p_i/2$ and $C_{ii} = p_i/3 - p_i^2/4$. For binary features, $\mu_i = p_i$ and $C_{ii} = p_i(1 - p_i)$. Importance weights multiply the error after these distributional quantities are specified. This is an exact place for rarity and importance to enter, unlike inserting p into Equation (2.3).

3.2 An isotropic example with no distributional advantage

Take centered, isotropic features $C = vI_n$, isotropic input noise $N = \sigma_x^2 I_n$ and equal importance. Let F be the Gram matrix of a unit-column tight frame, so

$$\text{diag } F = \mathbf{1}, \quad F^2 = \frac{n}{m}F, \quad \text{rank } F = m. \quad (3.3)$$

Use $G = gF$. Every feature sees aggregate squared overlap $E = n/m - 1$, and its reconstruction risk is

$$\ell(g) = v(1 - g)^2 + vg^2E + \sigma_x^2g^2(1 + E). \quad (3.4)$$

Differentiate, set the derivative to zero, and obtain

$$g_* = \frac{v}{(v + \sigma_x^2)(1 + E)}. \quad (3.5)$$

The minimized total risk is

$$L_* = nv - \frac{mv^2}{v + \sigma_x^2}. \quad (3.6)$$

Now compare a monosemantic linear representation retaining m coordinates, each with optimal shrinkage $a = v/(v + \sigma_x^2)$, and dropping the rest. Each retained coordinate has risk $v\sigma_x^2/(v + \sigma_x^2)$; each dropped coordinate costs v . Their sum is the same expression.

This is not evidence against the project. It shows that isotropic second moments with a linear decoder can erase the sparse, nonlinear mechanism of interest. The learned sparse-superposition phenomenon cannot be assumed to survive replacing its decoder by an arbitrary linear map.

3.3 A controlled way forward

Keep a threshold or ReLU decoder, retain ground-truth concepts, and include the risk of dropped features in the comparator. The next chapters do this. The linear identities remain useful as sanity checks and baselines, but the nonlinear feature-detection model produces an actual sparsity-dependent tradeoff.

Before moving on. Explain why two distributions with identical mean and covariance have identical risk in Equation (3.2), but can behave differently after a nonlinear decoder.

Project 1: an exact risk for arbitrary fixed geometry

Learning objective. Derive an operational error probability that contains feature frequency, importance and geometry explicitly.

4.1 A matched feature detector

Observe $h = Wb + \xi$ with $\xi \sim \mathcal{N}(0, \sigma^2 I_m)$. Read feature i with

$$z_i = w_i^\top h = (Gb)_i + \zeta_i, \quad \hat{b}_i = \mathbf{1}\{z_i > \theta_i\}. \quad (4.1)$$

The vector $\zeta = W^\top \xi$ is Gaussian with covariance $\sigma^2 G$. Its components are generally correlated. For a single detector, its marginal variance is $\sigma^2 G_{ii}$.

Let Φ denote the standard normal CDF. Conditional on a fixed binary state b , the error probability is

$$e_i(b) = \Phi\left(-\frac{(2b_i - 1)((Gb)_i - \theta_i)}{\sigma\sqrt{G_{ii}}}\right). \quad (4.2)$$

When $b_i = 1$, failure means the Gaussian score lies below threshold. When $b_i = 0$, failure means it lies above threshold. These are exactly the two signs in the formula.

Proposition 4.1 (Exact finite-state weighted risk). *For $\sigma > 0$ and all scored columns satisfying $G_{ii} > 0$,*

$$R(G, \theta; P) = \sum_{i=1}^n I_i \sum_{b \in \{0,1\}^n} P(b) e_i(b). \quad (4.3)$$

For an initially independent distribution, $P(b) = \prod_j p_j^{b_j} (1 - p_j)^{1-b_j}$.

Proof. Condition each indicator of error on b , use Equation (4.2), and apply linearity of expectation. No independence between detector errors is needed. Independence of input features is needed only for the product formula for P . $\square$

4.2 What the formula buys us

This is an exact model-based prediction, not a regression fit. For small n , enumerate states and calculate risk directly. For larger n , Monte Carlo estimates the same defined quantity. Importance decay enters as $I_i = i^{-\gamma}$, frequency enters through P and capacity enters through the admissible rank- m geometry.

A dropped feature requires separate treatment. A detector constantly predicting zero has error p_i ; a best constant detector has error $\min(p_i, 1 - p_i)$. The choice must be fixed when defining the comparison. We use the zero-output reconstruction convention below, appropriate to representing absence when a concept is omitted.

4.3 Rarity also affects the best decision threshold

The midpoint thresholds used in the explicit pair calculation are fixed detector choices, not a claim of Bayes optimality for every prior. For an isolated scalar code $h = ab + \xi$ with $b \sim \text{Bernoulli}(p)$, $0 < p < 1$, and Gaussian noise variance σ^2 , compare the two posterior probabilities. Their log-likelihood ratio is

$$\log \frac{\mathbb{P}(b = 1 \mid h)}{\mathbb{P}(b = 0 \mid h)} = \log \frac{p}{1-p} + \frac{ah}{\sigma^2} - \frac{a^2}{2\sigma^2}. \quad (4.4)$$

Under equal false-positive and false-negative costs, the optimal threshold is

$$\theta_{\text{Bayes}} = \frac{a}{2} + \frac{\sigma^2}{a} \log \frac{1-p}{p}. \quad (4.5)$$

A rare feature therefore requires stronger evidence before it is declared present. This can reduce total error while increasing the fraction of true rare occurrences that are missed. Report recall as well as average error. Re-optimizing thresholds is an important ablation: the fixed-threshold phase boundary below is exact for its detector family, not a universal Bayes-risk boundary.

4.4 Defining the phase boundary

Let $\mathcal{G}_{\text{mono}}$ and $\mathcal{G}_{\text{super}}$ be two admissible families with the same declared capacity and norm constraints. A population-optimal comparison would use

$$R_a^*(p, I, \sigma, m) = \inf_{G \in \mathcal{G}_a, \theta \in \Theta_a} R(G, \theta; P), \quad a \in \{\text{mono}, \text{super}\}. \quad (4.6)$$

The boundary is $R_{\text{super}}^* = R_{\text{mono}}^*$. This definition is precise, but solving the optimization for all loads and importance schedules is not accomplished merely by writing it. We next solve a fixed, interpretable two-feature family exactly. Learned-model sweeps then test whether the predicted local mechanism appears in trained geometries.

4.5 Distinguishing geometric overlap from statistical interference

The earlier expression $p(n/m) \leq \tau^{-2} - 1$ inserted an average active-feature count into a worst-case alignment formula. An unrestricted attacker can perturb inactive coordinates too; that substitution is not valid. In Equation (4.3), activation probabilities enter the actual distribution of states, so a sparsity-dependent risk has a legitimate derivation.

The project remains a phase-diagram project. Its phase variable is now a defined error comparison rather than an arbitrary cutoff on a scale-dependent proxy.

Status of this result. The finite-state expectation is exact under the detector model. It is a starting formula, not yet a solved global optimization over all learned dictionaries.

Project 1: a closed tradeoff for a superposed pair

Learning objective. Derive a genuine boundary, including rare concepts, importance weights and matched resource controls.

5.1 Two features, one dimension

Take independent binary features with probabilities p_1, p_2 . A superposed code is

$$h = a(b_1 - b_2) + \xi, \quad \xi \sim \mathcal{N}(0, \sigma^2), \quad a > 0. \quad (5.1)$$

Decode with $\hat{b}_1 = \mathbf{1}\{h > a/2\}$ and $\hat{b}_2 = \mathbf{1}\{h < -a/2\}$. This corresponds to a tied matched score with an appropriate threshold. A monosemantic comparator stores feature 1 as $h_m = cb_1 + \xi$, predicts it at threshold $c/2$ and outputs zero for feature 2.

The two encoders have the same latent dimension, but equal $a = c = 1$ gives different squared Frobenius norms: 2 versus 1. Therefore report two controls: equal stored-feature amplitude ($a = c = 1$), and equal total encoder energy ($a = 1/\sqrt{2}, c = 1$). This prevents an apparent robustness effect from being explained entirely by scale.

5.2 Enumerate the four states

Write

$$q = \Phi\left(-\frac{a}{2\sigma}\right), \quad q_3 = \Phi\left(-\frac{3a}{2\sigma}\right), \quad q_m = \Phi\left(-\frac{c}{2\sigma}\right). \quad (5.2)$$

Let $P_{00} = (1 - p_1)(1 - p_2)$, $P_{10} = p_1(1 - p_2)$, $P_{01} = (1 - p_1)p_2$ and $P_{11} = p_1p_2$. Conditional weighted errors are

State	Clean code	Expected weighted detection error
00	0	$(I_1 + I_2)q$
10	a	$I_1q + I_2q_3$
01	$-a$	$I_1q_3 + I_2q$
11	0	$(I_1 + I_2)(1 - q)$

For state 10, missing feature 1 needs $\xi \leq -a/2$, while falsely predicting feature 2 needs $\xi < -3a/2$. For state 11, both features should be present but the decoder can output at most one; each individual feature is missed with probability $1 - q$. This is the nonlinear interference cost.

Theorem 5.1 (Exact pair risk and boundary). *The superposed pair has risk*

$$R_p = (I_1 + I_2)P_{00}q + P_{10}(I_1q + I_2q_3) + P_{01}(I_1q_3 + I_2q) + P_{11}(I_1 + I_2)(1 - q). \quad (5.3)$$

The two zero-output monosemantic comparators have risks

$$R_{m1} = I_1 q_m + I_2 p_2, \quad R_{m2} = I_2 q_m + I_1 p_1. \quad (5.4)$$

For these specified codes and decoders, superposition is preferable exactly when $R_p < \min(R_{m1}, R_{m2})$. Equality defines the boundary.

Proof. Average the four conditional risks using their probabilities. For a stored mono feature, false-positive and false-negative probabilities both equal q_m . The omitted feature is wrong whenever it is active, giving the additional $I_j p_j$ term. $\square$

5.3 The clean boundary in closed form

As $\sigma \rightarrow 0$, $q, q_3, q_m \rightarrow 0$. Thus

$$R_p = (I_1 + I_2) p_1 p_2, \quad R_{m1} = I_2 p_2. \quad (5.5)$$

For $p_2 > 0$, the superposed pair beats the comparator retaining feature 1 precisely when

$$p_1 < \frac{I_2}{I_1 + I_2}. \quad (5.6)$$

For equal frequencies and importances, this is $p < 1/2$, or $s > 1/2$. The mathematical mechanism is explicit: the cost of losing an omitted feature scales with one activation probability, while collision costs scale with a joint activation probability.

If feature 1 is more important, increasing I_1/I_2 makes its protection more valuable and shrinks the region where sharing its direction is worthwhile. For an importance schedule $I_j = j^{-\gamma}$, pairing features i and j gives the corresponding threshold $p_i < j^{-\gamma}/(i^{-\gamma} + j^{-\gamma})$ against retaining i . This is a derived importance dependence, not an inferred slogan that rare concepts are always safe.

5.4 Noise moves the boundary

Equation (5.3) gives an exact boundary in terms of normal CDFs. It need not reduce to a short elementary polynomial to be analytic. Numerical root finding on a proved expression is different from fitting an unknown law to data.

At $p_1 = p_2 = 0.2$, equal importances and equal stored-feature amplitude, the risk crossing is $\sigma \simeq 0.658350$. Under equal total encoder energy, it occurs at $\sigma \simeq 0.279939$. Both comparisons have a tradeoff; its position depends strongly on the resource convention. This is exactly why that convention belongs in the theorem.

5.5 A reconstruction version, to show task dependence

Now let $x_i = b_i U_i$ with uniform amplitudes and noiseless $h = x_1 - x_2$. Decode with ReLUs. When only one feature is active, reconstruction is exact. When both are active, each coordinate's error magnitude is $\min(U_1, U_2)$, and

$$\mathbb{E}[\min(U_1, U_2)^2] = \int_0^1 2t \mathbb{P}(\min(U_1, U_2) > t) dt = \frac{1}{6}. \quad (5.7)$$

Therefore $L_p = (I_1 + I_2) p_1 p_2 / 6$, while dropping feature 2 costs $L_{m1} = I_2 p_2 / 3$. The clean threshold becomes $p_1 < 2I_2 / (I_1 + I_2)$, different from the binary detection threshold. A universal boundary independent of task would miss this difference.




Figure 5.1: Exact equal-energy pair phase diagram. Blue indicates lower weighted error for the superposed pair; red indicates lower error for the specified monosemantic comparator. The black curve is the equality of the two analytic risks. This is not a fitted phase boundary or a result about an arbitrary trained diffusion model.

5.6 What is solved and what is still research

This chapter solves a small, explicit family of codes and decoders. It does not prove that those codes globally minimize the learned autoencoder objective at every parameter setting. The next experiments should compare analytic predictions with measured trained geometries, and test whether optimization changes the boundary. The solved pair is an interpretable unit from which larger constructions and hypotheses can be built.

Status of this result. The risk formulas and pair boundaries are exact under their stated model. Their status as a novel contribution, and their transfer to learned networks, are not established by this derivation alone.

Chapter 6

Project 1: load, interference and a general bound

Learning objective. Connect arbitrary feature load to an error bound without replacing a worst-case norm by an average active count.

6.1 Centering interference

For independent input features, choose

$$\theta_i = \sum_{j \neq i} G_{ij} p_j + \frac{G_{ii}}{2}. \quad (6.1)$$

Then the distance from the conditional mean to threshold is $G_{ii}/2$ for either value of b_i . Define

$$S_i = \sum_{j \neq i} G_{ij} (b_j - p_j), \quad V_i = \sum_{j \neq i} G_{ij}^2 p_j (1 - p_j), \quad B_i = \max_{j \neq i} |G_{ij}|. \quad (6.2)$$

The nuisance fluctuation is $S_i + \zeta_i$, whose variance proxy is $V_i + \sigma^2 G_{ii}$. Unlike M_i , this expression has activation probabilities in a legitimate statistical role.

Theorem 6.1 (A sufficient feature-error bound). *For independent Bernoulli features and independent Gaussian code noise, the detector above satisfies*

$$e_i \leq \exp \left[- \frac{G_{ii}^2}{8 \{ V_i + \sigma^2 G_{ii} + B_i G_{ii} / 6 \}} \right]. \quad (6.3)$$

The bound applies separately to each class-conditional error and hence to their mixture. Degenerate zero-variance cases are interpreted by their deterministic limit.

Proof. Each centered summand $G_{ij}(b_j - p_j)$ has absolute magnitude at most B_i and total variance V_i . Its independent sum obeys the Bernstein exponential-moment bound with variance parameter V_i and scale $B_i/3$. Adding independent Gaussian ζ_i adds $\sigma^2 G_{ii}$ to the variance parameter. Chernoff optimization gives either one-sided tail bound

$$\mathbb{P}(\pm(S_i + \zeta_i) \geq t) \leq \exp \left[- \frac{t^2}{2(V_i + \sigma^2 G_{ii} + B_i t / 3)} \right]. \quad (6.4)$$

Set $t = G_{ii}/2$. The class-0 error uses the upper tail and the class-1 error the lower tail. Each is bounded by the same expression, so averaging them adds no factor two. $\square$

6.2 A robust region with an operational threshold

For a target error $e_i \leq \delta$, a sufficient condition is

$$V_i + \sigma^2 G_{ii} + \frac{B_i G_{ii}}{6} \leq \frac{G_{ii}^2}{8 \log(1/\delta)}. \quad (6.5)$$

This defines a guaranteed region for the specified detector. It is a sufficient bound, not an “if and only if” boundary: failure of the inequality does not imply failure of the detector. The exact pair calculation is sharper in its small setting; the bound covers many features.

6.3 How load enters

For unit columns, define mutual coherence $\mu = \max_{i \neq j} |G_{ij}|$. Then $G_{ii} = 1$, $B_i \leq \mu$ and

$$V_i \leq \mu^2 \sum_{j \neq i} p_j (1 - p_j). \quad (6.6)$$

For independent random unit directions in m dimensions, $\mathbb{E} G_{ij}^2 = 1/m$, so

$$\mathbb{E} V_i = \frac{1}{m} \sum_{j \neq i} p_j (1 - p_j). \quad (6.7)$$

For equal activation probability this is $(n-1)p(1-p)/m$. This is an expectation over random dictionaries, not a universal identity for a trained dictionary or a high-probability guarantee. A concentration argument or direct measured geometry is needed before using it in a theorem about a particular realization.

6.4 A capacity lower bound

For any unit-column dictionary, $\text{tr } G = n$ and $\text{rank } G \leq m$. Cauchy–Schwarz applied to its nonzero eigenvalues gives

$$\text{tr}(G^2) \geq \frac{(\text{tr } G)^2}{m} = \frac{n^2}{m}. \quad (6.8)$$

Subtracting the n diagonal squares and dividing by n yields

$$\frac{1}{n} \sum_i \sum_{j \neq i} G_{ij}^2 \geq \frac{n}{m} - 1. \quad (6.9)$$

This is the familiar frame-potential/Welch-type constraint. It proves unavoidable average overlap when all features have unit norm and $n > m$. It does not upper-bound the worst row, force every pair to overlap, or determine where the optimizer places interference.

6.5 Importance, correlations and learned geometry

Importance enters total risk as $\sum_i I_i e_i$, and also influences W when the training objective is importance-weighted. To predict a boundary purely from $p, n/m, \gamma$, one must characterize the learned geometry as a function of these controls. This remains a substantive theorem target; measuring G and predicting risk conditional on it is a valuable intermediate test, not the same final claim.

If features are correlated, use conditional distributions given b_i and the nuisance covariance. In general, a single centered threshold no longer gives symmetric margins. Under recursive training, initially independent features can become dependent after one generation. Do not keep applying the independent-Bernoulli bound without checking its conditions.

Block-diagonal dictionaries containing isolated columns and antipodal pairs provide an analytically tractable family for loads from one toward two, with risk summed across blocks. They do not cover arbitrary load above two. General state enumeration and the interference bound are the tools for that extension.

Before moving on. Explain the difference between an exact boundary, a sufficient region, an average random-dictionary prediction and a theorem about learned geometry.

Project 1: clean training selects the feature geometry

7.1 The training-to-geometry question

The mathematical task in this chapter is to connect the training objective to the representation that it selects. Once that representation and its decoder are specified, Gaussian feature-detection risk can be evaluated by conditioning on the input states. These are two different steps: predicting error for a supplied code does not establish which code training chooses.

We study the smallest nonlinear setting in which sharing and feature retention can compete. Let b_1, b_2 be independent Bernoulli variables with common activation probability p , and write $q = 1 - p$. There is one latent coordinate. The tied encoder and decoder are

$$\begin{aligned} W &= (u, v), \quad u^2 + v^2 = 1, \quad h = Wb, \\ G &= W^\top W = \begin{pmatrix} u^2 & uv \\ uv & v^2 \end{pmatrix}, \quad \hat{b}_i = \max\{(Gb)_i + \beta_i, 0\}. \end{aligned} \quad (7.1)$$

The total encoder energy is one for every admissible code. Decoder biases β_1, β_2 are free parameters. For positive importance weights, clean population training minimizes

$$L_p(W, \beta) = \sum_{i=1}^2 I_i \sum_{b \in \{0,1\}^2} P_p(b) (\max\{(Gb)_i + \beta_i, 0\} - b_i)^2, \quad (7.2)$$

where the probabilities in state order 00, 01, 10, 11 are q^2, pq, pq, p^2 . This is an exact four-state population objective; no sample approximation or fitted expression enters its definition.

Unless a result explicitly states otherwise, $0 < p \leq 1/2$. General importance weights are retained for the bias and local-allocation calculations. The frequency transition later specializes to $(I_1, I_2) = (1, 1/2)$. This chapter fixes the feature-to-dimension load at two; it does not prove a theorem for arbitrary loads.

7.2 Global decoder-bias minimization at fixed geometry

For a supplied W , the two bias minimizations separate. A positive importance weight multiplies an individual bias objective and therefore does not change its minimizer. It can change the geometry selected when W is optimized jointly.

Fix a feature and abbreviate its four scores by $z_b = (Gb)_i$, its target by $y_b = b_i$, and the state probability by P_b . Its objective is

$$\ell_i(\beta) = \sum_b P_b (\max\{z_b + \beta, 0\} - y_b)^2. \quad (7.3)$$

Its ReLU breakpoints are the distinct numbers $-z_b$. Between consecutive breakpoints the active set $\mathcal{A} = \{b : z_b + \beta > 0\}$ is fixed. Expanding the squared errors gives

$$\begin{aligned}\ell_i(\beta) &= A_0\beta^2 - 2B_0\beta + C_0, \\ A_0 &= \sum_{b \in \mathcal{A}} P_b, \quad B_0 = \sum_{b \in \mathcal{A}} P_b(y_b - z_b), \\ C_0 &= \sum_{b \in \mathcal{A}} P_b(y_b - z_b)^2 + \sum_{b \notin \mathcal{A}} P_b y_b^2.\end{aligned}\tag{7.4}$$

Proposition 7.1 (Finite global bias calculation). *Evaluate every breakpoint and, on every active interval with $A_0 > 0$, the vertex B_0/A_0 clipped to that interval's closure. The least objective value among these candidates is the global fixed-geometry minimum.*

Proof. Within each interval, Equation (7.4) is a quadratic with positive leading coefficient whenever at least one state is active. Its constrained minimum is its vertex if feasible and otherwise an endpoint. An interval with $A_0 = 0$ has constant loss. The far-left inactive interval is represented by its finite boundary; the far-right quadratic is coercive. The original loss is continuous at every breakpoint because a newly active prediction is zero there. Thus the candidates include a minimum on every interval closure and omit no global minimum. $\square$

The complete scalar objective need not be convex. At activation of a state whose target is one, the derivative jumps downward by $2P_b$. This is why identifying one feasible stationary bias is insufficient without comparing the other intervals.

Global sign reversal of W leaves G unchanged. Every geometry can therefore be parametrized, modulo that sign, by

$$W(\theta) = (\cos \theta, \sin \theta), \quad -\frac{\pi}{2} \leq \theta \leq \frac{\pi}{2}, \quad F_{p,I}(\theta) = \sum_i I_i \min_{\beta_i} \ell_i(\beta_i; W(\theta)).\tag{7.5}$$

The clean geometry problem is exactly $\min_{\theta} F_{p,I}(\theta)$. Its reduction to one dimension is exact; a finite angular grid still does not establish its global minimum.

7.3 Monosemantic retention under the same budget

For $W = (1, 0)$, the optimal retained-feature bias is zero and its reconstruction is perfect. The second decoder is a constant. Its optimum is $\beta_2 = p$, since

$$\mathbb{E}[(p - b_2)^2] = p(1 - p) = pq.\tag{7.6}$$

Consequently

$$L_{\text{mono},1} = I_2 pq, \quad L_{\text{mono},2} = I_1 pq.\tag{7.7}$$

The clean-selected comparator retains the more important feature. This choice is made by the training objective, rather than selected afresh at each noise level to obtain a favorable comparison.

Binary detection is a different outcome. With the strict decision $\hat{b}_i > 1/2$, an omitted concept has error p for the declared $p \leq 1/2$, including the tie at $p = 1/2$. A small frequency can therefore make an always-absent detector appear accurate while its recall is zero.

7.4 The equal antipodal pair with trained biases

Consider the supplied equal-energy pair

$$W = \left(\frac{1}{\sqrt{2}}, -\frac{1}{\sqrt{2}} \right). \quad (7.8)$$

For either feature, the scores, targets and probabilities can be ordered as

score	0	1/2	-1/2	0
target	0	1	0	1
probability	q^2	pq	pq	p^2

On $0 \leq \beta \leq 1/2$, the negative-score state is inactive and

$$\ell(\beta) = D\beta^2 - 2N\beta + \frac{pq}{4} + p^2, \quad D = 1 - p + p^2, \quad N = \frac{p(1+p)}{2}. \quad (7.9)$$

Proposition 7.2 (Clean-optimal equal-pair bias). *For $0 < p \leq 1/2$, the globally optimal bias and per-feature loss are*

$$\beta_{\text{pair}}^* = \frac{p(1+p)}{2(1-p+p^2)}, \quad \ell_{\text{pair}}^* = \frac{pq}{4} + p^2 - \frac{p^2(1+p)^2}{4D}. \quad (7.10)$$

The weighted loss is $(I_1 + I_2)\ell_{\text{pair}}^*$.

Proof. The nonconstant bias intervals are $(-1/2, 0)$, $(0, 1/2)$ and $(1/2, \infty)$. Their stationary points are respectively $1/2$, N/D and p . For $p \leq 1/2$, $N/D \leq 1/2$, the first point is to the right of its interval, and the last is at or left of its interval. The objective decreases to N/D and then increases. The inactive interval has loss p and cannot improve the displayed solution. Substitution in Equation (7.9) gives the loss. $\square$

For completeness, when $p \geq 1/2$ the same interval comparison selects the all-active bias $\beta = p$ with loss $pq/2$ per feature. The two expressions agree at $p = 1/2$.

The actual trained decoder uses threshold

$$\vartheta = \frac{1}{2} - \beta_{\text{pair}}^* \quad (7.11)$$

on the score before adding the bias. It is not a freely supplied midpoint threshold. For $p = 1/20$, it is approximately 0.4724409449; for $p = 1/5$, it is $5/14$. In particular, rare concepts can have a small positive detection margin even when reconstruction benefits from the pair.

To state the corruption calculation precisely, add Gaussian noise $\sigma\epsilon$ to the latent coordinate, with $\epsilon \sim \mathcal{N}(0, 1)$. The marginal score-noise standard deviation is $\nu = \sigma/\sqrt{2}$. Let Φ denote the standard normal cumulative distribution function. For $\sigma > 0$, the pair's per-feature detection error is

$$\begin{aligned} e_{\text{pair}}(\sigma) &= q^2\Phi(-\vartheta/\nu) + pq\Phi((\vartheta - 1/2)/\nu) \\ &\quad + pq\Phi((-1/2 - \vartheta)/\nu) + p^2\Phi(\vartheta/\nu). \end{aligned} \quad (7.12)$$

Each term is a conditional error for one row of the four-state table. No independence between the two decoder outputs is required to add their weighted marginal risks. At zero noise, evaluate the strict threshold directly rather than divide by $\nu = 0$; the per-feature error is p^2 . For a retained mono concept the error at positive noise is $\Phi(-1/(2\sigma))$, while the omitted concept's error stays p .

7.5 Importance changes the selected amplitudes

For an opposite-sign code, let

$$u = \sqrt{a}, \quad v = -\sqrt{1-a}, \quad r = \sqrt{a(1-a)}. \quad (7.13)$$

Near $a = 1/2$, for $0 < p < 1/2$, both feature decoders have three active states. For a feature with squared column norm d , that branch has

$$\begin{aligned} \beta_3(d) &= \frac{p(1-d+pr)}{D}, \\ \ell_3(d) &= pq(1-d)^2 + p^2(1-d+r)^2 - \frac{p^2(1-d+pr)^2}{D}. \end{aligned} \quad (7.14)$$

The bias lies strictly between its relevant breakpoints at the equal pair when $0 < p < 1/2$. Thus its local differentiation is legitimate. Because $r'(1/2) = 0$, differentiation of $I_1\ell_3(a) + I_2\ell_3(1-a)$ yields

$$\left. \frac{dF}{da} \right|_{a=1/2} = -(I_1 - I_2) \frac{p(1-p)^2(1+p)}{D}. \quad (7.15)$$

If $I_1 > I_2$, moving encoder energy toward concept 1 decreases clean loss. The equal-amplitude pair is therefore not even stationary under unequal importance in this range. This is a statement about training-selected geometry, rather than merely reweighting errors for a fixed code. It does not identify the global optimum at every probability.

7.6 Exact global optima at activation probability one half

Now set $p = 1/2$. Let $a \geq 1/2$ denote the larger squared column norm, $c = 1-a$ the smaller, and $r = \sqrt{ac}$. The words stronger and weaker in this section refer to column norms, not importance weights.

7.6.1 Bias profiles in both sign sectors

For opposite signs, the strong feature has breakpoints $[-a, r-a, 0, r]$. The stationary points on successive nonempty intervals are

$$c, \quad c+r/2, \quad (2c+r)/3, \quad (c+r)/2. \quad (7.16)$$

The first two lie beyond their interval's right boundary; the third lies in $[0, r]$ because $c \leq r$; the last lies at or left of r . Therefore the third is the global minimum. The weak opposite-sign breakpoints are $[-c, 0, r-c, r]$, with successive roots

$$a, \quad a/2, \quad (2a+r)/3, \quad (a+r)/2. \quad (7.17)$$

The first three lie to the right of their intervals: use $a/2 \geq r-c$, equivalent to $1+c \geq 2r$, and $a \geq r$. The last is feasible. Hence

$$\beta_s^- = (2c+r)/3, \quad \beta_w^- = (a+r)/2. \quad (7.18)$$

For same signs, the strong breakpoints are $[-a-r, -a, -r, 0]$ and roots

$$c-r, \quad c-r/2, \quad 2(c-r)/3, \quad (c-r)/2. \quad (7.19)$$

The first two are beyond their right boundaries; the third lies in $[-r, 0]$ and the last is to its left. The weak breakpoints are $[-c - r, -r, -c, 0]$ and roots

$$a - r, \quad (a - 2r)/2, \quad 2(a - r)/3, \quad (a - r)/2. \quad (7.20)$$

Again the first three are beyond their right boundaries, using $1 + c \geq 2r$, and the last is feasible. Consequently

$$\beta_s^+ = 2(c - r)/3, \quad \beta_w^+ = (a - r)/2. \quad (7.21)$$

In both sign sectors the globally profiled feature losses are identical:

$$\ell_s = \frac{c^2 + cr + r^2}{6} = \frac{c(1 + r)}{6}, \quad \ell_w = \frac{a}{4}. \quad (7.22)$$

Substitution of the biases into the population quadratic gives these expressions; the identity $c^2 + r^2 = c$ uses $a + c = 1$. Degenerate endpoints follow by direct evaluation or continuity.

7.6.2 Equal importance: unequal norms are globally selected

Theorem 7.3 (Dense equal-importance geometry). *At $p = 1/2$, $I_1 = I_2 = 1$, the global minimum over all energy-one rank-one codes and free biases is*

$$\begin{aligned} a^* &= \frac{1 + 1/\sqrt{2}}{2}, \quad c^* = \frac{1 - 1/\sqrt{2}}{2}, \\ L^* &= \frac{1}{4} - \frac{3 - 2\sqrt{2}}{48} = 0.246425565098879 \dots \end{aligned} \quad (7.23)$$

Both sign sectors and either assignment of the column norms attain it.

Proof. Equations (7.18)–(7.22) profile every sign sector and norm ordering. The total loss is

$$F(a) = \frac{a}{4} + \frac{c(1 + r)}{6} = \frac{1}{4} + \frac{c(2r - 1)}{12}. \quad (7.24)$$

The endpoints $a = 1/2$ and $a = 1$ have loss $1/4$; every interior point has smaller loss. Maximizing $c(1 - 2r)$ gives the stationary equation $c(4a - 1) = r$. Squaring on $1/2 < a < 1$, where both sides are positive, yields

$$(2a - 1)(8a^2 - 8a + 1) = 0. \quad (7.25)$$

The unique interior solution is the stated a^* . Substitution gives the loss, and the endpoint comparison makes it global. $\square$

Equal importance does not force equal encoder-column norms. The result also shows that antipodality is not uniquely required for this dense reconstruction optimum.

7.6.3 Importance ratio two: mono retention is global

Proposition 7.4. *At $p = 1/2$ and $(I_1, I_2) = (1, 1/2)$, the globally optimal geometry retains concept 1 alone, with clean loss $1/8$.*

Proof. If the important concept has the stronger column, the weighted loss is

$$F = \frac{c(1 + r)}{6} + \frac{a}{8} = \frac{1}{8} + c \left(\frac{1 + r}{6} - \frac{1}{8} \right) \geq \frac{1}{8}, \quad (7.26)$$

with equality only at $c = 0$. If it has the weaker column, the loss is $c(1 + r)/12 + a/4 \geq a/4 \geq 1/8$, with a positive additional term at $a = 1/2$. Both sign sectors are covered by Equation (7.22). $\square$

7.7 A computer-assisted global certificate at probability one fifth

Set $p = 1/5$ and $(I_1, I_2) = (1, 1/2)$. Parametrize all finite geometries by

$$W(t) = \frac{(1, t)}{\sqrt{1+t^2}}, \quad t \in \mathbb{R}. \quad (7.27)$$

The point at infinity is retention of concept 2. In the attaining opposite-sign region $-1 < t < 0$, write

$$a = \frac{1}{1+t^2}, \quad c = \frac{t^2}{1+t^2}, \quad r = \frac{-t}{1+t^2}. \quad (7.28)$$

The important feature has three active states; the weaker feature has only states 00, 01 active. Minimizing their population quadratics gives

$$\begin{aligned} \beta_1 &= \frac{p(c+pr)}{D}, & L_1 &= pqc^2 + p^2(c+r)^2 - \frac{p^2(c+pr)^2}{D}, \\ \beta_2 &= pa, & L_2 &= pq^2a^2 + p^2. \end{aligned} \quad (7.29)$$

The weak active set is feasible when $pa \leq r - c$, or

$$t^2 + t + \frac{1}{5} \leq 0, \quad \frac{-1 - 1/\sqrt{5}}{2} \leq t \leq \frac{-1 + 1/\sqrt{5}}{2}. \quad (7.30)$$

The strong bias is feasible on this interval. Substitution into $L_1 + L_2/2$ gives

$$F(t) = \frac{905t^4 - 320t^3 + 410t^2 + 441}{5250(1+t^2)^2}, \quad (7.31)$$

$$F'(t) = \frac{8t(20t^3 + 175t^2 - 60t - 59)}{2625(1+t^2)^3}. \quad (7.32)$$

Theorem 7.5 (Certified low-frequency optimum). *The global clean optimum in this declared case is $W(t^*)$, where*

$$\begin{aligned} 20t^{*3} + 175t^{*2} - 60t^* - 59 &= 0, \\ -0.442090239023400245 < t^* < -0.442090239023400243. \end{aligned} \quad (7.33)$$

Its biases and loss are

$$\begin{aligned} \beta_1^* &= \frac{t^*(5t^* - 1)}{21(1+t^{*2})}, \quad \beta_2^* = \frac{1}{5(1+t^{*2})}, \\ 0.07775208251629429 < L^* < 0.07775208251629430. \end{aligned} \quad (7.34)$$

The global comparison is a finite exact-rational computer-assisted certificate, in addition to the explicit attaining-branch derivation.

Proof. Proposition 7.1 supplies four breakpoints and four feasible active-interval vertices for each feature at a generic geometry. The inactive constant is represented by the first breakpoint. Score order changes only at $t = -1, 0, 1$, so four regions each have $8 \times 8 = 64$ joint bias candidates. Their losses are rational functions with positive denominators proportional to $(1+t^2)^2$; bias feasibility gives polynomial inequalities of degree at most two. Their stationary equations have degree at most four.

On a feasible interval a candidate function attains its finite minimum at a stationary point or a feasibility or region endpoint. Infinite endpoints are covered by the separately solved mono

geometry. Thus isolating these finitely many roots and comparing their objective values covers every geometry and bias minimum; a local angular optimum or a numerical grid is not used as the certificate.

The exact-rational implementation enumerated 256 joint branches and 1,031 raw root candidates. After feasibility checks and duplicate removal, 265 distinct function/root candidates remained. Feasibility signs were proved by rational interval evaluation and polynomial greatest common divisors for exact shared roots. Objective values were enclosed by rational interval arithmetic. The upper enclosure for the stated winner lies below every competing lower enclosure. The smallest such competitor lower bound exceeds 0.07827150837086602; mono retention losses are 0.08 and 0.16. No competitor enclosure overlaps the winner. The displayed root is strictly inside its active-set feasibility interval. This proves the global claim. $\square$

The competing regional minimum values, rounded only for exposition, are 0.0942857143 for $t \leq -1$, 0.0777520825 for $-1 \leq t \leq 0$, 0.08 for $0 \leq t \leq 1$, and 0.1569262176 for $t \geq 1$. The certificate compares rational enclosures rather than these decimals.

The code and full candidate ledger are preserved in `project1_toy/focused_bridge_2026-10-06/math/certify_low_p.py` and `project1_toy/focused_bridge_2026-10-06/math/low_p_certificate_candidates.json`. The summary contains the exact root and objective enclosures. The independent review reproduced the winner using a different active-subset enumeration.

Despite its lower clean squared error, the selected weaker feature never passes the actual decoder's detection threshold:

$$\max_b \hat{b}_2 = c + \beta_2 = \frac{t^{*2} + 1/5}{1 + t^{*2}} < \frac{1}{2}. \quad (7.35)$$

The important feature is detected correctly in all four clean states; the weighted clean detection error is therefore $\frac{1}{2}p = 0.1$, equal to mono's clean detection error. Reconstruction improvement cannot be substituted for feature-detection improvement.

7.8 The global frequency transition at importance ratio two

Fix $(I_1, I_2) = (1, 1/2)$ while allowing $0 < p \leq 1/2$.

Theorem 7.6 (Clean geometry-selection transition). *Let*

$$p_c = \frac{3 - \sqrt{5}}{2} = 0.3819660112501051518\dots \quad (7.36)$$

For $p_c \leq p \leq 1/2$, retention of concept 1 alone is the globally optimal geometry, with biases $(0, p)$ and loss $pq/2$. Every mixed code has strictly greater loss. For $0 < p < p_c$, a mixed opposite-sign code strictly improves on mono, so every global minimizer has both columns nonzero. This theorem does not assert a unique explicit sharing geometry at every $p < p_c$.

The proof requires global comparison, not merely a change of local stability. We give the complete bias profiles and then compare their losses. Let the larger squared norm be $a \geq 1/2$, set $c = 1 - a$, $r = \sqrt{ac}$, and introduce

$$k = \sqrt{c/a} \in [0, 1], \quad a = (1 + k^2)^{-1}, \quad c = k^2(1 + k^2)^{-1}, \quad r = k(1 + k^2)^{-1}. \quad (7.37)$$

7.8.1 Excluding all same-sign sharing

The strong feature has breakpoints $[-a-r, -a, -r, 0]$. Its successive stationary biases at general p are

$$c-r, \quad c-pr, \quad \frac{c-r}{2-p}, \quad p(c-r). \quad (7.38)$$

The first two lie past the right boundaries, the third lies in $[-r, 0]$, and the last is to its left. The global strong-feature solution is

$$\beta_s^+ = \frac{c-r}{2-p}, \quad S_+ = \frac{pq c(1+2qr)}{2-p}. \quad (7.39)$$

For the weak feature, any bias $\beta \leq -c$ misses its single-feature state and has loss at least pq . The all-active candidate $\beta_w^+ = p(a-r) \geq 0$ has loss $W_+ = pqa \leq pq$. For $\beta \geq -c$, the three-active vertex $(a-r)/(2-p)$ lies at or beyond zero, so that interval decreases toward the all-active interval. The global weak solution is therefore

$$\beta_w^+ = p(a-r), \quad W_+ = pqa. \quad (7.40)$$

This comparison covers possible inferior negative-bias local minima rather than discarding them by a convexity assumption.

Since $2r \leq 1$, $S_+ \leq pqc \leq W_+$. Assigning the larger importance to the stronger column is therefore the better orientation. Even it has

$$S_+ + \frac{1}{2}W_+ - \frac{1}{2}pq = \frac{pq c(p/2 + 2qr)}{2-p} > 0 \quad (7.41)$$

for every mixed code. Thus same-sign sharing never improves mono in this importance setting, for the full declared probability range.

7.8.2 Opposite-sign strong-feature profile

The strong feature's breakpoints are $[-a, r-a, 0, r]$. The interval roots are

$$c, \quad c+pr, \quad \frac{p(c+pr)}{D}, \quad p(c+r). \quad (7.42)$$

The first two lie to the right of their intervals. The third lies in $[0, r]$ because $pc \leq qr$ follows from $p \leq q$ and $c \leq r$. The last root lies at or left of r by the same inequality. Hence

$$\begin{aligned} \beta_s^- &= \frac{p(c+pr)}{D}, \\ S_- &= pqc - \frac{pq}{D}(qr-pc)^2 = \frac{pqc}{D} [p + (1-2p)c + 2pqr]. \end{aligned} \quad (7.43)$$

The second equality follows by expanding and using $r^2 = ac$, $a+c=1$.

7.8.3 Every competing weak-feature bias branch

The weak feature's breakpoints are $[-c, 0, r-c, r]$. All possible global minima are among the following feasible candidates:

$$\beta_{w,2} = pa, \quad W_2 = pq^2a^2 + p^2, \quad (7.44)$$

$$\beta_{w,3} = \frac{p(a+pr)}{D}, \quad W_3 = pqa - \frac{pq}{D}(qr-pa)^2, \quad (7.45)$$

$$\beta_{w,4} = p(a+r), \quad W_4 = pqa. \quad (7.46)$$

The two-active branch in Equation (7.44) is feasible when $p \leq k - k^2$; the all-active branch in Equation (7.46) is feasible when $k \leq p/q$. The three-active branch is feasible when $p + Dk^2 - qk \geq 0$ and $k \geq p/q$. It also equals

$$W_3 = \frac{pqa}{D} [p + (1 - 2p)a + 2pqr]. \quad (7.47)$$

Here is why the candidate list is complete. The first interval's single target-one stationary bias lies beyond its right boundary. The two-active quadratic has vertex pa ; the three-active and all-active vertices are as displayed. A target-one activation has a strictly downward derivative jump and cannot be a minimum at its kink. Target-zero activation has a continuous derivative and its possible minimum is covered by adjacent interval vertices on their closures. The far-left inactive loss is p ; for $c > 0$, the feasible bias zero has loss $pqa^2 + p^2 < p$, and for a dropped column the optimal constant has loss $pq < p$. Thus the constant cannot win. At low probabilities two- and three-active minima can coexist; taking only one of them without comparison would be unjustified.

For the upper-phase proof, $p \geq p_c > 1/4$. Equation (7.44) is then impossible, since $k - k^2 \leq 1/4$. For $k \geq p/q$, the remaining three-active lower-feasibility condition follows from

$$pa + Dc - qr = [pa - (r - c)] + p(r - qc) > 0. \quad (7.48)$$

The first bracket is positive because the excluded two-active branch is infeasible; the second is nonnegative. Therefore above p_c the weak profile is precisely W_4 for $k \leq p/q$ and W_3 for $k \geq p/q$.

Both profiles are at least the strong loss. For W_4 , use $S_- \leq pqc \leq pqa$. For W_3 , subtraction gives

$$W_3 - S_- = \frac{pq(a - c)}{D} [q + 2pqr] \geq 0. \quad (7.49)$$

Thus assigning importance one to the stronger column is the better orientation in the upper phase; the reverse assignment is covered without assuming it away.

7.8.4 Global comparison above the threshold

In the all-active weak branch, the mixed-code difference from mono is

$$\Delta_A = S_- + \frac{1}{2}pqa - \frac{1}{2}pq = \frac{pqc}{D} \left[p - \frac{D}{2} + (1 - 2p)c + 2pqr \right]. \quad (7.50)$$

For $p_c \leq p \leq 1/2$, all bracketed terms are nonnegative because $p - D/2 = -(p^2 - 3p + 1)/2 \geq 0$. For every mixed code, the bracket is strictly positive, including at $p = p_c$.

For the three-active weak branch, direct substitution gives

$$\Delta_3 = \frac{pq}{2D(1 + k^2)^2} H_p(k), \quad (7.51)$$

$$H_p(k) = (1 - p - p^2)k^4 + 4pqk^3 + (-2 + 5p - 2p^2)k^2 + 2pqk - p^2. \quad (7.52)$$

Set $A = 1 - p - p^2$ and $B = -2 + 5p - 2p^2$. On $[p_c, 1/2]$, $A \geq 1/4$ and $B + p = -2(p^2 - 3p + 1) \geq 0$. For $k \geq p/q$,

$$H'_p(k) = 4Ak^3 + 2pk(6qk - 1) + 2pq + 2(B + p)k > 0, \quad (7.53)$$

because $qk \geq p \geq p_c > 1/6$. At $k = p/q$, the weak branches agree and H_p is positive by Equation (7.50). It remains positive throughout the three-active branch. This covers every mixed opposite-sign code above the threshold. Same-sign codes and importance swaps have already been covered, while retention of the unimportant concept has loss $pq > pq/2$. Mono is therefore globally optimal, not merely locally stable.

7.8.5 Strict sharing improvement below the threshold

For a sufficiently small opposite-sign perturbation of important-concept mono, $k > 0$ lies in the weak all-active branch. Since $c = k^2 + O(k^4)$ and $r = k + O(k^3)$, Equation (7.50) gives

$$\Delta_A = pq \left(\frac{p}{D} - \frac{1}{2} \right) k^2 + O(k^3). \quad (7.54)$$

For $p < p_c$ the leading coefficient is strictly negative, so a sufficiently small mixed code has lower clean loss than mono. A global optimum exists: the energy-one code circle is compact and score magnitudes are at most $\sqrt{2}$; a minimizing bias can be restricted to $[-3, 3]$ without losing an optimum. Loss is continuous on this compact parameter set. Therefore the global optimum cannot be mono below p_c , and cannot be same-sign by Equation (7.41). This completes the proof of Theorem 7.6.

7.9 What is established and what remains a research question

The clean allocation threshold is global: a competing off-axis optimum does not preempt the local stability threshold. In the sparsity convention $s = 1 - p$, sharing is selected for $s > (\sqrt{5} - 1)/2$, while important-concept retention is selected at and below that sparsity value.

This is a clean-training phase, not automatically a robustness phase. Corruption risk must be calculated using the selected W and its actual trained biases. The exact fixed-state Gaussian expression

$$e_i(\sigma) = \sum_b P_p(b) \Phi \left[-\frac{(2b_i - 1)((Gb)_i + \beta_i - 1/2)}{\sigma \sqrt{G_{ii}}} \right], \quad \sigma > 0, \quad (7.55)$$

applies to stored features; a dropped feature is evaluated as a constant decision. The following corruption analysis uses this distinction to connect selected geometry to operational errors rather than infer robustness from a clean reconstruction label.

The finite-bias and dense-case formulas were independently reviewed. The low-frequency certificate used exact root isolation and received a separate independent active-subset check. The frequency proof's identities are recorded in `project1_toy/frequency_boundary_2026-10-06/math/verification_results.json`; an independent review checked the population quadratics and branch coverage. These checks substantiate the declared toy mathematics. They do not by themselves certify research novelty, prove a variable-load law, demonstrate real-model transfer, or guarantee conference acceptance.

Project 1: robustness of the clean-selected representation

Learning objective. Evaluate the decoder selected by clean reconstruction, and prove exactly what its noise comparison establishes.

8.1 A fixed evaluation rule after clean training

Let the independent bits b_1, b_2 have common activation probability p , put $q = 1 - p$, and retain the encoder-energy constraint $\|W\|_F^2 = 1$. The clean objective is weighted squared reconstruction with weights $(1, 1/2)$. Once W and β have been selected by that objective, they remain fixed during noise evaluation. For scalar code noise $\xi \sim \mathcal{N}(0, \sigma^2)$, independent of the bits,

$$h = Wb + \xi, \quad \hat{b}_i^{\text{amp}} = \text{ReLU}((Gb)_i + \beta_i + w_i \xi), \quad G = W^\top W. \quad (8.1)$$

The evaluated binary decision is the actual trained decoder's strict rule

$$\hat{b}_i = \mathbf{1}\{\hat{b}_i^{\text{amp}} > 1/2\} = \mathbf{1}\{(Gb)_i + \beta_i + w_i \xi > 1/2\}. \quad (8.2)$$

The equality follows because the threshold is positive: ReLU exceeds $1/2$ exactly when its preactivation does. The weighted detection error is

$$R(W, \beta; p, \sigma) = r_1 + \frac{1}{2}r_2, \quad r_i = \mathbb{P}\{\hat{b}_i \neq b_i\}. \quad (8.3)$$

Neither the detector nor the mono comparator is reselected to improve noisy performance. This is average Gaussian code-noise robustness of a specified decision rule; it is not a worst-case input-attack theorem, a Bayes-optimal detector comparison, or an AI-safety evaluation.

8.2 Deriving the exact four-state Gaussian risk

Define the clean threshold gap $m_i(b) = (Gb)_i + \beta_i - 1/2$. Conditional on b , $w_i \xi$ has zero mean and variance $\sigma^2 w_i^2 = \sigma^2 G_{ii}$. If $w_i \neq 0$ and $\sigma > 0$, the two conditional errors are

$$\mathbb{P}(\hat{b}_i = 1 \mid b_i = 0, b_{-i}) = \Phi\left(\frac{m_i(b)}{\sigma|w_i|}\right), \quad (8.4)$$

$$\mathbb{P}(\hat{b}_i = 0 \mid b_i = 1, b_{-i}) = \Phi\left(\frac{-m_i(b)}{\sigma|w_i|}\right). \quad (8.5)$$

For the first identity, use $\mathbb{P}(\xi' > -m) = \Phi(m/s)$ for $\xi' \sim \mathcal{N}(0, s^2)$. For the second, use $\mathbb{P}(\xi' \leq -m) = \Phi(-m/s)$. Both are captured by

$$r_i(p, \sigma) = \sum_{b \in \{0,1\}^2} P_p(b) \Phi\left(\frac{(1 - 2b_i)m_i(b)}{\sigma|w_i|}\right). \quad (8.6)$$

The two output errors can be correlated through their shared scalar noise. Independence of those errors is unnecessary: linearity of expectation gives $R = r_1 + r_2/2$ regardless. For $w_i = 0$ or $\sigma = 0$, evaluate the deterministic decision directly instead of dividing by zero.

For an opposite-sign encoder $W = (1, t)/\sqrt{1 + t^2}$ with $t < 0$, write $a = 1/(1 + t^2)$, $c = t^2/(1 + t^2)$, and $d = t/(1 + t^2) < 0$. The complete state enumeration is

State b	Probability	$m_1(b)$	$m_2(b)$
00	q^2	$\beta_1 - 1/2$	$\beta_2 - 1/2$
10	pq	$a + \beta_1 - 1/2$	$d + \beta_2 - 1/2$
01	pq	$d + \beta_1 - 1/2$	$c + \beta_2 - 1/2$
11	p^2	$a + d + \beta_1 - 1/2$	$c + d + \beta_2 - 1/2$

This table supplies all eight state/feature CDF terms. It makes explicit how negative overlap can suppress false positives while simultaneously cancelling a coactive concept's positive signal.

8.3 The clean-selected mono control

For $0 < p \leq 1/2$, mono retention of the important concept is $W = (1, 0)$ and $\beta = (0, p)$. The retained concept's threshold gaps are $-1/2$ and $+1/2$, so its error is $\Phi(-1/(2\sigma))$ for either class. The omitted concept has the constant amplitude prediction p and is never declared present under the strict rule, including the tie at $p = 1/2$. Consequently

$$\boxed{R_{\text{mono}}(p, \sigma) = \Phi\left(-\frac{1}{2\sigma}\right) + \frac{p}{2}, \quad \sigma > 0,} \quad R_{\text{mono}}(p, 0) = \frac{p}{2}. \quad (8.7)$$

The retained-feature error's lack of dependence on p is a property of its supplied symmetric threshold gaps. A detector optimized separately for rare-class detection would generally use a different threshold.

8.4 Why reconstruction and detection can disagree

At $p = 1/5$, the globally certified clean optimum has $t_* \approx -0.4420902390$ and weighted reconstruction loss 0.0777520825163, below mono's 0.08. Its weak concept's maximum reconstruction is

$$c + \beta_2 = \frac{t_*^2 + 1/5}{1 + t_*^2} < \frac{1}{2}. \quad (8.8)$$

The inequality is checked at the isolated algebraic root, not inferred from a rounded angle. Thus the weak concept is never detected, even when present alone. The important concept is detected correctly in every clean state, and the total clean detection error equals mono's $p/2 = 0.1$. A nonzero stored column and improved amplitude estimation therefore do not imply improved concept detection. All primary positive-noise settings tested at this frequency yield greater detection error for sharing than for mono.

8.5 The infinite-noise comparison

Let the encoder be genuinely shared, meaning both $w_i \neq 0$, and keep its finite biases fixed. Every argument in Equation (8.6) tends to zero as $\sigma \rightarrow \infty$. Since $\Phi(0) = 1/2$ and the state

probabilities sum to one, $r_i \rightarrow 1/2$. Therefore

$$R_{\text{shared}}(p, \sigma) \rightarrow \frac{3}{4}, \quad (8.9)$$

$$R_{\text{mono}}(p, \sigma) \rightarrow \frac{1}{2} + \frac{p}{2}, \quad (8.10)$$

$$\boxed{\Delta(p, \sigma) := R_{\text{shared}} - R_{\text{mono}} \rightarrow \frac{1}{4} - \frac{p}{2} > 0} \quad (0 < p < 1/2). \quad (8.11)$$

A sharing advantage at a finite noise level cannot persist to arbitrarily large noise under this fixed rule. The omitted mono feature receives no code noise, whereas every nonzero shared column eventually randomizes. This asymptotic conclusion does not assert which finite noise levels favor sharing.

To derive the possible noncommutation of limits explicitly, fix $p < 1/2$ and consider the supplied family

$$W_\varepsilon = (\sqrt{1 - \varepsilon^2}, -\varepsilon), \quad \beta = (0, p), \quad 0 < \varepsilon < 1. \quad (8.12)$$

At any fixed finite $\sigma > 0$, the weak column's score and noise tend to zero as $\varepsilon \downarrow 0$, leaving a constant prediction $p < 1/2$ and weak error p . The strong column converges to mono's Gaussian detector. Hence $R_\varepsilon(p, \sigma) \rightarrow R_{\text{mono}}(p, \sigma)$. For every fixed $\varepsilon > 0$, both columns are nonzero, so $R_\varepsilon(p, \sigma) \rightarrow 3/4$ as $\sigma \rightarrow \infty$. It follows that

$$\lim_{\sigma \rightarrow \infty} \lim_{\varepsilon \downarrow 0} R_\varepsilon = \frac{1}{2} + \frac{p}{2} \neq \frac{3}{4} = \lim_{\varepsilon \downarrow 0} \lim_{\sigma \rightarrow \infty} R_\varepsilon. \quad (8.13)$$

This proves the limit distinction for a supplied family; it does not claim that every member is clean-optimal.

8.6 A rigorous example with at least two crossings

The fixed $p = 1/20$ case has an exactly certified global clean-selected mixed geometry. Its parameter is the unique root in the rational interval

$$\left[-\frac{591234926798208727}{913174160928329335}, -\frac{93541350199299932}{144476484910883523} \right] \quad (8.14)$$

of

$$380t^3 + 14800t^2 - 1140t - 6839 = 0, \quad t_* \approx -0.6474503463798862. \quad (8.15)$$

The attaining branch has biases

$$\beta_1 = \frac{20t_*^2 - t_*}{381(1 + t_*^2)}, \quad \beta_2 = \frac{1}{20(1 + t_*^2)}. \quad (8.16)$$

These follow from the three-active strong-feature and two-active weak-feature population quadratics. Global selection follows from the finite-branch certificate described in the preceding chapter; the cubic alone would establish only a stationary candidate.

To obtain the cubic explicitly, substitute $p = 1/20$, $a = 1/(1 + t^2)$, $c = t^2/(1 + t^2)$, $r = -t/(1 + t^2)$ and $D = 381/400$ into the attaining branch in Equation (7.29). Its weighted clean objective and derivative simplify to

$$F_{1/20}(t) = \frac{296420t^4 - 28880t^3 + 30440t^2 + 145161}{6096000(1 + t^2)^2}, \quad (8.17)$$

$$F'_{1/20}(t) = \frac{19t(380t^3 + 14800t^2 - 1140t - 6839)}{1524000(1 + t^2)^3}. \quad (8.18)$$

The weak two-active branch requires $p \leq -t - t^2$, and the isolated winner is strictly feasible. The same complete region/bias/root comparison used for $p = 1/5$ certifies this winner against every other branch and both mono endpoints. Thus differentiation explains the polynomial, while the separate global comparison explains why its specified root is selected.

Theorem 8.1 (Two-crossing existence for the certified rare-concept code). *For this fixed geometry, the actual-decoder weighted detection-risk difference from mono obeys*

$$\Delta(0) = \frac{1}{400} > 0, \quad \Delta(3/10) < 0, \quad \lim_{\sigma \rightarrow \infty} \Delta(\sigma) = \frac{9}{40} > 0. \quad (8.19)$$

It has at least one zero in $(0, 3/10)$ and a distinct zero in $(3/10, \infty)$.

Proof. The eight clean gaps in the state table are strictly nonzero at the root enclosure. The important concept errs only in state 11, whose probability is $p^2 = 1/400$. The weak concept is never detected, contributing weighted error $p/2 = 1/40$. Hence $R_{\text{shared}}(p, 0) = 11/400$, while $R_{\text{mono}}(p, 0) = 1/40$, and their difference is $1/400$.

At $\sigma = 3/10$, the rational CDF enclosure derived below proves a strictly negative difference. Equation (8.11) gives the limiting difference $1/4 - (1/20)/2 = 9/40$. For positive σ , the error is a finite sum of continuous CDF terms. Strict clean gaps imply that each term tends to its deterministic clean error as $\sigma \downarrow 0$, so the positive clean limit agrees with $\Delta(0)$. There is therefore a positive point below $3/10$ with positive difference and a finite point above $3/10$ with positive difference. Apply the intermediate value theorem on the two disjoint intervals with the negative value at $3/10$ as their shared endpoint. $\square$

The theorem proves existence, not exactly two roots, their locations, or a result for every frequency. In particular, a single permanent “robust phase” after one noise crossing is not an adequate description of this example.

8.7 Deriving the finite-noise rational certificate

8.7.1 From an algebraic root to interval probabilities

Represent the isolated root by a rational interval $T = [\ell, u]$. Evaluate $1 + t^2$, the Gram entries, biases, and nonzero column norms by outward interval arithmetic. For intervals $A = [a, b]$, $B = [c, d]$,

$$A + B = [a + c, b + d], \quad -A = [-b, -a], \quad (8.20)$$

$$AB = [\min(ac, ad, bc, bd), \max(ac, ad, bc, bd)]. \quad (8.21)$$

If $0 \notin B$, reciprocation maps its endpoints to $B^{-1} = [1/d, 1/c]$, and division uses multiplication by this interval. Repeated occurrences of t can widen the result but cannot invalidate containment.

For a nonnegative rational x and $S = 10^{30}$, let

$$n = \left\lfloor \sqrt{xS^2} \right\rfloor. \quad (8.22)$$

Integer square root applied to $\lfloor xS^2 \rfloor$ computes n exactly. Then $n/S \leq \sqrt{x} < (n+1)/S$, with coincident endpoints allowed for an exact grid square. Apply monotonicity at both endpoints of a nonnegative interval to enclose its square root. This gives rational bounds for $|w_i|$ and $1/\sqrt{2\pi}$ without relying on floating-point square roots.

For each state and feature, the resulting interval encloses

$$x_{i,b} = \frac{(1 - 2b_i)m_i(b)}{(3/10)|w_i|}. \quad (8.23)$$

Round its lower endpoint downward and upper endpoint upward to a 10^{-12} grid. Monotonicity of Φ reduces the remaining task to bounding Φ at two rational endpoints.

8.7.2 An exact enclosure of the normalizing constant

For $0 < x < 1$, write

$$A_{30}(x) = \sum_{j=0}^{30} \frac{(-1)^j x^{2j+1}}{2j+1}, \quad A_{31}(x) = A_{30}(x) - \frac{x^{63}}{63}. \quad (8.24)$$

The alternating-series remainder gives $A_{31}(x) \leq \arctan x \leq A_{30}(x)$. Machin's identity is

$$\pi = 16 \arctan(1/5) - 4 \arctan(1/239). \quad (8.25)$$

To justify its quadrant as well as its tangent, doubling $\tan \alpha = 1/5$ gives $\tan(2\alpha) = 5/12$ and $\tan(4\alpha) = 120/119$. Hence

$$\tan(4\alpha - \arctan(1/239)) = \frac{120/119 - 1/239}{1 + 120/(119 \cdot 239)} = 1. \quad (8.26)$$

Also $x/(1+x^2) < \arctan x < x$ for $x > 0$, by the integral definition. Thus the angle is positive and below $4/5$. Since $\arctan 1 > 1/2$, $\pi/2 = 2 \arctan 1 > 1 > 4/5$. The angle lies in $(0, \pi/2)$ and must equal $\pi/4$. Combining the rational arctangent bounds encloses π ; monotone square root and reciprocation then enclose $1/\sqrt{2\pi}$.

8.7.3 Rational upper and lower Gaussian CDF bounds

For $x \geq 0$,

$$\Phi(x) = \frac{1}{2} + \frac{1}{\sqrt{2\pi}} \int_0^x e^{-u^2/2} du. \quad (8.27)$$

Taylor's theorem at zero for e^{-v} , $v \geq 0$, gives

$$e^{-v} = \sum_{j=0}^N \frac{(-v)^j}{j!} + \frac{(-1)^{N+1} e^{-\eta}}{(N+1)!} v^{N+1}, \quad 0 \leq \eta \leq v. \quad (8.28)$$

An even N therefore gives an upper bound and the next odd degree a lower bound, regardless of whether the initial term magnitudes decrease. Substituting $v = u^2/2$ and integrating term by term yields the rational polynomial

$$J_N(x) = \sum_{j=0}^N \frac{(-1)^j x^{2j+1}}{2^j j! (2j+1)}, \quad J_{81}(x) \leq \int_0^x e^{-u^2/2} du \leq J_{80}(x). \quad (8.29)$$

Combine this interval with the normalizing-constant interval. Negative arguments use $\Phi(-x) = 1 - \Phi(x)$; intersecting a valid enclosure with $[0, 1]$ preserves validity. Every operation uses rational arithmetic.

8.7.4 The certified sign and what is independently checked

Multiply each enclosed CDF by the exact state probability and feature weight and sum the eight contributions. Evaluate the mono term at argument $-5/3$ with the same CDF enclosure and add $p/2$. Interval subtraction yields a wholly negative enclosure for $\Delta(3/10)$, with displayed approximate endpoints

$$[-0.014474506712828471, -0.014474506712749220]. \quad (8.30)$$

The decimal endpoints summarize exact rational endpoints saved in the certificate; the sign decision is made using those rational endpoints. An independent 100-digit reference value is approximately $-0.0144745067127706407821$ and lies inside the enclosure. That reference is a numerical crosscheck, not the sign proof.

Independent review checked the root and bias mapping, all eight strict clean gaps, the outward operations, Taylor and arctangent recurrences, and the continuity argument. Fresh reproduction in a separate directory matched the saved certificate except runtime. The full rational per-state records are in `project1_toy/frequency_boundary_2026-10-06/math/noise_sign_certificate_results.json`; the implementation is `certify_noise_sign.py` in that directory.

8.8 The fixed numerical checks and their interpretation

Five frequencies were declared before the verification: .05, .20, .35, .375, .40. The first four use global population certificates; the last uses the analytic mono theorem. Selected clean losses and weak-column energies are

p	Weak-column energy	Selected MSE	Mono MSE
.05	.295373681287	.017742880541	.023750000000
.20	.163490565429	.077752082516	.080000000000
.35	.002595127216	.113740582191	.113750000000
.375	.000121567823	.117187402718	.117187500000
.40	0	.120000000000	.120000000000

These are discrete certified selections, not a sampled proof of the continuous transition. Near-transition MSE gains are small: approximately 9.42×10^{-6} at .35 and 9.73×10^{-8} at .375.

The original primary noise settings remain 0, .05, .15, .30, .60. The actual-decoder selected-minus-mono risk differences are

p	0	.05	.15	.30	.60
.05	.002500	.002500	.003472	-.014475	-.003661
.20	0	.012228	.017397	.011652	.007769
.35	0	≈ 0	.000053	-.000144	-.000863
.375	0	≈ 0	-.000001	-.000135	-.000223
.40	0	0	0	0	0

Negative values favor sharing. Approximate zeros can reflect tiny Gaussian tails or finite-precision CDF saturation, so they do not certify exact equality. At .40 the selected geometry is exactly mono, so equality holds independently of floating-point evaluation.

A prespecified 240-point logarithmic display grid from noise .001 to 4 and Brent refinement of reliable adjacent sign changes gave numerical roots

p	Lower crossing	Upper crossing
.05	.2079551941	.6252753374
.35	.2698862926	1.0932765901
.375	.1277396994	3.3971473354

No crossing was found for .20 in the display interval. This is not a proof of no crossing at all noise strengths, nor a complete root count for the other cases. The separate rational theorem proves existence of at least two roots only at .05.

8.9 Verification failures and stopping scope

The preceding projected-Adam campaign preserved all 36 final iterates. Twenty-seven passed its loss-stability criterion and nine failed; passing loss stability did not ensure stationarity, and several iterates retained substantial tangent gradients. The separate two-case angular solver reached gradient tolerance at the already specified .05 and .20 unequal-importance cases. Later global certificates establish population optima, not a guarantee that every solver or initialization learns them. The frozen training files were not replaced by the certificate outputs.

The independent numerical review checked 48 raw hashes, 150 primary risk rows, 1,200 display values, and six numerical root brackets. It recovered the four sharing optima using a different 225-active-subset construction, while explicitly sharing the reviewed rational-root backend. Maximum detection and display discrepancies were below 1.4×10^{-16} .

The initial independent MSE checker failed because unbounded quadrature missed central Gaussian mass when a weak-feature ReLU cutoff lay far into a tail. Its source and discrepancies remain archived. The repaired checker integrates over $[-12, 12]$, splitting at zero and any relevant ReLU cutoff, and bounds the omitted tails. If a conditional preactivation is $\mu + sZ$ and the target bit is b , then

$$|\text{ReLU}(\mu + sZ) - b|^2 \leq 2(|\mu| + 1)^2 + 2s^2Z^2. \quad (8.31)$$

For $A > 0$, Mills' bound and integration by parts give

$$Q(A) \leq \frac{\phi(A)}{A}, \quad \mathbb{E}[Z^2 \mathbf{1}\{|Z| > A\}] = 2\{A\phi(A) + Q(A)\}. \quad (8.32)$$

Therefore the omitted two-sided contribution is bounded by

$$4(|\mu| + 1)^2 Q(A) + 4s^2\{A\phi(A) + Q(A)\}. \quad (8.33)$$

At $A = 12$, the largest bound over the reviewed cases was below 3.90×10^{-31} ; corrected quadrature agreed with saved MSE values to 1.39×10^{-16} . Scientific outputs and validation acceptance thresholds were not modified.

The result now connects clean training to certified geometry and operational noisy detection in one fixed-load family. It does not finish arbitrary-load or arbitrary-importance theory, establish a real-model result, or certify scientific novelty. Further toy work should resolve a specific prediction or transfer blocker, rather than accumulate more special-case proofs. The two global analytic statements, their computer-assisted examples, and the unfavorable detection outcomes are the evidence to carry forward together.

8.10 The remaining paper-critical decisions

The completed mechanism is clean population reconstruction objective $\rightarrow$ selected geometry $\rightarrow$ fixed decoder $\rightarrow$ noisy detection. The research direction remains the sparsity/load/importance

tradeoff, but the proved transition fixes load at two and importance ratio at two. Independent pair-block replication would merely add their risks; it would not derive a new load-dependent crossing or prove global optimality among unrestricted larger codes.

The motivating monosemanticity paper uses different learning objectives, continuous toy amplitudes, and corruption locations. This binary, fixed-energy, code-noise calculation is a controlled extension of that motivation, not an automatic strict generalization of all its results. A focused comparison of these exact assumptions and claims with direct prior theory remains necessary. Verification does not settle novelty.

The next real-model test must specify the feature measurement, corruption location, decoder/outcome and comparator before substantial compute is allocated. An SAE dictionary Gram matrix is not automatically the tied autoencoder map in this chapter. A transfer experiment may test the qualitative prediction that clean reconstruction selection and downstream corrupted detection need not agree, without claiming that the toy CDF formula is exact in a diffusion network. New root-count proofs, additional noise families, or recursive diffusion retraining are not needed to finish the current bounded mathematical task.

8.11 Does the reconstruction reversal survive decoder calibration?

The clean-optimal representation and its clean-optimal decoder are two different objects. A corrupted input distribution may make a previously optimal decoder bias inappropriate even when its representation is fixed. To distinguish this effect from a representation cost, allow the same decoder adjustment for both shared and monosemantic codes, and compare the same continuous reconstruction loss before and after adjustment.

8.11.1 A fixed, globally solved comparison

Let b_1, b_2 be independent Bernoulli(1/2) variables, with equal feature importance. There is one code dimension and two concepts. Write $w = (w_1, w_2)$, $\|w\|_2^2 = 1$, and $G = w^\top w$. The clean-optimal shared geometry already derived is

$$a_* = \frac{1 + 1/\sqrt{2}}{2}, \quad c_* = \frac{1 - 1/\sqrt{2}}{2}, \quad r_* = \sqrt{a_* c_*}. \quad (8.34)$$

We retain both relative-sign representatives $w_+ = (\sqrt{a_*}, \sqrt{c_*})$ and $w_- = (\sqrt{a_*}, -\sqrt{c_*})$. The monosemantic comparator is $w_{\text{mono}} = (1, 0)$; exchanging the equally important concepts gives the same population outcome. Its omitted concept is estimated by its prior, not omitted from the loss.

Their clean-optimal biases, ordered as strong then weak column, are

$$\beta_{+,1} = \frac{2(c_* - r_*)}{3}, \quad \beta_{+,2} = \frac{a_* - r_*}{2}, \quad (8.35)$$

$$\beta_{-,1} = \frac{2c_* + r_*}{3}, \quad \beta_{-,2} = \frac{a_* + r_*}{2}, \quad (8.36)$$

$$\beta_{\text{mono}} = (0, 1/2). \quad (8.37)$$

The globally selected shared clean loss is $(9 + 2\sqrt{2})/48$, strictly below mono's 1/4. This is a population optimum statement, not a statement that every projected-Adam trajectory reaches it.

Freeze each encoder and add scalar Gaussian code noise:

$$h = wb + \sigma Z, \quad Z \sim \mathcal{N}(0, 1), \quad \hat{b}_i = \max\{w_i h + \beta_i, 0\}. \quad (8.38)$$

The primary outcome throughout is

$$L(w, \beta; \sigma) = \sum_{i=1}^2 \mathbb{E}[(\hat{b}_i - b_i)^2]. \quad (8.39)$$

No binary decision threshold enters this comparison. Noise is in code space, and its projected standard deviation for feature i is $s_i = \sigma|w_i|$.

8.11.2 Gaussian moments and the scalar calibration objective

Condition on a binary state b , put $z_{i,b} = (Gb)_i$, $y_{i,b} = b_i$, and $\mu_{i,b} = z_{i,b} + \beta_i$. For $s > 0$, define the first and second rectified Gaussian moments

$$M_1(\mu, s) = \mathbb{E}[(\mu + sZ)_+] = s\phi(\mu/s) + \mu\Phi(\mu/s), \quad (8.40)$$

$$M_2(\mu, s) = \mathbb{E}[(\mu + sZ)_+^2] = (\mu^2 + s^2)\Phi(\mu/s) + \mu s\phi(\mu/s). \quad (8.41)$$

Here ϕ and Φ are the standard normal density and distribution function. These formulas follow by integrating over $Z > -\mu/s$: $\int_a^\infty z\phi(z) dz = \phi(a)$ and $\int_a^\infty z^2\phi(z) dz = 1 - \Phi(a) + a\phi(a)$. At $s = 0$, use $M_1 = \mu_+$ and $M_2 = \mu_+^2$ directly.

For one feature, its scalar population loss is

$$f_i(\beta) = \sum_b P(b) \left[M_2(z_{i,b} + \beta, s_i) - 2y_{i,b}M_1(z_{i,b} + \beta, s_i) + y_{i,b}^2 \right], \quad P(b) = \frac{1}{4}. \quad (8.42)$$

Since the feature biases are independent optimization coordinates, the oracle-calibrated fixed-geometry loss is

$$L_{\text{cal}}(w; \sigma) = \min_{\beta_1} f_1(\beta_1) + \min_{\beta_2} f_2(\beta_2). \quad (8.43)$$

Oracle population calibration uses the true noisy distribution and is not a finite-data deployment procedure.

8.11.3 Derivatives and a uniform curvature bound

Differentiation of the moments gives $\partial_\mu M_1 = \Phi(\mu/s)$ and $\partial_\mu M_2 = 2M_1$. Thus, for $s_i > 0$,

$$f'_i(\beta) = 2 \sum_b P(b) [(\mu_{i,b} - y_{i,b})\Phi(\mu_{i,b}/s_i) + s_i\phi(\mu_{i,b}/s_i)], \quad (8.44)$$

$$f''_i(\beta) = 2 \sum_b P(b) \left[\Phi(\mu_{i,b}/s_i) - \frac{y_{i,b}}{s_i} \phi(\mu_{i,b}/s_i) \right]. \quad (8.45)$$

The second derivative can be negative, so a single bounded local minimizer does not establish the global optimum. Using $0 \leq y_{i,b} \leq 1$, $\Phi \leq 1$, and $\phi \leq 1/\sqrt{2\pi}$ yields the globally valid bound

$$|f''_i(\beta)| \leq H_i := 2 \left(1 + \frac{1}{s_i\sqrt{2\pi}} \right). \quad (8.46)$$

This deliberately conservative constant applies at every finite bias.

For an interval with midpoint m and radius q , Taylor's theorem and the curvature bound imply

$$f_i(\beta) \geq f_i(m) + f'_i(m)(\beta - m) - \frac{H_i}{2}(\beta - m)^2 \geq f_i(m) - |f'_i(m)|q - \frac{H_i q^2}{2}. \quad (8.47)$$

The second inequality minimizes the concave quadratic lower expression over the interval, whose minimum lies at an endpoint. Nonnegativity of MSE supplies another lower bound, zero.

8.11.4 Explicitly excluding the two infinite tails

Let $U_i = 1 - \min_b z_{i,b}$. For $\beta \geq U_i$ all conditional means satisfy $\mu_{i,b} \geq 1 \geq y_{i,b}$. Every term in $f'_i(\beta)$ is positive when $s_i > 0$. Hence the right half-line has minimum at U_i .

For the left half-line choose, before optimization,

$$B_i = -\max_b z_{i,b} - 12s_i - 1. \quad (8.48)$$

The rectified first moment is nondecreasing in its mean, while the second moment is nonnegative. Since $y_{i,b}^2 = y_{i,b}$ for binary labels, for every $\beta \leq B_i$,

$$f_i(\beta) \geq p_i - 2 \sum_b P(b) y_{i,b} M_1(z_{i,b} + B_i, s_i), \quad p_i = \sum_b P(b) y_{i,b} = \frac{1}{2}. \quad (8.49)$$

The selected finite boundary is useful only if this lower bound exceeds an actually attainable loss. That comparison is checked and saved; it is not assumed from the size of the multiplier 12. If either exclusion fails, the prescribed control remains unresolved rather than silently enlarging its search.

For an exactly zero encoder column, $z_{i,b} = s_i = 0$, so its decoder is constant. The exact optimum is $\beta_i = p_i$ with loss $p_i(1 - p_i)$. No noisy scalar solver is needed for that column.

8.11.5 The declared numerical branch-and-bound calculation

A bounded local minimizer, the clean bias, and the two finite boundary values provide feasible incumbent losses. Local minimization is used only to obtain an upper bound. The interval $[B_i, U_i]$ is then partitioned adaptively, always expanding an interval with the smallest current lower bound from Equation (8.47). Each midpoint is a new feasible incumbent candidate; intervals whose lower bounds exceed the best upper bound are discarded. Both excluded-tail bounds remain in the global comparison.

The declared floating-point slack is 10^{-12} : numerical lower bounds are reduced and feasible upper values increased by that amount. Stop when the numerical global upper-minus-lower gap is at most 10^{-8} per feature, or after 100,000 interval expansions. An unresolved gap or excluded-tail check is preserved as a failure. Every interval, midpoint loss and gradient, lower bound, pruning action, remaining interval and boundary comparison is archived.

What the numerical bound does and does not mean. The floating slack is not outward-rounded interval arithmetic and does not make this an exact-rational noisy-optimum theorem. It is a numerical global-gap diagnostic, accompanied by independent gradient and Gaussian-quadrature checks. The sign of a model comparison must be separated from the sum of its global gaps and integration discrepancies before being interpreted as evidence.

8.11.6 Why both clean-optimal relative signs have the same profile

For the positive-sign code, a feature-1 score is $a_*b_1 + r_*b_2 + \beta_1$. Replace the other fair bit by $b'_2 = 1 - b_2$; the score becomes $a_*b_1 - r_*b'_2 + (\beta_1 + r_*)$, with target b_1 unchanged. Likewise replacing b_1 by $1 - b_1$ in the feature-2 score converts $r_*b_1 + c_*b_2 + \beta_2$ into $-r_*b'_1 + c_*b_2 + (\beta_2 + r_*)$, with target b_2 unchanged. These transformations preserve the marginal conditional population, and Gaussian projected noise has the same variance. Consequently, feature by feature,

$$f_{i,-}(\beta + r_*) = f_{i,+}(\beta). \quad (8.50)$$

The clean and globally calibrated summed losses therefore agree for both signs. The two transformations need not be applied as one joint state transformation: the outcome is a sum of marginal feature losses. Exchanging equally important features similarly preserves the total.

8.11.7 Prespecified outcome

Only $\sigma = 0, .30, .60$ were evaluated in this control. The zero-noise row is the clean analytical anchor; the two positive levels were already present in the frozen-decoder record. The following values are summed population reconstruction MSE:

Code and decoder policy	0	.30	.60
Sharing, frozen clean biases	.246425565099	.318682818682	.512993808968
Mono, frozen clean biases	.250000000000	.317463880602	.505449771089
Sharing, oracle calibrated	.246425565099	.317930513671	.491134733704
Mono, oracle calibrated	.250000000000	.312505410963	.467092806999
Calibrated sharing minus mono	-.003574434901	+.005425102707	+.024041926705

Both sign variants of sharing agree to floating precision. All ten nonzero-column scalar calibrations meet the per-feature gap target; each required 47–58 interval expansions. The two zero-column calibrations are exact. Summed numerical upper-minus-lower gaps are:

Code	$\sigma = .30$	$\sigma = .60$
Same-sign sharing	$1.142385357 \cdot 10^{-8}$	$1.354560977 \cdot 10^{-8}$
Opposite-sign sharing	$1.142385345 \cdot 10^{-8}$	$1.354560977 \cdot 10^{-8}$
Mono	$6.987186149 \cdot 10^{-9}$	$9.931263173 \cdot 10^{-9}$

The sum gap for two features may exceed 10^{-8} because the stopping target was declared per feature. Twenty analytical-gradient versus central-difference checks have maximum discrepancy $9.1433 \cdot 10^{-11}$. All excluded-region checks pass. The independently integrated risk and tail verification is recorded separately from these student global-gap results.

The independent checker imports neither the Gaussian-moment implementation nor the calibration solver. It integrates each conditional squared loss directly over $Z \in [-12, 12]$, splitting at zero and any rectifier cutoff inside that interval, and bounds the omitted tails by the inequality derived earlier in this chapter. All 85 archived hashes match. Its 36 per-feature frozen/calibrated loss checks differ from the saved values by at most 1.943×10^{-16} ; the largest omitted-tail bound is 3.913×10^{-31} . It also checks the curvature constants, interval lower bounds, excluded-region inequalities and final active-interval global bounds. These checks independently corroborate the calculation; they do not convert the floating-point optimization bounds into an exact interval proof.

The resulting numerical comparison brackets are

$$.0054250912846 \leq L_{\text{cal}}(w_+; .30) - L_{\text{cal}}(w_{\text{mono}}; .30) \leq .0054251096957, \quad (8.51)$$

$$.0240419131602 \leq L_{\text{cal}}(w_+; .60) - L_{\text{cal}}(w_{\text{mono}}; .60) \leq .0240419366371. \quad (8.52)$$

The numerical sign margins are much larger than the optimization gaps and integration discrepancies.

Oracle calibration improves both representations, but mono improves more. The same-outcome noisy disadvantage of sharing therefore survives this control rather than disappearing. This supports a cost within the fixed geometry and tied-ReLU decoder class; it does not prove robustness against every possible readout.

Scope of the result. No encoder was retrained for corruption, no arbitrary-load law was established, and no real-model transfer is represented by this oracle population exercise. The monosemantic code estimates its dropped concept by its prior, with that coverage loss included in the MSE. The result is neither a novelty certificate nor a conference acceptance prediction. Its next use is a reviewed mechanism claim and, only under a separate protocol, one focused real-model test.

8.11.8 A linear reference separates generic overlap from rectification

Removing the rectifier changes the decoder class. In that reference class allow a mean-correcting intercept $d = (I - G)\mathbb{E}[b]$ and predict $\hat{b} = Gb + w^\top \sigma Z + d$. Since $\|w\|^2 = 1$, $G^2 = G$ and $\text{tr } G = 1$; also $\text{Cov}(b) = I/4$. Independence and zero-mean noise eliminate the cross term, giving

$$L_{\text{lin}} = \text{tr}\{(I - G) \text{Cov}(b)(I - G)^\top\} + \sigma^2 \|w\|^2 \quad (8.53)$$

$$= \frac{1}{4} \text{tr}(I - G)^2 + \sigma^2 = \frac{1}{4} + \sigma^2. \quad (8.54)$$

Thus all energy-one rank-one geometries tie in this linear dense reference. Rectification is necessary for the difference in the specified nonlinear comparison. This is analytical context, not another experimental arm and not proof that rectification uniquely explains every contribution to the nonlinear difference. An arbitrary freely learned decoder is outside the present comparison.




Figure 8.1: Summed population reconstruction MSE at the three prescribed code-noise settings. Sharing has a clean advantage but a corrupted disadvantage under both frozen and oracle-calibrated biases. Both sharing signs give the same population loss. Lines connect saved points for readability; their intersections do not estimate a crossing location or establish a complete robustness boundary.




Figure 8.2: Certified clean-selected geometries at five frequencies. Left: encoder energy allocated to the weaker concept. Right: selected and mono clean weighted reconstruction loss. The dashed line is the analytically proved clean-storage transition. The discrete markers do not supply a continuous formula for the selected sharing geometry.




Figure 8.3: Actual-decoder selected-minus-mono weighted detection error versus Gaussian code-noise standard deviation. Negative values favor sharing. Panels use different vertical scales. Dotted lines are saved numerical crossings; they do not certify exact locations or complete counts. The .40 panel is identically zero because both compared codes are mono.

8.12 A joint storage and robustness boundary near the clean transition

The outcome in this section is weighted reconstruction MSE, not detection error. The two Bernoulli features are independent with activation probability p , weights $(1, 1/2)$, and one code coordinate of total encoder energy one. Training minimizes clean population loss with tied ReLU reconstruction and free decoder biases. Testing adds independent Gaussian noise of standard deviation σ to the code. Geometry is frozen at its clean-selected value. These assumptions are essential: the result is not an adversarial guarantee.

Write $q = 1 - p$, $D = 1 - p + p^2$ and

$$p_c = \frac{3 - \sqrt{5}}{2}, \quad \epsilon = p_c - p > 0.$$

The preceding global clean-selection theorem gives a unique mono geometry at p_c , up to the irrelevant code sign. Below p_c , the optimum shares opposite-sign directions. Parameterize the nearby branch by

$$a = \frac{1}{1 + k^2}, \quad c = \frac{k^2}{1 + k^2}, \quad r = \frac{k}{1 + k^2}, \quad W = (\sqrt{a}, -\sqrt{c}).$$

The strong-feature clean bias is $p(c + pr)/D$; the weak-feature bias is $p(a + r)$. The latter output is active on every state for sufficiently small k , since its branch condition is $k \leq p/q$.

8.12.1 Deriving the clean critical law

Subtract the matched mono loss $pq/2$ from the exact branch loss:

$$\Delta_0(k, p) = \frac{pqc}{D} \left[p - \frac{D}{2} + (1 - 2p)c + 2pqr \right]. \quad (8.55)$$

Let $A(p) = p - D/2$. Since $A(p_c) = 0$ and $A'(p_c) = \sqrt{5}/2$,

$$A(p_c - \epsilon) = -\frac{\sqrt{5}}{2}\epsilon - \frac{1}{2}\epsilon^2, \quad \Delta_0 = \frac{pq}{D} \{ A(p)k^2 + 2pqk^3 + O(k^4) \}.$$

For $k > 0$, divide the stationarity equation by k :

$$0 = 2A(p) + 6pqk + O(k^2).$$

Its derivative in k at $(0, p_c)$ is $6p_c q_c > 0$. The implicit function theorem therefore gives the small stationary solution

$$k_*(p) = K\epsilon + O(\epsilon^2), \quad K = \frac{\sqrt{5}}{6p_c q_c}. \quad (8.56)$$

This local calculation alone would not identify the trained optimum. The global clean theorem supplies the missing step: compactness of the energy-constrained geometry and continuity of its bias-minimized loss imply that every minimizing sequence as $p \uparrow p_c$ approaches the unique mono optimum. Thus the important feature is the strong column and the minimizer enters the stated branch. In that neighborhood the nonzero stationary solution is unique and is a minimum; comparison with the endpoint selects it because its loss difference is negative. Consequently

$$c_* = K^2 \epsilon^2 + O(\epsilon^3), \quad \Delta_0^* = -C\epsilon^3 + O(\epsilon^4), \quad C = \frac{5\sqrt{5}}{216D_c p_c q_c}. \quad (8.57)$$

The quadratic disappearance of weak storage and cubic disappearance of the clean sharing advantage are different quantities.

8.12.2 Frozen-bias corrupted risk and its boundary

Consider $\sigma = t\epsilon^{3/2}$ with bounded t . In the selected sharing code, the strong output's two zero-target preactivations have gate distances of order $k_* = \Theta(\epsilon)$; the two unit-target preactivations stay bounded away from zero. The weak output is positive on all four states. Gaussian gate-switching corrections are bounded by a polynomial prefactor times $\exp(-d/\epsilon)$ for fixed bounded $t > 0$, and vanish at $t = 0$. The strong active-state probability is $q^2 + p = D$. Hence the sharing noise increment, with gates held in their clean pattern, is $(Da + c/2)\sigma^2$.

For mono, the retained output has bias zero. On target zero, the half-Gaussian second moment is exactly $\sigma^2/2$. On target one, the contribution is σ^2 up to exponentially small clipping corrections. The omitted output remains the constant p . Mono's increment is $(p + q/2)\sigma^2$. Subtracting gives

$$\Delta_{\text{frozen}} = -C\epsilon^3 + B\sigma^2 + o(\epsilon^3), \quad B = D_c - \frac{1 + p_c}{2} = \frac{q_c(1 - 2p_c)}{2} > 0. \quad (8.58)$$

Uniformly on compact t intervals, division by ϵ^3 converges to $-C + Bt^2$. Thus any fixed brackets on either side of $\sqrt{C/B}$ have opposite signs for sufficiently small ϵ . Continuity implies a crossing in those brackets, with asymptotic location

$$\sigma_{\times} = \sqrt{C/B} \epsilon^{3/2} [1 + o(1)]. \quad (8.59)$$

This asserts a local boundary, not the unique crossing over all noise levels or a uniform finite- ϵ error bound.

8.12.3 A second frozen-bias ordering reversal

To see why a single monotone boundary cannot be assumed, examine $\sigma = k_*x$ for fixed $x > 0$. Put $t_0 = p_c^2/D_c = p_c/2$ and define the rectified Gaussian second moment

$$H(\mu, s) = (\mu^2 + s^2)\Phi(\mu/s) + \mu s\phi(\mu/s).$$

In the strong target-zero states, preactivation divided by k_* tends to $t_0 - b_2 + xZ$. In the target-one states, direct squared-error expansion gives $k_*^2\{(t_0 - p_c)^2 + p_cq_c + x^2\}$. This squared-error expression is conditional on $b_1 = 1$ and receives probability p_c in the full sum. The weak output's clean risk difference from its mono prior is $-p_cq_ck_*^2$; its noise variance is of order k_*^4 . That weak difference receives its importance weight $1/2$. Before cancellation the limit is

$$\begin{aligned} & q_c^2 H(t_0, x) + p_c q_c H(t_0 - 1, x) \\ & + p_c \{(t_0 - p_c)^2 + p_c q_c + x^2\} - p_c q_c / 2 - (p_c + q_c / 2) x^2. \end{aligned}$$

Using $q_c^2 t_0^2 + p_c \{(t_0 - p_c)^2 + p_c q_c\} - p_c q_c / 2 = 0$ gives

$$\frac{\Delta_{\text{frozen}}}{k_*^2} \longrightarrow F(x) = q_c \left[q_c \{H(t_0, x) - t_0^2\} + p_c H(t_0 - 1, x) - \frac{x^2}{2} \right]. \quad (8.60)$$

As $x \downarrow 0$, $F(x) = Bx^2$ plus exponentially small terms, so it is positive for sufficiently small positive x . As $x \rightarrow \infty$, $H(\mu, x) = x^2/2 + 2\mu x/\sqrt{2\pi} + \mu^2/2 + O(x^{-1})$, and therefore

$$F(x) = \frac{2q_c(t_0 - p_c)}{\sqrt{2\pi}} x + O(1) < 0 \quad \text{eventually.}$$

Choose fixed small and large x with these signs. Together with the first crossing at scale $\epsilon^{3/2} \ll k_*$, continuity establishes at least two frozen-bias reconstruction-MSE crossings sufficiently near p_c . This is not the preceding detection-error crossing theorem. No exact root count or unique second crossing is claimed.

8.13 How symmetric oracle bias calibration changes the critical law

Keep the same frozen clean-selected geometry, but independently minimize each output's population corrupted MSE over its bias, for both comparators. For mono's retained feature, write $\beta = \sigma z$. Target-one clipping is negligible near zero noise; the leading coefficient is

$$v(p) = \min_{z \in \mathbb{R}} \{p(1 + z^2) + qH(z, 1)\}. \quad (8.61)$$

Differentiation gives $pz + q\{z\Phi(z) + \phi(z)\} = 0$. The derivative of this equation is $p + q\Phi(z) > 0$, so the solution is unique and negative. Evaluating at $z = 0$ gives $p + q/2$; the derivative there is positive, so $v(p) < p + q/2$.

For sharing at $\sigma = O(\epsilon^{3/2})$, global localization follows from continuity on bounded bias intervals. Positive bias tails are uniformly coercive; the far-negative plateau has loss p , strictly above the attained limiting clean risks. Thus the strong optimum approaches zero and the weak optimum approaches p_c . Here the target-one means stay above $d_0 > 0$. The zero-target function $(u_+)^2$ has a 2-Lipschitz derivative, so zero-mean Gaussian averaging changes its expectation by at most its noise variance. The target-one terms have the same bound plus exponentially small clipping corrections. Uniformly, $|f_\sigma(\beta) - f_0(\beta)| \leq L\sigma^2$. The clean strong-feature objective in this neighborhood is strongly convex with curvature at least $2p$: its target-one states are active quadratics and its target-zero terms are convex. Consequently $|\beta_{\text{cal}} - \beta_{\text{clean}}| \leq L'\sigma = o(k_*)$. This localizes the optimum strictly inside the clean gate branch. There $f''_\sigma \geq p$ for sufficiently small noise, and the gradient at the clean bias is a polynomial factor times $\exp(-dk_*^2/\sigma^2)$. The mean-value theorem gives the exponential bias shift and loss correction. The weak-feature argument uses its nonvanishing gate distances. For mono, the local clean lower bound is $p\beta^2$, while the feasible noisy increase at bias zero is $O(\sigma^2)$. Thus its calibrated bias is $O(\sigma)$; minimizing the uniform limiting objective over bounded $z = \beta/\sigma$ justifies $v(p)$ globally.

This proves the leading sharing coefficient remains D_c .

Consequently the calibrated boundary is

$$\Delta_{\text{cal}} = -C\epsilon^3 + B_{\text{cal}}\sigma^2 + o(\epsilon^3), \quad B_{\text{cal}} = D_c - v(p_c) > B, \quad \sigma_\times = \sqrt{C/B_{\text{cal}}}\epsilon^{3/2}[1 + o(1)]. \quad (8.62)$$

The exponent survives calibration, but the coefficient changes. The frozen-bias second-crossing derivation is not automatically a calibrated second-crossing theorem.

8.14 An arbitrary-load bound for fixed geometries under bias calibration

Let $X_i \geq 0$ have finite third moments, let W be any fixed finite encoder, and let independent isotropic Gaussian noise be added to its code. The tied output is $g_i = (w_i^\top W X + \beta_i + \sigma\|w_i\|Z)_+$. Dependence among features is allowed. Define $r_i(\sigma) = \inf_{\beta_i} \mathbb{E}(g_i - X_i)^2$.

If $w_i \neq 0$, conditional on the complete X , the density of the preactivation is at most $1/(\sigma\|w_i\|\sqrt{2\pi})$. The positive improvement over predicting zero occurs only for $0 < g < 2X_i$. Thus

$$\mathbb{E}[X_i^2] - \mathbb{E}(g_i - X_i)^2 \leq \mathbb{E}(2X_i g_i - g_i^2)_+ \quad (8.63)$$

$$\leq \frac{\mathbb{E} \int_0^{2X_i} (2X_i t - t^2) dt}{\sigma\|w_i\|\sqrt{2\pi}} = \frac{4\mathbb{E}X_i^3}{3\sigma\|w_i\|\sqrt{2\pi}}. \quad (8.64)$$

This is uniform in the bias. Sending the bias to minus infinity approaches the zero prediction, giving

$$\mathbb{E}X_i^2 - \frac{4\mathbb{E}X_i^3}{3\sigma\|w_i\|\sqrt{2\pi}} \leq r_i(\sigma) \leq \mathbb{E}X_i^2. \quad (8.65)$$

For $w_i = 0$, the output is an uncorrupted constant, and the optimal constant is $\mathbb{E}X_i$, giving $r_i = \text{Var}(X_i)$. Therefore, for positive importance weights,

$$\lim_{\sigma \rightarrow \infty} R_{\text{cal}}(W, \sigma) = \sum_i I_i \text{Var}(X_i) + \sum_{i:w_i \neq 0} I_i (\mathbb{E}X_i)^2. \quad (8.66)$$

This limit measures support exposed to noise, not coherence alone. Even a full-support orthogonal geometry obeys it. Freely learned output gains or noise-dependent geometry are outside the theorem.

Compare full-support sharing with a matched mono code retaining a subset S . Put $P = \sum_{i \notin S} I_i (\mathbb{E}X_i)^2$ and $Q = 4 \sum_i I_i \mathbb{E}X_i^3 / \|w_i\| / (3\sqrt{2\pi})$ for the sharing code. The sharing lower bound and mono upper bound imply

$$R_{\text{share,cal}} - R_{\text{mono,cal}} \geq P - Q/\sigma. \quad (8.67)$$

When $P > 0$, $\sigma > Q/P$ is a sufficient mono-favored region. It is a conservative bound, not an exact phase boundary or a clean optimizer for arbitrary load.

Finally, continuity of calibrated risk does not require an attained minimizer. Coupling the same X, Z , the ReLU Lipschitz inequality and the reverse triangle inequality in L^2 give, uniformly in β_i ,

$$|\sqrt{f_i(\beta_i, \sigma)} - \sqrt{f_i(\beta_i, \tau)}| \leq |\sigma - \tau| \|w_i\|.$$

Taking infima in both directions proves the same inequality for $\sqrt{r_i}$. Hence a strict clean sharing advantage and $P > 0$ imply at least one finite calibrated ordering crossing. For the two Bernoulli features the high-noise difference is exactly $p^2/2$ in the limit.

8.15 Verification status and contribution scope

The constants are $K \simeq 1.5786893$, $C \simeq 0.2870182$, $\sqrt{C/B} \simeq 1.9835585$ and $\sqrt{C/B_{\text{cal}}} \simeq 1.3187498$. Four prescribed 70-digit checks used $\epsilon = .02, .01, .005, .0025$. The ratios k_*/ϵ decrease from 1.590019 to 1.580230; the clean gain divided by ϵ^3 decreases from 0.288519 to 0.287327, toward the derived constants. The prescribed lower critical bracket was negative in all four cases. The upper bracket was positive in the last three, but negative at .02. Both prescribed gate-scale samples were negative in all four cases: the positive small- x limit did not overcome finite- ϵ corrections at $x = .05$. These failed finite brackets are retained. They neither verify the full two-crossing map nor refute the asymptotic theorem. No calibrated critical toy experiment has yet been run.

The research candidate is the composition of clean-selected geometry with an operational corrupted-risk boundary, its critical scaling and its calibration dependence. The general support bound extends the declared fixed-geometry risk model to arbitrary load and nonbinary nonnegative features; it does not solve arbitrary-load training selection. Neither the existence of superposition phases nor a generic mono/noisy advantage is claimed as new. Publication novelty of these precise additions remains a targeted prior-art question, not a consequence of this derivation.

8.16 An exact global certificate for the prescribed algebraic frequencies

The critical numerical frequencies are $p = p_c - \epsilon$, where $p_c = (3 - \sqrt{5})/2$ and each prescribed ϵ is rational. They are algebraic probabilities; replacing them by nearby rationals would change the problem. A short global branch argument covers all six exact probabilities, without repeating an exhaustive rational-coefficient root enumeration.

Theorem 8.2 (Global clean geometry on a finite critical strip). *In the two-feature, one-code-dimensional Bernoulli model with importance $(1, 1/2)$, total encoder energy one, tied ReLU reconstruction and free clean biases, let*

$$\frac{9}{25} \leq p < p_c.$$

Up to global encoder sign, the unique globally minimizing geometry is

$$W_* = (\sqrt{a_*}, -\sqrt{c_*}), \quad a_* = (1 + k_*^2)^{-1}, \quad c_* = k_*^2(1 + k_*^2)^{-1}.$$

Here k_ is the unique root in $(0, 1/8)$ of*

$$J_p(k) = A(p) + 3pqk + \frac{3 - 5p - p^2}{2}k^2 - pqk^3 = 0, \quad A(p) = p - \frac{1 - p + p^2}{2}. \quad (8.68)$$

Its globally optimal biases are

$$\beta_1 = \frac{p(c_* + pr_*)}{D}, \quad \beta_2 = p(a_* + r_*), \quad r_* = \frac{k_*}{1 + k_*^2}, \quad D = 1 - p + p^2.$$

At $p = p_c$, the global optimum is important-feature mono retention with $k_ = 0$; no mixed geometry ties it.*

Proof. Let $a \geq 1/2$ be the larger squared column norm, $c = 1 - a$ and $r = \sqrt{ac}$. The parameter $k = \sqrt{c/a}$ lies in $[0, 1]$. The finite bias profiling calculation in Proposition 7.1 and the subsequent sign-sector profiles cover all free biases, including the all-off plateau and nonconvex activation kinks. We use those profiles to exclude every competing geometry, rather than assume that a stationary cubic root is global.

Same-sign geometries and mono orientation. The globally profiled same-sign strong and weak losses are

$$S_+ = \frac{pqc(1 + 2qr)}{2 - p}, \quad W_+ = pqa.$$

Since $S_+ \leq pqc \leq W_+$, assigning importance one to the strong column is the better orientation. Even that orientation has excess over important-feature mono

$$S_+ + \frac{1}{2}W_+ - \frac{1}{2}pq = \frac{pqc(p/2 + 2qr)}{2 - p} > 0$$

for every mixed same-sign code. Retention of the unimportant feature has loss pq , larger than the important-retention loss $pq/2$.

Complete opposite-sign weak profile on the strip. The globally profiled strong loss and bias are

$$S_- = pqc - \frac{pq}{D}(qr - pc)^2, \quad \beta_s = \frac{p(c + pr)}{D}.$$

The possible weak minima are the two-, three- and all-active vertices derived in Equations (7.44)–(7.46). The two-active vertex requires $p \leq k - k^2 \leq 1/4$, so it is impossible on this strip. The all-active vertex is feasible for $k \leq p/q$ and has loss $W_4 = pqa$. The three-active vertex is feasible for $k \geq p/q$ if $p + Dk^2 - qk \geq 0$; its loss is

$$W_3 = pqa - \frac{pq}{D}(qr - pa)^2.$$

Its lower feasibility restriction is automatic here, because

$$pa + Dc - qr = [pa - (r - c)] + p(r - qc) > 0.$$

Indeed, the first bracket is $a\{p - (k - k^2)\} > 0$ and the second term is nonnegative. The first weak interval's stationary point is beyond its right boundary. Target-one activation kinks have downward derivative jumps and cannot be minima; target-zero kink minima are covered by adjacent quadratic vertices on their closures. The all-off weak loss p is worse than the feasible zero-bias loss $pqa^2 + p^2 < p$ for $c > 0$. Thus no additional weak-bias minimum has been omitted.

For both surviving profiles the weak loss is at least the strong loss:

$$W_4 \geq pqc \geq S_-, \quad W_3 - S_- = \frac{pq(a - c)}{D}\{q + 2pqr\} \geq 0.$$

Therefore placing importance one on the stronger column is always the better orientation. This covers importance swaps rather than assuming them away.

Excluding the three-active weak region globally. For $k \geq p/q$, the weighted excess over mono is

$$\Delta_3 = \frac{pq}{2D(1 + k^2)^2} H_p(k),$$

where

$$H_p(k) = (1 - p - p^2)k^4 + 4pqk^3 + (-2 + 5p - 2p^2)k^2 + 2pqk - p^2.$$

Write $U(p) = 1 - p - p^2$ and $V(p) = -2 + 5p - 2p^2$. On this strip,

$$pq \geq \frac{144}{625}, \quad V(p) + p = -2(p^2 - 3p + 1) \geq -\frac{62}{625}, \quad U(p) > 0.$$

Moreover $qk \geq p \geq 9/25 > 1/6$ and $k \leq 1$. Hence

$$\begin{aligned} H'_p(k) &= 4U(p)k^3 + 2pk(6qk - 1) + 2pq + 2\{V(p) + p\}k \\ &\geq \frac{288}{625} - \frac{124}{625} = \frac{164}{625} > 0. \end{aligned}$$

At $k = p/q$, the three-active and all-active weak profiles agree. For the all-active profile the excess is

$$\Delta_A = \frac{pqc}{D}\{A(p) + (1 - 2p)c + 2pqr\}.$$

At this boundary $r = pq/(p^2 + q^2)$. Since

$$A(p) \geq -\frac{31}{1250}, \quad r \geq \frac{144}{337}, \quad 1 - 2p > 0,$$

its bracket is bounded below by

$$-\frac{31}{1250} + 2\frac{144}{625}\frac{144}{337} = \frac{72497}{421250} > 0.$$

Thus $H_p(p/q) > 0$ and strict increase excludes every three-active weak geometry. The global minimum can only be in the all-active weak branch $0 \leq k \leq p/q$.

A unique global minimum in the remaining branch. Substitute a, c, r into the exact remaining loss difference:

$$\Delta_A(k, p) = \frac{pq}{D} \frac{k^2 \{A(p) + 2pqk + E(p)k^2\}}{(1 + k^2)^2}, \quad E(p) = \frac{1 - p - p^2}{2}. \quad (8.69)$$

Direct differentiation gives

$$\partial_k \Delta_A = \frac{2pq}{D} \frac{k J_p(k)}{(1 + k^2)^3}. \quad (8.70)$$

For $p < p_c$, $J_p(0) = A(p) < 0$. Its derivative is

$$J'_p(k) = 3pq(1 - k^2) + (3 - 5p - p^2)k > 0 \quad (0 \leq k \leq 1),$$

because $p \leq p_c < 2/5$ implies $3 - 5p - p^2 > 0$. At $k = 1/8$, dropping the positive quadratic term yields the exact lower bound

$$J_p(1/8) \geq -\frac{31}{1250} + \frac{3}{8} \frac{144}{625} - \frac{1}{512} \frac{1}{4} = \frac{78223}{1280000} > 0.$$

There is therefore exactly one root $k_* \in (0, 1/8)$. Since $p/q \geq 9/16 > 1/8$, it lies strictly within the all-active weak branch. The loss strictly decreases before this root and strictly increases after it. Its minimum is negative, since the loss initially decreases from $\Delta_A(0, p) = 0$. All other sign, importance and weak-bias sectors have been excluded above, so this is the unique global geometry up to global sign. The unique minimizing biases follow from their strict interior quadratic vertices. At p_c , $A = 0$ and $J_p(k) > 0$ for $k > 0$, which gives the stated mono endpoint. $\square$

8.16.1 Exact root brackets without rationalizing the probability

The six prescribed ϵ values .02, .01, .005, .0025, .001, .0005 are rational, and $p_c - .02 > 9/25$; hence every prescribed frequency lies within the theorem's strip. Each coefficient of J_p belongs to the exact quadratic field $\mathbb{Q}(\sqrt{5})$.

Represent a field element as $u + v\sqrt{5}$ with rational u, v , reducing $(\sqrt{5})^2$ to five in every multiplication. At a rational dyadic k , evaluate $J_p(k)$ exactly in this representation. Its sign needs no floating approximation: if u and v have the same sign, that sign is decisive; if $u > 0, v < 0$, compare u^2 with $5v^2$; if $u < 0, v > 0$, reverse that comparison. Zero coefficients are handled directly. Nonzero rational u, v cannot satisfy $u^2 = 5v^2$.

Bisection of $[0, 1/8]$ using these exact signs gives a dyadic bracket for the unique globally selected root at any declared width. Preserve the field coefficients, endpoint signs and bracket width with the numerical geometry. High-precision evaluation inside that bracket gives approximate weights and risks; the probability itself has never been replaced by a rational surrogate. The strip theorem establishes global selection, while the bracket establishes which algebraic root the numerical evaluation approximates.

8.17 Prescribed finite reconstruction-risk maps

Six fixed distances below the clean storage transition were evaluated:

$$\epsilon \in \{.02, .01, .005, .0025, .001, .0005\}.$$

The same opposite-sign branch geometry and clean biases were used in each case. At execution, the numerical solver supplied a local stationary point, with residual below 10^{-55} and valid branch

inequalities. The global geometry justification must come from the separate analytical branch proof; a residual is not a global optimization certificate.

For each geometry, the population weighted reconstruction MSE was evaluated at zero and 121 logarithmic noise levels from 10^{-7} to 2 using 70-decimal arithmetic. Every state of the two-feature Bernoulli population was included. Adjacent sign changes were bisected to relative bracket width below 10^{-5} . All 732 population rows and all 726 adjacent grid intervals are retained. Same-sign endpoints mean no detected bracket, not proof that the interval contains no root.

ϵ	Detected brackets	First $\sigma/\epsilon^{3/2}$	Second $\sigma/\epsilon^{3/2}$
.02	0	—	—
.01	2	2.016567	5.030085
.005	2	1.954383	7.509814
.0025	2	1.960828	10.851459
.001	2	1.974055	17.358848
.0005	2	1.978801	24.640265

The first detected boundary approaches the analytical coefficient $\sqrt{C/B} = 1.98355848\dots$. No exponent was fitted. The farthest prescribed case has no detected crossing on this grid; that result is retained. The earlier gate-scale samples at $x = .05$ and 10 missed the positive interval, while the prescribed noise grid resolves it in five cases. This is added evidence, not replacement of those failed samples. The full positive-noise root count is not certified.

An independent implementation integrated the Gaussian squared-error function directly at three fixed case/grid-index combinations, for both comparators. Six risk comparisons agree to within 2.78×10^{-17} ; the omitted $|Z| > 12$ tail bounds are below 1.38×10^{-32} . The quadrature routine's reported errors are also retained. This is an implementation check of selected population risks, not directed-rounding certification of every root.

8.18 Numerical calibration bounds at critical precision

The following computation preserves the existing scalar bias branch-and-bound method while replacing float64 arithmetic by 70-decimal arithmetic and declaring a numerical slack of 10^{-40} . For one output with offsets o_j , binary targets y_j and probabilities P_j , let $s > 0$ be its noise standard deviation and $\mu_j = o_j + \beta$. Its objective satisfies

$$f''(\beta) = 2 \sum_j P_j \left[\Phi(\mu_j/s) - \frac{y_j}{s} \phi(\mu_j/s) \right].$$

On a bias interval $[l, u]$, define $d_j = \text{dist}(0, [o_j + l, o_j + u])$. Since $0 \leq \Phi \leq 1$, $\sum_j P_j = 1$, and ϕ decreases with absolute argument,

$$|f''(\beta)| \leq H_{[l,u]} = 2 \left[1 + \sum_j P_j y_j \frac{\phi(d_j/s)}{s} \right] \quad (l \leq \beta \leq u). \quad (8.71)$$

This specializes the preceding global curvature bound without changing the risk or the permitted biases. At the midpoint m and radius $r = (u - l)/2$, Taylor's theorem therefore gives the numerical lower bound

$$\max\{0, f(m) - |f'(m)|r - H_{[l,u]}r^2/2 - 10^{-40}\}.$$




Figure 8.4: Saved population MSE curves with frozen clean biases. Positive means mono has lower risk. The horizontal coordinate is scaled by $\epsilon^{3/2}$ and the vertical coordinate by ϵ^3 ; these scalings are derived rather than fitted. Dashed lines show the predicted first-crossing coefficient. The display window clips values outside $[-2, 2]$; the full numerical curves are preserved in the raw CSV. Connecting sampled points does not certify root completeness.

Any evaluated feasible bias supplies an upper bound with the same declared slack. The outside half-lines are bounded using the preceding tail and monotonicity arguments. Thus no local bias optimizer is mistaken for a global result.

The prescribed total weighted-loss gap target is $\min(10^{-10}, C\epsilon^3/100)$. Each output's gap target is this quantity divided by $1 + 1/2$, so the weighted sum obeys the total target. A maximum of 100,000 interval expansions per output is retained. Differences between model losses are reported as intervals, and a sign is claimed only when that interval excludes zero. High-precision arithmetic with declared slack remains a numerical bound check, not exact or directed-rounding interval arithmetic. Calibration results are a separate execution record; the frozen map above does not supply them.

8.19 Finite calibrated boundary comparisons and numerical resolution

The same six globally selected clean geometries were evaluated under symmetric oracle bias calibration. The prescribed noise brackets are 0.5 and 2 times the derived calibrated coefficient $1.3187498488 \dots \epsilon^{3/2}$. All 12 comparisons met their declared global numerical loss-gap targets. At every lower bracket sharing has lower weighted MSE; at every upper bracket mono has lower weighted MSE. The difference was evaluated using the sharing lower bound minus the mono upper bound and the sharing upper bound minus the mono lower bound. These intervals exclude zero for the 12 prescribed endpoints.

Only these successful brackets were bisected. The loss-gap target was left unchanged, and execution stopped when a midpoint's difference interval contained zero. No additional calibration tolerance was chosen to make a narrow boundary plot. The surviving sign-certified numerical

brackets are:

ϵ	Lower $\sigma/\epsilon^{3/2}$	Upper $\sigma/\epsilon^{3/2}$
.02	1.260636	1.260757
.01	1.280920	1.281402
.005	1.296857	1.300720
.0025	1.277539	1.339355
.001	1.308447	1.323901
.0005	1.308447	1.323901

The final two brackets contain the derived limiting coefficient. The wider .0025 bracket reflects numerical resolution under the unchanged stopping rule, not evidence that the crossing location jumps. No calibrated second-crossing or complete noise-map claim follows. All intermediate bias bounds, active queues, source snapshots and input geometry records are retained. These are high-precision numerical global-gap checks with declared slack, not exact interval-arithmetic certificates.




Figure 8.5: Saved calibrated crossing brackets and the analytically derived limiting coefficient. Error bars show the actual sign-certified numerical brackets; plotted midpoints are visual locations, not estimates with greater precision. All six bisections stopped at an ambiguous midpoint under the unchanged numerical loss-gap target.

The completed two-feature result now combines global clean selection, critical reconstruction-risk predictions, prescribed finite frozen maps, and symmetric calibration controls. Arbitrary-load trained toy validation and real-representation transfer remain subsequent tests. Nothing here establishes universal monosemantic robustness or guarantees a venue.

8.20 The prescribed eight-feature clean test and its limitations

The broader clean-training stage used eight independent Bernoulli concepts with $p = .2$, four code dimensions, encoder energy four, and importance $I_i = 2^{-i/4}$ for $i = 0, \dots, 7$. The existing projected Adam settings were preserved: learning rate .01, momentum parameters .9, .999, stabilizer 10^{-8} , and 5,000 steps for each of seeds $0, \dots, 4$. All 256 population states were used.

Mono retained the four most important concepts on orthogonal unit columns and predicted prior means for the others. Its clean weighted MSE is 0.2514085403.

Seed	Final clean MSE	Tangent gradient norm	Q/P
0	.2494283	.102105	461784.60
1	.3509482	.133606	15.49270
2	.2970402	.126647	11.29595
3	.3259854	.225744	11.90185
4	.2982442	.147117	11.33140

Only seed zero’s final iterate beats mono. All five fail the existing loss-stability diagnostic, which compares the mean losses in the last two 100-step windows at tolerance 10^{-5} . Their encoder tangent gradients are materially nonzero. No result is labeled converged or globally optimal. Every trace and all 51 checkpoints per seed are saved; earlier lower losses are not substituted for the prescribed final iterates. Bias-only clean profiling does not materially fix these encoder results.

All columns remain mathematically nonzero, so the support penalty is $P = 0.0628521351$. Seed zero has two column norms near 2.60×10^{-6} and 2.18×10^{-6} , yielding an enormous sufficient bound scale through Q ’s inverse norm. These columns are not retrospectively rounded to zero. The result exposes both optimizer stability and bound-tightness limitations. The two-feature global selection proof remains valid, but does not certify these larger runs. No larger-toy noisy calibration or nonbinary run has yet been performed.

8.21 A bounded repair of the larger-toy optimizer

The eight-feature test uses independent Bernoulli inputs x , the enumerated population probabilities P_x , importance weights I_i , and the objective

$$L(W, \beta) = \sum_x P_x \sum_i I_i [\text{ReLU}(z_i(x)) - x_i]^2, \quad z(x) = W^\top W x + \beta, \quad \|W\|_F^2 = E = 4.$$

The initial fixed-step projected-Adam runs are retained without replacing their final iterates by earlier checkpoints. This subsection describes a single bounded optimization diagnostic applied to all five saved final iterates; it supplies neither a global training theorem nor an extra scientific setting.

8.21.1 Population gradients and the energy constraint

Let X contain population states as rows, $H = XW^\top$, and $Z = HW + \mathbf{1}\beta^\top$. Define the matrix

$$D_{xi} = 2P_x I_i (\text{ReLU}(Z_{xi}) - X_{xi}) \mathbf{1}_{\{Z_{xi} > 0\}}.$$

At nonzero preactivations, differentiation gives

$$dZ = dH W + H dW + \mathbf{1} d\beta^\top, \quad dH = X dW^\top,$$

and therefore

$$\nabla_W L = H^\top D + (DW^\top)^\top X, \quad \nabla_\beta L = D^\top \mathbf{1}.$$

Both encoder appearances matter in the first expression. At exact ReLU kinks the implementation uses the selected derivative zero. This convention does not by itself supply a complete nonsmooth stationarity test.

For the Frobenius sphere, a tangent displacement U satisfies $\langle U, W \rangle_F = 0$. The orthogonal projection of the encoder gradient is

$$T = \nabla_W L - \frac{\langle \nabla_W L, W \rangle_F}{E} W.$$

At a smooth constrained stationary point, the Lagrange multiplier condition $\nabla_W L = \lambda W$ is equivalent to $T = 0$; the free bias also requires $\nabla_\beta L = 0$. Neither condition proves a global minimum.

8.21.2 Retraction and actual-loss backtracking

For a trial step $\alpha > 0$, form

$$V = W - \alpha T, \quad W_{\text{new}} = \sqrt{E} \frac{V}{\|V\|_F}, \quad \beta_{\text{new}} = \beta - \alpha \nabla_\beta L.$$

Since T is tangent, $\|V\|_F^2 = E + \alpha^2 \|T\|_F^2 > 0$. The normalization is well defined, preserves energy exactly in real arithmetic, and expands as $W_{\text{new}} = W - \alpha T + O(\alpha^2)$. Away from activation kinks,

$$\left. \frac{d}{d\alpha} L(W_{\text{new}}, \beta_{\text{new}}) \right|_{\alpha=0} = -\|T\|_F^2 - \|\nabla_\beta L\|_2^2.$$

Thus at a smooth point with a nonzero selected gradient, a sufficiently small step satisfies the Armijo condition

$$L(W_{\text{new}}, \beta_{\text{new}}) \leq L(W, \beta) - c\alpha \left(\|T\|_F^2 + \|\nabla_\beta L\|_2^2 \right), \quad c = 10^{-4}.$$

The algorithm checks the actual population loss, rather than assume a step will decrease it. It starts each step at $\alpha = .01$ and permits at most twenty halvings. A failed search stops that run and remains visible. At a kink, the smooth derivative argument may not apply; bounded failure is an unresolved outcome rather than a proof of optimality.

8.21.3 Why merely decaying projected Adam is inconclusive

The existing solver forms Adam moments from the full Euclidean gradient, updates with a coordinatewise preconditioner, and then rescales to the Euclidean sphere. If a local preconditioned direction is written Mg , normalization removes its radial component. A fixed-point condition of the form $P_W(Mg) = 0$ need not imply the necessary smooth stationarity condition $P_W g = 0$, where P_W is the Euclidean tangent projection. Coordinatewise preconditioning and Euclidean projection do not generally commute. Small learning rates can reduce movement without resolving that distinction. This is a reason for the direct-gradient diagnostic; it is not a proof that preconditioning is the sole cause of every observed failure.

The repair starts from every prescribed final iterate, uses at most 2,000 accepted updates per seed and a total wall-clock budget of 120 seconds, and stops on a selected joint gradient norm at most 10^{-5} , a failed search, a nonfinite value, or a budget limit. Full trial and loss traces, periodic checkpoints and final iterates are retained. The minimum preactivation distance to zero is recorded: a small selected gradient at kinks requires a separate nonsmooth interpretation. No favorable seed, checkpoint or second optimizer is selected.

8.21.4 Recorded repair outcomes and the stopping decision

All five runs reached the fixed limit of 2,000 accepted updates. Every proposal satisfied the actual-loss Armijo test at $\alpha = .01$; there were 10,000 recorded trials, no rejected backtracks and no nonfinite failure. Their final losses are shown against the common mono loss 0.2514085403:

seed	original final loss	repaired final loss	selected joint gradient norm
0	.2494283236	.2362549117	.001033501
1	.3509481697	.2819765217	.03570603
2	.2970402276	.2600155475	.01560052
3	.3259853695	.2505002456	.01755209
4	.2982442175	.2502017777	.01517120

Three final dictionaries now beat mono on clean loss, compared with one before repair, but *none* reaches the predeclared gradient threshold. Only seed zero passes the loss-stability diagnostic. The larger-toy runs therefore remain finite-step optimization outputs, not converged or globally optimal training-selected representations. The actual objective decreases demonstrate that the original unfavorable losses were not evidence of an architectural impossibility; they do not prove general training success.

The sufficient support-bound scales Q/P remain approximately

$$(617.511, 12.5307, 12.3378, 11.4811, 13.0136)$$

for seeds zero through four. No tiny column was thresholded to zero. The single-repair stopping rule now closes this optimizer arm: no second repair, favorable checkpoint selection or huge-noise calibration is launched. Both the original and repaired outputs are retained as limitations alongside the separately proved two-feature critical-boundary result.

8.22 The critical law persists over an interval of feature importances

Fix two independent Bernoulli features with common activation probability p , one code dimension, encoder energy one, tied ReLU reconstruction, and positive importance weights $(1, \eta)$. Clean training minimizes

$$L_0(W, \beta; p, \eta) = \mathbb{E} \left[(\text{ReLU}((W^\top W X)_1 + \beta_1) - X_1)^2 + \eta (\text{ReLU}((W^\top W X)_2 + \beta_2) - X_2)^2 \right].$$

The importance interval and probability rectangle in this section are fixed:

$$\eta \in [12/25, 13/25], \quad p \in [7/20, 21/50].$$

No claim about arbitrary importance ratios, loads or neural architectures is inferred from this rectangle. Set $q = 1 - p$, $D = 1 - p + p^2$ and

$$p_c(\eta) = \frac{1 + \eta - \sqrt{1 + 2\eta - 3\eta^2}}{2\eta} = \frac{2\eta}{1 + \eta + \sqrt{1 + 2\eta - 3\eta^2}}.$$

This is the root of $p = \eta D$ lying in the rectangle. Indeed $p - \eta D < 0$ at $p = .36$ for every allowed η , and $p - \eta D > 0$ at $p = .4$ for every allowed η ; hence $p_c(\eta) \in (.36, .4)$.

Theorem 8.3 (Global clean storage transition on the declared rectangle). *For $p \geq p_c(\eta)$ in the rectangle, the unique global clean geometry retains the important concept alone, up to global*

encoder sign. Its loss is ηpq and its biases are $(0, p)$. For $p < p_c(\eta)$, the unique global geometry, up to global sign, is

$$W_* = (\sqrt{a}, -\sqrt{c}), \quad a = (1 + k_*^2)^{-1}, \quad c = k_*^2(1 + k_*^2)^{-1}, \quad r = k_*(1 + k_*^2)^{-1},$$

where k_* is the unique root in $(0, 1/8)$ of

$$J_{p,\eta}(k) = A + 3pqk + (2 - 3p - \eta D)k^2 - pqk^3 = 0, \quad A = p - \eta D. \quad (8.72)$$

The globally optimal clean biases are

$$\beta_1 = \frac{p(c + pr)}{D}, \quad \beta_2 = p(a + r).$$

Proof. Let $a \geq 1/2$ be the stronger column's squared norm, $c = 1 - a$, $r = \sqrt{ac}$, and $k = \sqrt{c/a} \in [0, 1]$. The earlier finite activation-interval bias proof covers all vertices, interval boundaries and the all-off plateau. Positive importance weights do not change a fixed feature's optimal bias; they weight its minimized loss. We cover both sign sectors and both assignments of the two importances.

Same signs. The globally profiled losses are $S_+ = pqc(1 + 2qr)/(2 - p)$ and $W_+ = pqa$, with $S_+ \leq W_+$. Since $\eta < 1$, assigning importance one to the stronger column is the better assignment. Even that assignment has excess

$$S_+ + \eta W_+ - \eta pq = pqc \left(\frac{1 + 2qr}{2 - p} - \eta \right) > 0,$$

because $1/(2 - p) \geq 20/33 > 13/25 \geq \eta$. The opposite mono orientation has loss $pq > \eta pq$.

Opposite signs: complete weak-bias coverage. The strong feature has exact globally profiled loss

$$S_- = pqc - \frac{pq}{D}(qr - pc)^2.$$

The two-active weak-bias candidate requires $p \leq k - k^2 \leq 1/4$, and is therefore impossible on this rectangle. The all-active weak branch has loss $W_4 = pqa$ for $k \leq p/q$. For $k \geq p/q$, the three-active candidate has loss

$$W_3 = pqa - \frac{pq}{D}(qr - pa)^2.$$

Its lower feasibility condition is automatic, since

$$pa + Dc - qr = a\{p - (k - k^2)\} + p(r - qc) > 0.$$

Here $p > 1/4$ and $r \geq c$. The one-active candidate lies beyond its interval. Target-one kinks have downward derivative jumps and cannot be minima; target-zero kink minima are included by adjacent interval vertices. The all-off loss p is exceeded downward by the feasible zero-bias loss $pqa^2 + p^2 < p$ for a mixed geometry. Thus the two remaining branches exhaust weak-bias global minima. Moreover

$$W_4 \geq pqc \geq S_-, \quad W_3 - S_- = \frac{pq(a - c)}{D}(q + 2pqr) \geq 0.$$

The importance-one strong assignment is therefore no worse than its swap, in both branches.

Excluding the three-active weak branch. Its weighted excess is

$$\Delta_3 = \frac{pq}{D(1+k^2)^2} N_{p,\eta}(k),$$

where direct substitution gives

$$N = (q - \eta D)k^4 + 2pqk^3 + \{p(1 + \eta) - 2\eta D\}k^2 + 2\eta pqk - \eta p^2.$$

On the entire rectangle,

$$pq \geq 91/400, \quad D \leq 309/400, \quad A \geq -517/10000, \quad q - \eta D \geq 1783/10000 > 0.$$

Regroup the derivative as

$$N' = 4(q - \eta D)k^3 + 2pk\{3qk - (1 - \eta)\} + 4Ak + 2\eta pq.$$

For $k \geq p/q$, $3qk \geq 3p \geq 21/20 > 1 - \eta$. The first two terms are nonnegative, and $k \leq 1$ gives

$$N' \geq -4 \frac{517}{10000} + 2 \frac{12}{25} \frac{91}{400} = \frac{29}{2500} > 0.$$

At $k = p/q$, this branch matches the all-active loss difference

$$\Delta_A = \frac{pq c}{D} \{A + (1 - 2p)c + 2pqr\}.$$

At the matching boundary, $r = pq/(1 - 2pq) \geq 91/218$, and $1 - 2p > 0$. Thus its bracket is at least

$$- \frac{517}{10000} + 2 \frac{91}{400} \frac{91}{218} = \frac{9417}{68125} > 0.$$

Strict increase then excludes every three-active weak geometry.

Global minimum in the remaining branch. In the all-active weak branch define $E = q - \eta D$. Algebra gives

$$\Delta_A = \frac{pq}{D} \frac{k^2(A + 2pqk + Ek^2)}{(1 + k^2)^2}, \quad \partial_k \Delta_A = \frac{2pq}{D} \frac{k J_{p,\eta}(k)}{(1 + k^2)^3}. \quad (8.73)$$

Since $2E - A = 2 - 3p - \eta D \geq 3383/10000 > 0$,

$$J'_{p,\eta}(k) = 3pq(1 - k^2) + 2(2 - 3p - \eta D)k > 0 \quad (0 \leq k \leq 1).$$

For $A < 0$, $J(0) < 0$, whereas

$$J(1/8) \geq - \frac{517}{10000} + \frac{3}{8} \frac{91}{400} - \frac{1}{512} \frac{1}{4} = \frac{42399}{1280000} > 0.$$

The unique root lies in $(0, 1/8)$, strictly below $p/q \geq 7/13 > 1/8$. The loss decreases then increases, with a negative minimum below mono. All competing branches have been excluded, and the strong/weak losses are strictly ordered at this root, so the geometry is unique up to global sign. If $A \geq 0$, every term in the bracket of Δ_A is nonnegative and the positive overlap term makes the excess strictly positive at every mixed geometry. This includes equality $A = 0$ and proves unique mono at the transition. The displayed biases are the strictly interior quadratic minima for $p < p_c$. $\square$

8.22.1 Uniform clean critical expansion

Put $\epsilon = p_c(\eta) - p > 0$ and let

$$p_0 = p_c(\eta), \quad q_0 = 1 - p_0, \quad D_0 = 1 - p_0 + p_0^2, \quad s_\eta = \sqrt{1 + 2\eta - 3\eta^2}.$$

The slope at the critical point is $\partial_p(p - \eta D)|_{p_0} = 1 + \eta - 2\eta p_0 = s_\eta$, so exactly

$$A(p_0 - \epsilon) = -s_\eta \epsilon - \eta \epsilon^2.$$

The monotone cubic at $(k, \epsilon) = (0, 0)$ has $\partial_k J = 3p_0 q_0 > 0$. These derivatives stay bounded away from zero over the compact importance interval. The implicit-function expansion, with uniform remainder, yields

$$k_* = K_\eta \epsilon + O(\epsilon^2), \quad K_\eta = \frac{s_\eta}{3p_0 q_0}. \quad (8.74)$$

Substituting k_* into the exact global branch gives

$$\begin{aligned} \Delta_0 &= \frac{p_0 q_0}{D_0} \left[-s_\eta K_\eta^2 + 2p_0 q_0 K_\eta^3 \right] \epsilon^3 + O(\epsilon^4) \\ &= -C_\eta \epsilon^3 + O(\epsilon^4), \quad C_\eta = \frac{s_\eta^3}{27D_0 p_0 q_0} > 0. \end{aligned}$$

The weak column energy is $c_* = K_\eta^2 \epsilon^2 + O(\epsilon^3)$. This is a proved parameter-family law, not a fitted exponent. The finite rectangle supplies global selection, so the implicit stationary solution is not mistaken for an unverified branch minimum.

8.22.2 Gaussian population risks and the frozen-bias law

The corruption is independent code noise $h + \sigma Z$, $Z \sim N(0, 1)$. For output i , the conditional mean and standard deviation are

$$\mu_i(X) = (W^\top W X)_i + \beta_i, \quad s_i = \sigma |w_i|.$$

For $s > 0$, Gaussian integration gives

$$M_1(\mu, s) = \mu \Phi(\mu/s) + s \phi(\mu/s), \quad M_2(\mu, s) = (\mu^2 + s^2) \Phi(\mu/s) + \mu s \phi(\mu/s).$$

Thus the exact population risk is

$$R = \sum_X P_X \sum_i I_i \{M_2(\mu_i, s_i) - 2X_i M_1(\mu_i, s_i) + X_i^2\}, \quad (I_1, I_2) = (1, \eta).$$

At zero standard deviation use the corresponding deterministic ReLU moments. All risk comparisons below use this same weighted MSE.

For clean-selected sharing, the strong feature's offsets are $0, -r, a, a - r$ on states $(0, 0), (0, 1), (1, 0), (1, 1)$. Its bias is $\beta_1 = p(c + pr)/D$. As $\epsilon \rightarrow 0$, the first mean is $t_0 k + O(k^2)$ and the second is $-(1 - t_0)k + O(k^2)$, where $t_0 = p_0^2/D_0 = \eta p_0 \in (0, 1)$. The other two means remain near one. The strong output is therefore clean-active in three states, with total probability D , and their distances from zero are uniformly at least a constant times k . The weak output bias $p(a + r)$ places every mean uniformly away from zero; it is all-active.

At $\sigma = t\epsilon^{3/2}$, for bounded $t \geq 0$, $\sigma/k = O(\sqrt{\epsilon})$. Gaussian gate corrections are bounded by polynomial factors times $\exp(-d/\epsilon)$, uniformly in the importance interval and bounded t . Strong active-state variance contributes $Da\sigma^2$; weak variance contributes $\eta c\sigma^2$. Mono with

clean bias zero has important-output variance coefficient $p + q/2$ (the zero state contributes half a Gaussian second moment), plus the unchanged omitted-feature loss ηpq . It follows that the frozen difference obeys

$$\Delta_{\text{fr}} = -C_\eta \epsilon^3 + B_{\text{fr}}(\eta) \sigma^2 + o(\epsilon^3), \quad B_{\text{fr}} = D_0 - \frac{1 + p_0}{2} = \frac{q_0(1 - 2p_0)}{2} > 0. \quad (8.75)$$

The remainder is uniform for bounded t and for the declared importance interval. The coefficient is not an attacker margin or a universal monosemanticity metric; it is a population variance coefficient for this decoder and corruption.

8.22.3 Global bias localization under calibration

Both models now receive free population bias minimization at each noise level, with their encoders fixed. The following argument ensures that a local bias expansion describes the global infimum.

Positive large biases have uniformly coercive squared loss; far-left biases have loss approaching p , which is separated from the relevant minima by a positive amount on the compact probability interval. Thus global bias minimizers can be localized to a common compact set. ReLU is Lipschitz, and coupling clean and noisy scores gives uniform convergence of their objectives there as $k, \sigma \rightarrow 0$. The limiting strong-feature optimum is bias zero, uniquely; the weak-feature optimum is bias p_0 , uniquely. Consequently global minimizers are in neighborhoods of these biases.

In a small fixed neighborhood of the strong clean bias, target-one states have mean bounded away from zero. Their squared-error losses have curvature $2p$ and their Gaussian gate errors are exponentially small. Target-zero terms are squared ReLUs, whose derivative is globally two-Lipschitz and whose noise-induced objective change is at most the score variance. Therefore $|f_\sigma - f_0| \leq L\sigma^2$ plus exponentially small terms, uniformly in that neighborhood. The clean objective is strongly convex there, with $f_0(\beta) - f_0(\beta_*) \geq p(\beta - \beta_*)^2$. Comparing a global noisy minimizer with the clean bias gives $|\beta_\sigma - \beta_*| = O(\sigma) = o(k)$.

The clean strong bias has a uniform positive distance proportional to k from both nearby target-zero activation boundaries. Localization therefore places the global noisy minimizer in the same active branch. In that branch the clean derivative is $2D(\beta - \beta_*)$, while the noise correction to it is exponentially small because each gate is many Gaussian standard deviations away. The noisy curvature stays bounded away from zero. The optimal bias displacement and its improvement over the frozen clean bias are therefore exponentially small, not order σ^2 .

For the weak output, all conditional means remain bounded away from zero near its globally localized bias p_0 . The linear-quadratic variance is independent of its bias; gate corrections and bias changes are again exponentially small. The positive importance factor η does not change the scalar minimizing bias and is bounded above and below uniformly. Thus calibration of sharing leaves its leading noise coefficient D_0 unchanged. These localization steps exclude a distant competing bias minimum rather than assume local root finding establishes a global optimum.

8.22.4 Mono calibration and the uniform calibrated law

For the retained mono feature, the clean bias loss in a fixed neighborhood of zero is $p\beta^2 + q\text{ReLU}(\beta)^2 \geq p\beta^2$. The same variance bound gives $\beta_\sigma = O(\sigma)$ for its global noisy optimum. Write $\beta = \sigma z$ with bounded z . The target-one gate correction is exponentially small, while the target-zero moment is exact:

$$f_\sigma(\sigma z) = \sigma^2 \{p(1 + z^2) + qM_2(z, 1)\} + \text{exponentially small terms.}$$

Define

$$v(p) = \min_{z \in \mathbb{R}} \{p(1 + z^2) + qM_2(z, 1)\}.$$

Differentiating the bracket yields the root equation

$$pz + q\{z\Phi(z) + \phi(z)\} = 0,$$

whose derivative is $p + q\Phi(z) > 0$. Its unique root is negative, since the left side is positive at zero and negative as $z \rightarrow -\infty$. It stays uniformly bounded on the compact probability interval. At $z = 0$ the bracket is $p + q/2$, and its positive derivative there means a small negative bias strictly improves it, so $v(p) < p + q/2$.

The dropped mono output remains an exact prior-mean predictor with loss ηpq at every noise level. Combining global sharing localization and global mono calibration gives

$$\Delta_{\text{cal}} = -C_\eta \epsilon^3 + B_{\text{cal}}(\eta) \sigma^2 + o(\epsilon^3), \quad B_{\text{cal}} = D_0 - v(p_0) > B_{\text{fr}} > 0, \quad (8.76)$$

uniformly over the same family and bounded $\sigma/\epsilon^{3/2}$. Calibrating both models changes the coefficient rather than eliminating the critical tradeoff. The ratio η affects the clean birth slope and gain; the leading noisy coefficient depends on its selected critical probability through D_0 and the mono bias problem.

8.22.5 What boundary is proved

For either policy $j \in \{\text{fr}, \text{cal}\}$, define $\kappa_j(\eta) = \sqrt{C_\eta/B_j(\eta)}$. The coefficients are positive and continuous on the closed importance interval. Uniform expansions give opposite risk signs at any fixed multipliers below and above one times $\kappa_j \epsilon^{3/2}$ for sufficiently small ϵ . Risk is continuous in noise: coupling X, Z and using ReLU Lipschitzness gives continuity for fixed biases, and the localized compact minimization gives continuity of the calibrated minimum. The intermediate value theorem therefore provides at least one crossing with

$$\sigma_{*,j}/\epsilon^{3/2} \longrightarrow \kappa_j(\eta).$$

The limit and sign statement are uniform over the declared importance interval. No unique global crossing count is asserted. A finite validation uses predetermined noise brackets and must preserve failed brackets; the asymptotic theorem does not promise every finite bracket succeeds.

The coefficient distinction predicts an earlier mono-favored critical crossing under symmetric calibration, since $\kappa_{\text{cal}} < \kappa_{\text{fr}}$. This prediction is proved inside the controlled model and importance rectangle, not transferred automatically to arbitrary-load training or vision networks.

8.23 Fixed finite validation of the importance-family law

The question was fixed before execution: does the proved critical storage and corrupted-MSE law survive changes in relative feature importance, under the identical resource constraint and decoder classes? Only four additional geometries were permitted: $\eta = .48, .52$ at $\epsilon = .005, .001$. The original $\eta = .5$ records remain unchanged. Every new probability is $p = p_c(\eta) - \epsilon$ in the proved rectangle. No fitted exponent, extra geometry or stochastic training optimizer enters this validation.

The globally selected cubic root was bracketed with 200 bisections in the exact quadratic field whose radicand is $1 + 2\eta - 3\eta^2$. Rational-pair arithmetic preserved the exact probability, all coefficients, and opposite endpoint signs. The Gaussian risk was evaluated at 70 decimal digits.

Frozen-bias brackets began at one half and twice the derived coefficient times $\epsilon^{3/2}$; only successful brackets were bisected, stopping at relative width 10^{-5} or 100 iterations. All four brackets succeeded. The following values summarize the preserved records:

η	ϵ	k_*/ϵ	K_η	$-\Delta_0/\epsilon^3$	C_η
.48	.005	1.616475	1.613807	.296776	.296493
.48	.001	1.614353	1.613807	.296562	.296493
.52	.005	1.548666	1.545309	.278139	.277290
.52	.001	1.545991	1.545309	.277470	.277290

These are finite numerical comparisons to derived constants; convergence of a theorem is not inferred by fitting four points. The frozen crossing brackets, after division by $\epsilon^{3/2}$, are

η	ϵ	lower	upper	$\kappa_{\text{fr}}(\eta)$
.48	.005	1.862749	1.862760	1.888010
.48	.001	1.879613	1.879623	1.888010
.52	.005	2.063543	2.063555	2.097025
.52	.001	2.086005	2.086017	2.097025

The finite boundaries approach the analytic values from below in these prescribed cases. This does not prove monotone convergence or a complete positive-noise root count.

For calibration, both models received independent population bias minimization at half and twice their analytically derived calibrated critical coefficient. The predeclared total weighted model-gap target was

$$\min\{10^{-10}, C_\eta \epsilon^3 / 100\}.$$

Each output's target was this value divided by $1 + \eta$, using the existing interval-specific curvature bound, 70-digit arithmetic, numerical slack 10^{-40} and at most 100,000 interval expansions. All eight comparisons met their targets. Sharing won at every lower endpoint, and mono won at every upper endpoint. No additional calibrated bisection was performed.

The reused independent 90-digit audit checked 45 archived hashes, all 32 scalar bias ledgers, 2,684 interval nodes, tree partitions, active and pruned leaves, exterior bounds, feasible incumbents, weighted gaps and difference signs. Its maximum recomputation discrepancy was below 1.42×10^{-63} . These are high-precision numerical-gap checks with declared slack, not directed-rounding interval certificates. The full global mathematical selection proof and its independent review are separate from this numerical agreement.

The records and reproduction scripts are in the joint-phase experiment directory. Source snapshots precede the calculations; fresh output directories are required. Successful validation supports the mechanism and implementation within this model. It does not establish publication novelty or replace the still-needed vision transfer.

8.24 A singular endpoint separates frozen and calibrated robustness

This result addresses a specific limit of the critical noise boundary: the frozen-bias coefficient vanishes at importance $\eta = 2/3$. It does not replace the corruption model or claim a universal robustness law. Let X_1, X_2 be independent Bernoulli(p) variables, let $W \in \mathbb{R}^{1 \times 2}$ satisfy $\|W\|^2 = 1$, and reconstruct with $\text{ReLU}(W^\top W X + \beta)$. Loss is population squared error with weights $(1, 2/3)$ and free clean-training biases. At test time add independent σZ , $Z \sim N(0, 1)$, to the scalar code. In the frozen policy keep both models' clean biases. In the calibrated policy minimize each model's population risk over both biases, keeping each encoder fixed. Define $\Delta = R_{\text{sharing}} - R_{\text{mono}}$.




Figure 8.6: Derived importance-dependent critical coefficients in the proved two-feature family. The horizontal axis is relative importance η ; the vertical axis is crossing noise divided by $\epsilon^{3/2}$. Curves are the analytic frozen and calibrated limits. The shaded region predicts different policy orderings asymptotically. Symbols and error bars are all saved finite frozen brackets at the two declared distances, including the unchanged $\eta = .5$ control. This is not an empirical calibrated crossing map or a proof about arbitrary neural networks.

8.24.1 Global clean selection, including competing sign sectors

Put $q = 1 - p$, $D = 1 - p + p^2$, $\epsilon = 1/2 - p$, and

$$a = (1 + k^2)^{-1}, \quad c = k^2(1 + k^2)^{-1}, \quad r = k(1 + k^2)^{-1}, \quad 0 \leq k \leq 1.$$

The earlier complete activation-interval bias profiling applies unchanged: positive importance weights do not alter a scalar bias minimizer. For same signs the strong and weak profiled losses are

$$S_+ = \frac{pqc(1 + 2qr)}{2 - p}, \quad W_+ = pqa.$$

For opposite signs the strong loss is

$$S_- = pqc - \frac{pq}{D}(qr - pc)^2.$$

For $p > 1/4$, the two-active weak vertex is infeasible because it requires $p \leq k - k^2 \leq 1/4$. The one-active vertex is beyond its interval; the all-off plateau is worse than a feasible zero-bias predictor. Target-one kinks have downward derivative jumps and cannot minimize; target-zero kinks are included in adjacent interval closures. The remaining weak minima are

$$W_4 = pqa \quad (k \leq p/q), \quad W_3 = pqa - \frac{pq}{D}(qr - pa)^2 \quad (k \geq p/q).$$

The lower feasibility condition for W_3 is automatic since $pa + Dc - qr = a\{p - (k - k^2)\} + p(r - qc) > 0$. Both sign sectors have strong loss no larger than weak loss: $S_+ \leq pqc \leq W_+$, $S_- \leq pqc \leq W_4$, and $W_3 - S_- = pq(a - c)(q + 2pqr)/D \geq 0$. Thus assigning the larger importance to the stronger column is optimal; this covers the swapped-importance sector rather than discarding it. The reversed mono orientation costs pq , versus $2pq/3$ for retaining the important feature.

At $p = 1/2$ the all-active opposite branch covers every $k \leq 1$. Both mixed sign sectors have excess $cr/6 > 0$ over important mono. Consequently mono is the unique clean geometry up to global encoder sign. Global biases are $(0, 1/2)$. Continuity and compactness also imply that any globally minimizing geometries converge to this mono geometry as $p \uparrow 1/2$; a remote geometry cannot supply a hidden critical branch.

For the finite checks we can strengthen localization to the entire strip $99/200 \leq p < 1/2$, or $0 < \epsilon \leq 1/200$. In the all-active opposite branch the exact excess is

$$F_-(k) = \frac{pq}{D} \frac{k^2(A + 2pqr + Ek^2)}{(1 + k^2)^2}, \quad A = -\epsilon - \frac{2}{3}\epsilon^2, \quad E = \epsilon - \frac{2}{3}\epsilon^2. \quad (8.77)$$

Its derivative has the sign of

$$J(k) = A + 3pqr + (2E - A)k^2 - pqk^3.$$

Here $J' = 3pq(1 - k^2) + 2(3\epsilon - 2\epsilon^2/3)k > 0$ on $[0, 1]$, $J(0) < 0$, and $J(1/8) > -.0051 + 3(.2499)/8 - .25/512 > 0$. There is a unique minimum $k_* \in (0, 1/8)$, inside $k < p/q$.

The three-active opposite region is excluded globally. Its excess is $pqN/[D(1 + k^2)^2]$, with

$$N = Ek^4 + 2pqrk^3 + \{5p/3 - 4D/3\}k^2 + 4pqr/3 - 2p^2/3.$$

For $k \geq p/q$, rewrite its derivative as

$$N' = 4Ek^3 + 2pk(3qr - 1/3) + 4Ak + 4pqr/3 > 0,$$

using $E \geq 0$, $3qr \geq 3p > 1/3$, $A \geq -.0051$ and $pq > .2499$. At the matching point $k = p/q$,

$$A + (1 - 2p)c + 2pqr > -.0051 + 2(.2499)(.49) > 0.$$

Thus this entire competing activation region has positive excess.

Same-sign excess is exactly

$$F_+(k) = \frac{2pq\epsilon}{3/2 + \epsilon}(qr - \epsilon/3). \quad (8.78)$$

If it is negative, $r < \epsilon/(3q) \leq 2\epsilon/3$. Since $r \geq k/2$, initially $k < 4\epsilon/3 \leq 1/150$; substitution in $k = r(1 + k^2)$ sharpens this to $k < (167/250)\epsilon$. Hence, using $pq \leq 1/4$,

$$F_+(k) \geq -\frac{\epsilon}{9}k^2 > -.05\epsilon^3.$$

Meanwhile the feasible opposite geometry $k = \epsilon$ has

$$F_-(\epsilon) = \frac{pq}{D} \frac{\epsilon^3[-1 + 2pq + \epsilon(E - 2/3)]}{(1 + \epsilon^2)^2} < -.16\epsilon^3,$$

since the square bracket is at most $-1/2$ and its positive prefactor exceeds .32. The opposite minimum therefore beats every same-sign geometry throughout the declared strip. Mono has zero excess. This proves global selection, not merely a favorable stationary root. Its biases are

$$\beta_1 = p(c + pr)/D, \quad \beta_2 = p(a + r).$$

8.24.2 The two clean critical branches

The implicit-function theorem applied to J gives

$$k_* = \frac{4}{3}\epsilon + O(\epsilon^2), \quad F_-(k_*) = -\frac{16}{81}\epsilon^3 + O(\epsilon^4).$$

For comparison, substituting $k = u\epsilon$ into the same-sign branch gives

$$F_+(u\epsilon) = \epsilon^3 \left(-\frac{u^2}{9} + \frac{u^3}{6} \right) + O(\epsilon^4).$$

Its local minimum has $k = 4\epsilon/9 + O(\epsilon^2)$ and gain $16\epsilon^3/2187 + O(\epsilon^4)$: the opposite gain is 27 times larger. This branch comparison alone is not a global proof; the preceding sign, orientation and activation-region exclusions supply that proof.

8.24.3 Frozen risk at the singular endpoint

Fix finite $t \geq 0$ and set $\sigma = t\epsilon^{3/2}$. The selected strong feature's four clean means have signs positive, negative, positive, positive, and the nearest distances to zero are constant multiples of k_* . Its clean active probability is D . The weak feature's means remain bounded away from zero and positive. Since $\sigma/k_* = O(\sqrt{\epsilon})$, Gaussian gate corrections are smaller than every algebraic power of ϵ for bounded t . Using Gaussian ReLU first and second moments, the sharing variance coefficient is $Da + (2/3)c$. Frozen mono's coefficient is $(1 + p)/2$; its omitted output has constant loss $2pq/3$. At the endpoint $D_0 = 3/4 = (1 + p_0)/2$, so

$$\frac{\Delta_{\text{fr}}(t\epsilon^{3/2})}{\epsilon^3} \rightarrow -\frac{16}{81}. \quad (8.79)$$

The difference of variance coefficients is $O(\epsilon)$, and multiplying it by σ^2 contributes only $O(\epsilon^4)$ on this scale. This says sharing wins at every fixed finite scaled noise for sufficiently small ϵ . It proves neither all-noise ordering nor a different frozen crossing exponent; a larger noise scale changes the gate analysis.

8.24.4 Symmetric global bias calibration

For completeness, the leading calibration argument requires global rather than local minimization. For each scalar output the objective is coercive on the right, while on the far left it tends to p , separated from its clean minimum. Global minimizers therefore lie in a common compact set. ReLU Lipschitzness and Gaussian integrability give uniform convergence there as $(k, \sigma) \rightarrow 0$, so global strong and weak biases approach their unique limiting values zero and $1/2$.

Near the strong optimum, target-one scores remain uniformly positive and give clean curvature at least $2p$. Target-zero squared ReLU has two-Lipschitz derivative; independent centered noise changes its expected loss by at most its variance. Target-one gate errors are exponentially small. Comparison with the clean bias and strong convexity consequently localizes the optimal displacement to $O(\sigma) = o(k_*)$. The clean strong bias lies a distance proportional to k_* from both nearby target-zero gates. The global noisy optimum is thus in the same activation region, where the clean derivative is $2D(\beta - \beta_*)$ and its Gaussian correction is exponentially small. Its bias improvement does not change the leading $D_0\sigma^2$ coefficient. The weak means are bounded away from zero; their quadratic variance is bias independent, and calibration changes their risk only exponentially at this scale.

For mono, localization and the inequality $f_0(\beta) - f_0(0) \geq p\beta^2$ give $\beta = O(\sigma)$. Writing $\beta = \sigma z$ yields the global coefficient

$$v(p) = \min_z \{p(1 + z^2) + q[(z^2 + 1)\Phi(z) + z\phi(z)]\}.$$

Its derivative vanishes exactly when $pz + q[z\Phi(z) + \phi(z)] = 0$, a strictly increasing function with derivative $p + q\Phi(z) > 0$. Thus its root is unique and negative. At $p = 1/2$ it satisfies $z + z\Phi(z) + \phi(z) = 0$. The derivative at zero is positive, so a negative bias strictly improves on $v(1/2) < 3/4$. The omitted output stays at its prior mean. Consequently

$$\frac{\Delta_{\text{cal}}(t\epsilon^{3/2})}{\epsilon^3} \longrightarrow -\frac{16}{81} + [3/4 - v(1/2)]t^2. \quad (8.80)$$

Numerically $v(1/2) \simeq .6956312935673$, so the positive coefficient is $.05436870643268$ and

$$t_c = \sqrt{(16/81)/(3/4 - v(1/2))} \simeq 1.9060881518.$$

The positivity and formulas are analytic; these decimals are numerical evaluations, not exact constants. For fixed $t > t_c$, sufficiently near the clean endpoint, the identical encoder and noise favor sharing with frozen biases but favor mono with symmetric population bias calibration. This establishes evaluation-policy dependence of a representation-risk ordering under a fully specified task. It does not establish a universal failure of monosemanticity or a result about adversarial attacks.

8.24.5 Prospective finite validation and stopping scope

The numerical protocol fixes $\epsilon = .005, .001$ and $t = 2t_c$ before comparison. Both probabilities lie within the proved global strip. Use exact rational cubic brackets with 200 bisections and the existing 70-digit global-gap bias profiler; preserve all ledgers, unresolved comparisons and source hashes. Stop after those two cases. Finite population calculations are numerical evidence under the declared slack, not directed-rounding interval proofs. No novelty or conference-acceptance claim follows from the analytic identities alone.

8.25 The two prescribed endpoint-policy comparisons

The paper-critical gap was the singularity of the frozen variance coefficient, not the search for another favorable toy. The endpoint proof predicts opposite policy orderings at identical geometry and noise. Before execution, the protocol fixed exactly two distances $\epsilon = .005, .001$ and the same scaled-noise multiplier $t = 2t_c$. Both probabilities lie in the proved global-selection strip. No additional importance ratio, noise point, tolerance adjustment or optimizer was used.

For each case, 200 rational bisections bracketed the monotone clean cubic with exact endpoint signs. The clean-selected encoder and biases were then held fixed for the frozen test. Both encoders remained fixed for the calibrated test, with both models receiving independent population bias minimization. The predeclared 70-digit profiler used numerical slack 10^{-40} , at most 100,000 expansions per output, and total weighted gap

$$\min\{10^{-10}, (16/81)\epsilon^3/100\}.$$

Every retained calibration ledger met its assigned gap target. The observed raw differences are

ϵ	σ	Δ_{fr}	$\Delta_{\text{cal lower}}$	$\Delta_{\text{cal upper}}$
.005	.001347807858	$-2.760446858 \times 10^{-8}$	$7.342184482 \times 10^{-8}$	$7.352032194 \times 10^{-8}$
.001	.000120551600	$-1.906926819 \times 10^{-10}$	$6.020884857 \times 10^{-10}$	$6.040952719 \times 10^{-10}$

For both prescribed cases the frozen difference is negative, favoring sharing, while the calibrated difference is positive, favoring mono. These are comparisons at the same encoder, feature distribution and noise within each row. Calibration can lower both models' absolute risks while changing their relative ordering; a positive calibrated difference does not mean that calibrating sharing made its risk worse.

An independent 90-digit implementation audited 18 archived hashes, all eight bias ledgers and 666 interval nodes, including global-domain partition coverage, exterior bounds, feasible incumbents, weighted model gaps and the difference signs. Frozen risks were separately recomputed using the independent Gaussian moment formulas. The maximum discrepancy was below 5.37×10^{-64} . These numerical bounds use declared slack and are not directed-rounding interval certificates. Independent mathematical review separately checked the full global selection and calibration proof.

This strengthens the interpretation of Project 1: representation geometry and the permitted decoder adaptation jointly determine a robustness ordering. The result neither disputes all monosemanticity benefits nor establishes an adversarial guarantee. It remains a controlled two-feature mechanism. Real learned-representation evidence and a precise novelty comparison are separate paper requirements; no acceptance or spotlight claim follows from these checks.

8.26 Continuous-target profiling for the learned-vision pilot

This section supplies the implementation adaptation required to evaluate the same reconstruction mechanism on saved learned activations. It adds no new architecture or empirical result. Fix N observations $x^{(j)} \in [0, \infty)^2$, finite and normalized using training-only RMS factors, an encoder $w \in \mathbb{R}^2$ with $\|w\|^2 = 1$, and positive feature importances $(1, 2/3)$. For output i let

$$o_j = w_i(w^\top x^{(j)}), \quad y_j = x_i^{(j)} \geq 0, \quad F_i(\beta, s) = \frac{1}{N} \sum_{j=1}^N \mathbb{E}_Z[(\text{ReLU}(o_j + \beta + sZ) - y_j)^2].$$

The output standard deviation induced by scalar code noise is $s = \sigma|w_i|$. Each feature's positive importance multiplies its objective and does not change its minimizing bias. All results below are conditional on this finite activation dataset. Empirical optimality is distinct from optimality over the unknown image population.




Figure 8.7: All two predeclared endpoint comparisons. The horizontal axis is distance $\epsilon = 1/2 - p$ below the storage transition; the vertical axis is the sharing-minus-mono MSE difference divided by ϵ^3 . Negative values favor sharing and positive values favor mono. Within each case geometry and noise are identical across policies. Calibrated error bars are saved numerical loss-gap brackets, not sampling confidence intervals. The dashed limits are $-16/81$ and $16/27$ at $t = 2t_c$. Finite values do not demonstrate monotone convergence to those limits.

8.26.1 Complete clean bias profiling

At $s = 0$, sort the thresholds $t_j = -o_j$ and group equal thresholds. For an open interval between consecutive distinct thresholds, the active set $A = \{j : o_j + \beta > 0\}$ is constant. The objective there is

$$F_i(\beta, 0) = \frac{1}{N} \left[\sum_{j \in A} (o_j + \beta - y_j)^2 + \sum_{j \notin A} y_j^2 \right]. \quad (8.81)$$

With $n_A = |A| > 0$, its unconstrained quadratic vertex is

$$\beta_A = \frac{\sum_{j \in A} (y_j - o_j)}{n_A}.$$

On the interval closure, clip this vertex to its lower and upper endpoints. In the last interval there is no upper clipping. Enumerate every distinct threshold as a boundary candidate as well. At any boundary the newly activating outputs equal zero, so both neighboring quadratic expressions give the same objective value. Derivative continuity is not assumed: an activating observation with target $y_j > 0$ causes a downward derivative jump of $2y_j/N$. Retaining the boundary candidates covers these nonconvex kinks without treating the entire bias objective as convex.

Below the smallest threshold all outputs are off, with loss $N^{-1} \sum_j y_j^2$ independent of bias. Retain this plateau explicitly, using any finite representative below that threshold. The union of the plateau, all interval closures and the last half-line is the whole real line. Each nonempty active interval is a coercive quadratic on its closure; its clipped vertex attains its minimum. Comparing the finite candidate list therefore yields the global empirical clean bias minimum in exact arithmetic. Grouping equal thresholds avoids omitting or double-counting an activation change; evaluating every candidate with the direct ReLU loss also avoids selecting solely from cancellation-prone prefix sums.

The code uses prefix sums of $o_j - y_j$, $(o_j - y_j)^2$ and y_j^2 to construct these candidates. Floating-point evaluation is a numerical implementation of the exact finite enumeration, not a directed-rounding certificate. This scalar completeness does not certify the encoder angle: the prescribed 256/512 grids and one best-cell refinement at each resolution give a bounded numerical resolution check, not a theorem of global empirical encoder selection.

8.26.2 Gaussian conditional moments for continuous targets

For $s > 0$ and $\mu = o_j + \beta$, define $z = \mu/s$. Gaussian integration gives

$$\begin{aligned} M_1(\mu, s) &= \mathbb{E}[\text{ReLU}(\mu + sZ)] = \mu\Phi(z) + s\phi(z), \\ M_2(\mu, s) &= \mathbb{E}[\text{ReLU}(\mu + sZ)^2] = (\mu^2 + s^2)\Phi(z) + \mu s\phi(z). \end{aligned}$$

Indeed, integration on $Z > -\mu/s$ uses $\mathbb{E}[Z\mathbf{1}_{Z>-z}] = \phi(z)$ and $\mathbb{E}[Z^2\mathbf{1}_{Z>-z}] = \Phi(z) - z\phi(z)$. Hence

$$F_i(\beta, s) = \frac{1}{N} \sum_j \{M_2(o_j + \beta, s) - 2y_j M_1(o_j + \beta, s) + y_j^2\}. \quad (8.82)$$

The formula holds for continuous nonnegative y_j ; replacing targets with Bernoulli values is unnecessary and would change the task. At $s = 0$ use the deterministic clean profile above. Differentiating the Gaussian integrals gives

$$\begin{aligned} F'_i(\beta, s) &= \frac{2}{N} \sum_j \{(o_j + \beta - y_j)\Phi((o_j + \beta)/s) + s\phi((o_j + \beta)/s)\}, \\ F''_i(\beta, s) &= \frac{2}{N} \sum_j \left\{ \Phi((o_j + \beta)/s) - \frac{y_j}{s}\phi((o_j + \beta)/s) \right\}. \end{aligned}$$

The second derivative can be negative. A local stationary bias or a bounded numerical optimizer therefore supplies an incumbent only; global calibration requires objective bounds over all remaining biases.

8.26.3 Finite interval and exterior bounds

Let $y_{\max} = \max_j y_j$. For

$$\beta \geq u = y_{\max} - \min_j o_j,$$

every conditional mean satisfies $o_j + \beta \geq y_{\max} \geq y_j$. Each term in F'_i is nonnegative; thus the entire right half-line has lower bound $F_i(u, s)$. This uses the actual continuous target maximum, not the binary constant one.

Choose the finite left endpoint $\ell = -\max_j o_j - 12s - 1$. For every $\beta \leq \ell$, monotonicity of M_1 in its mean and $y_j \geq 0$ imply

$$\begin{aligned} F_i(\beta, s) &\geq \frac{1}{N} \sum_j \{y_j^2 - 2y_j M_1(o_j + \beta, s)\} \\ &\geq \frac{1}{N} \sum_j \{y_j^2 - 2y_j M_1(o_j + \ell, s)\}. \end{aligned}$$

The first step drops nonnegative second moments. This is a valid left-exterior lower bound, not an assertion that the Gaussian tail is exactly zero. Together with the right bound, it reduces certification of a global empirical bias gap to the compact interval $[\ell, u]$.

For any subinterval $[a, b]$ define

$$d_j = \text{dist}(0, [o_j + a, o_j + b]), \quad H_{a,b} = 2 \left[1 + \frac{1}{N} \sum_j \frac{y_j}{s} \phi(d_j/s) \right].$$

The Gaussian density is even and decreases with absolute argument, so $\phi((o_j + \beta)/s) \leq \phi(d_j/s)$ for $\beta \in [a, b]$. The displayed second derivative and $0 \leq \Phi \leq 1$ then give $|F''_i(\beta, s)| \leq H_{a,b}$. At midpoint $m = (a + b)/2$ and radius $r = (b - a)/2$, Taylor's theorem consequently supplies

$$F_i(\beta, s) \geq F_i(m, s) - |F'_i(m, s)|r - \frac{1}{2}H_{a,b}r^2 \quad (\beta \in [a, b]). \quad (8.83)$$

The maximum of this bound and zero is also valid, since squared-error risk is nonnegative. Midpoint evaluation or any other feasible bias provides an upper bound on the global minimum. Split the lowest-bound interval, retain both children, and discard an interval only when its lower bound is no smaller than the current incumbent upper bound. The minimum over retained interval bounds and both exterior bounds is a global lower bound in exact arithmetic. The incumbent minus that minimum is an explicit empirical optimality gap. The interval tree and active heap must remain archived so this logic can be checked.

8.26.4 Numerical tolerances, matched policies and scope

The implemented branch-and-bound subtracts numerical slack 10^{-12} from lower bounds and adds it to feasible upper bounds. Gaussian tail moments occasionally have tiny negative roundoff values; values within the declared slack are clipped to zero, whereas more negative values raise an error. This preserves physical nonnegativity but is not directed rounding, and the slack is not a formal machine-arithmetic error proof. Accordingly, report a floating-point numerical gap check, not an interval-arithmetic theorem.

Feature gaps are targeted at $10^{-7}/(1 + 2/3)$; their importance-weighted sum is at most 10^{-7} per model when every output resolves. Stop after 20000 expansions per output or the shared 300-second budget, preserving unresolved gaps. Do not treat these empirical calibration bounds as uncertainty intervals for the held-out image population.

Both primary frozen policies use biases minimized on the calibration split at zero noise. Noisy policies reoptimize biases on that same split with encoders fixed. Thus their fitted models coincide at zero noise; ordinary training/calibration mismatch cannot itself create the baseline policy contrast. The test split evaluates conditional Gaussian losses without fitting, and paired image resampling quantifies only finite test-sample variability conditional on the fitted models.

Continuous activations do not inherit the Bernoulli clean cubic, critical probability, global storage transition or $\epsilon^{3/2}$ law. Their support, coactivation, amplitudes and selected gate distances differ. The pilot tests only whether decoder calibration changes a learned-feature reconstruction-risk comparison under the specified corruption. A resolved result would be evidence of that mechanism outside Bernoulli inputs; it would not make two arbitrary channels semantic concepts or establish universal monosemanticity robustness.

8.27 A fixed learned-vision check of decoder-policy dependence

The continuous-target pilot was completed using genuine frozen ImageNet-pretrained ResNet18 weights (IMAGENET1K_V1) and CIFAR-10 images. This is an activation reconstruction task, not image generation, image-input corruption, adversarial evaluation or proof that the chosen channels are semantic concepts. All choices below were fixed before feature-risk outcomes; no pair or noise-grid replacement was made.

8.27.1 Inputs, selected geometry and evaluation settings

NumPy seed 20261006 shuffled the 50000 CIFAR-10 training indices. The first 256 supplied training targets, the next 256 calibration targets, and the next 512 held-out test targets. The official checkpoint transforms resize to 256, center crop to 224 and apply ImageNet normalization. Inference used the post-ReLU layer4 output at the predetermined central spatial cell, before average pooling. The first two channels by index with nonzero training mean and variance were channels 0 and 1. Class labels were not used to fit or select the reconstruction targets.

The training RMS factors were 2.1767477055 and 2.4060630488; division without centering preserved nonnegativity and exact zeros. Training zero fractions were .48828125 and .30859375, with coactivation .37109375. The encoder had one dimension and energy one, tied ReLU decoding, free biases and importance weights (1, 2/3). The two predeclared angle resolutions, 256 and 512, followed by one bounded best-cell refinement each, selected the same numerical solution:

$$w_{\text{sharing}} = (.9674172554, .2531873890), \quad L_{\text{sharing,train}} = .3670650955.$$

The train-selected mono comparator retained channel 0, with $L_{\text{mono,train}} = .3849481174$. The empirical clean gain .0178830219 passed the prescribed gate; the resolution-loss difference was zero. This is numerical clean selection, not a globally certified encoder optimum. Both exact-zero support and a distinctly mixed numerical solution were present, so the predefined noisy evaluation proceeded.

Both encoders were frozen. Their primary frozen biases were fitted on the calibration split at zero noise; noisy calibration reoptimized biases on that same split for both models. The primary policies therefore coincided at zero noise. The common code-noise reference was

$$s_{\text{ref}} = \sqrt{\frac{1}{2}\{\text{Var}_{\text{train}}(x_1) + \text{Var}_{\text{train}}(x_2)\}} = .8082696314.$$

The five prescribed noise multipliers were 0, .05, .1, .2, .4, applied to this reference, with identical absolute scalar Gaussian code noise for both encoders. Held-out risks were deterministic conditional Gaussian expectations averaged over test images. Calibration targeted total weighted empirical objective gap 10^{-7} per model, floating slack 10^{-12} , at most 20000 expansions per output and a shared 300-second budget. Every saved calibration profile resolved within its target.

8.27.2 All prescribed held-out comparisons

Define $\Delta_j = R_{\text{sharing},j} - R_{\text{mono},j}$ for policy $j \in \{\text{fr}, \text{cal}\}$, and $T = \Delta_{\text{cal}} - \Delta_{\text{fr}}$. Positive Δ favors mono; positive T means calibration shifts the relative comparison toward mono. The table reports means and 95% paired image-bootstrap percentile intervals using 2000 resamples and fixed seed 20261006. Intervals are pointwise, conditional on the fitted models, and do not include training/calibration uncertainty or give simultaneous coverage over all noise levels. Entries are rounded only for display; full precision and per-image losses remain in the archived outputs.

σ	Δ_{fr}	Δ_{cal}	T
.00000	.009029	.009029	.000000
	[−.015775, .037170]	[−.015775, .037170]	[.000000, .000000]
.04041	.008768	.008866	.000098
	[−.016023, .036947]	[−.015941, .037035]	[.000054, .000140]
.08083	.008118	.008518	.000399
	[−.016766, .036377]	[−.016426, .036760]	[.000208, .000577]
.16165	.006141	.007739	.001598
	[−.018991, .034632]	[−.017634, .036274]	[.000843, .002319]
.32331	.001454	.007401	.005947
	[−.023961, .030110]	[−.019016, .037169]	[.003554, .008243]

The learned sharing model’s training advantage did *not* establish a held-out clean advantage: the primary zero-noise sharing risk was .3792952475 versus mono .3702662291. At every prescribed level both policy difference means were positive, but every individual difference interval included zero. Thus this pilot supplies neither a resolved held-out mono/share winner nor a held-out ordering reversal or crossing. These unfavorable and inconclusive facts are retained explicitly.

By contrast, T was positive at each positive-noise level, with its pointwise intervals excluding zero. At the largest prescribed noise,

$$T = .0059465507, \quad 95\% \text{ interval } [.0035536102, .0082431507].$$

This supports decoder-policy dependence of the *relative risk comparison* on the fixed learned-feature example. It does not show that the full asymptotic Bernoulli endpoint ordering transfers. The exact decomposition

$$T = (R_{\text{mono},\text{fr}} - R_{\text{mono},\text{cal}}) - (R_{\text{sharing},\text{fr}} - R_{\text{sharing},\text{cal}})$$

identifies differential calibration benefit. At this noise the held-out benefits were .0122366692 for mono and .0062901185 for sharing. Their difference equals the reported contrast. This is an operational explanation of the observed contrast; attributing it causally to a specific gate subset requires checking the saved per-image and gate records rather than assuming the Bernoulli explanation automatically.

8.27.3 Audit, relevance and stopping decision

An independent saved-artifact audit verified the extraction provenance and indices, normalized arrays, checkpoint/data hashes, bootstrap and five evaluation cases, and 978 calibration nodes.




Figure 8.8: Held-out sharing-minus-mono reconstruction MSE against absolute Gaussian code-noise standard deviation, with conditional paired-image bootstrap intervals. These are operational learned activation channels, not validated semantic concepts. The two policy curves have positive point estimates but their individual risk-difference intervals cross zero throughout. Their paired contrast is more precisely estimated than either difference alone.




Figure 8.9: Paired policy contrast from the same saved five cases, without additional evaluation. Its pointwise conditional bootstrap intervals exclude zero at positive noise. The contrast measures differential calibration benefit; it does not resolve either individual mono/share risk ordering or demonstrate a reversal.

Recomputed per-image Gaussian losses agreed within 5.33×10^{-15} . The independent audit did not duplicate the ResNet forward inference. Numerical slack and floating arithmetic are not directed-rounding certification.

The empirical contribution is narrow but relevant: real learned nonnegative activations show a measurable relative-risk policy effect with the encoder and corruption held fixed. The effect coexists with an unresolved held-out risk ordering and a failed transfer of the training clean advantage. Arbitrary channels remain an unresolved link to semantic monosemanticity, and no real-model critical exponent, all-noise boundary or adversarial claim follows.

Stop the fixed pilot after these five levels. Do not replace the channel pair, increase noise until a crossing appears, or retune selection. The saved-record mechanism check below uses only the existing inputs and outputs; it cannot turn this inconclusive ordering into a resolved reversal. Stronger semantic or full-model claims require separate evidence and remain outside this completed pilot.

8.27.4 Saved-record gate and contribution audit

The predeclared diagnostic partition was target channel 0 exactly zero versus positive. All 512 test images and all five saved noise levels were retained; no new fit, corruption value or bootstrap was generated. There were 271 zero-target and 241 positive-target images. For the zero group, mono’s clean output-0 means were exactly at zero. Sharing’s median absolute distance from that gate was .16482873, with minimum .00101901; 32.47% of those means were positive despite zero target. This supports the presence of the proposed gate-placement distinction, but sharing distances are heterogeneous rather than uniformly separated.

For each fixed policy contrast, the group contribution is the sum of its per-image contrasts divided by 512, not by the group size. Thus the two contributions add to the original overall mean. Their complete decomposition is

Noise multiplier	Target 0 zero	Target 0 positive
0	.000000000	.000000000
.05	.000084991	.000013015
.10	.000320524	.000078932
.20	.001180940	.000416779
.40	.004083858	.001862693

At the largest noise the zero group contributes approximately 68.7% of the aggregate contrast. Its mean calibration benefits were .0270890 for mono and .0193734 for sharing. The positive group had negative benefits, $-.0044645$ for mono and $-.0084218$ for sharing: calibration increased this group’s held-out conditional error while improving aggregate risk. Sharing was harmed more on that group, so it too contributes positively to the mono-favoring policy contrast. No unfavorable subgroup effect has been discarded.

Per-output recomputation matched the saved per-image total losses exactly in floating arithmetic. Output 0 supplies the contrast; output 1’s aggregate calibration benefit was zero for mono and of order 10^{-10} for sharing, not a separate robust effect. The records support a gate-related explanation observationally, but do not isolate its causal role: amplitudes, target-zero status and coactivation differ jointly. No subgroup significance test or new confidence interval was computed. The original unresolved risk ordering remains unchanged. This completes the fixed pilot’s mechanism audit; stop without new pairs or noise searches.

8.28 Class-mapped evidence transfer: a resolved sharing advantage

This is a fixed operational transfer check, not a new theorem or a claim that hidden ResNet features are monosemantic. Its question is whether clean-trained scalar sharing of two independently named model-evidence coordinates retains an advantage over coordinate retention under the same Gaussian code corruption and decoder policies as the preceding pilot. The rejected external SAE export did not provide the required concept and model mapping; it was not used for this run.

8.28.1 Prospective feature definition and controls

The official frozen ImageNet ResNet18 IMAGENET1K_V1 head has category 281 **tabby** and category 207 **golden retriever**. These exact identifiers and names were verified before score inference. The two targets are $\text{ReLU}(\ell_{281})$ and $\text{ReLU}(\ell_{207})$, where ℓ is the genuine full classifier-head logit vector, not the previous single-spatial-cell hidden activation. Raw logits for all 1,000 categories are preserved. The zero threshold was fixed before outcomes and never changed.

Class supervision supplies a documented evidence association, not a faithful causal concept decomposition. Raw logits have an additive softmax gauge: shifting every logit leaves classification probabilities unchanged but changes these rectified coordinates. The fixed checkpoint gives a reproducible coordinate system; zeros do not certify concept absence. CIFAR cat/dog labels do not certify the tabby or golden-retriever subclasses. This test does not intervene on image inputs or measure full-network adversarial robustness.

One seeded permutation (NumPy seed 20261007) of CIFAR10 training indices provided disjoint 256 training, 256 calibration and 4,096 test images. Official preprocessing and the previously verified checkpoint were reused; no new model or SAE was trained. The complete forward pass took 304.453 seconds. A prespecified training-only association gate required AUROC greater than .65 for both named scores against the respective coarse cat/dog labels. The values .9049523 and .8081940 passed. Held-out descriptive values were .8700349 and .8172911; test labels did not select or fit any reconstruction.

Each nonnegative target was divided by its training RMS without centering. Importances were (1, 2/3); scalar encoder squared energy was one; decoding was tied ReLU with free biases. Mono could retain either coordinate and its omitted coordinate's best constant error was included. The existing exact empirical bias-interval profiler and 256/512 angle grids, with one best-cell refinement per grid, selected the clean geometry. This finite search is not a global-optimum certificate. Zero fractions were .17578125 and .22265625, and coactivation was .640625.

Sharing's weights were approximately (.83541738, .54961605); mono retained the first coordinate. Training risks were .300735002 and .371810660, respectively: sharing gain .071075658. Both refined losses agreed, both refinements succeeded, and the original positive-gain/mixed-energy gates passed. The selected same-sign geometry differs from the theorem's opposite-sign Bernoulli critical branch.

8.28.2 Held-out risks and the unfavorable transfer limit

Encoders were frozen. Both models' zero-noise biases were fitted on the calibration split and used for the frozen policy; the calibrated policy refitted biases on that same split at each noise. Noise standard deviation was the fixed multiplier times .7308510286, the square root of the mean training variance of the two RMS-normalized targets. Both models received the same independent scalar Gaussian code corruption. Conditional Gaussian moments, scalar bias localization

and their assumptions are those already derived in this volume. No fitted scaling expression was introduced.

At zero noise, held-out risks were .333790049 (sharing) and .404845579 (mono). Their difference was $-.071055530$ with paired 95% bootstrap interval $[-.08898894, -.05201580]$. Thus the clean training advantage transferred in this operational test. Every prescribed noisy difference also favored sharing. The table reports total importance-weighted test MSE for sharing minus mono; positive values would favor mono. The prespecified primary outcome was the calibrated difference at multiplier .4; the frozen difference and policy contrast at .4 were secondary, with other levels retained as the fixed descriptive curve:

Multiplier	Frozen difference	Calibrated difference	Calibrated 95% interval
0.00	-0.07105553	-0.07105553	$[-.08898894, -.05201580]$
0.05	-0.07113327	-0.07115480	$[-.08908896, -.05210769]$
0.10	-0.07130387	-0.07137251	$[-.08931704, -.05230341]$
0.20	-0.07185323	-0.07193311	$[-.08995778, -.05278479]$
0.40	-0.07388619	-0.07326888	$[-.09160315, -.05388520]$

All ten model/noise calibration comparisons met the unchanged numerical gap target 10^{-7} with arithmetic slack 10^{-12} . These are audited floating-arithmetic bounds, not directed-rounding proof certificates. Bootstrap intervals use 2,000 paired image resamples and are pointwise, conditional on fitted models; they omit training/calibration uncertainty.

At multiplier .4, calibrated risks were .398223291 and .471492166, respectively. There was *no ordering reversal* on the fixed grid; sharing’s point advantage was slightly larger than at zero noise. This does not reproduce a fragile sharing phase or the critical exponent. It is evidence that clean benefits of packing named evidence coordinates can persist under this declared corruption, preventing a universal interpretation that monosemantic retention must win whenever noise is added.

8.28.3 Calibration contrast, independent audit and stop

The contrast $\Delta_{\text{calibrated}} - \Delta_{\text{frozen}}$ was negative at multipliers .05 and .10 with intervals below zero, unresolved at .20, and positive at .40: .000617314 with interval $[-.000346820, .000899737]$. Calibration therefore helped sharing relatively more at the lowest positive levels and mono relatively more at the largest level. It never changed the winner. At .40, both models’ held-out risks improved under calibration; improvements were .002236486 for sharing and .002853800 for mono.

Independent checks verified 9 extraction and 38 evaluation manifest files, fixed image/label/-category mapping, raw-logit rectification, RMS scaling, training AUROC by pairwise tie-aware counting, clean grid losses, all 896 calibration nodes, conditional per-image losses and paired bootstrap records. Maximum independently recomputed risk discrepancy was 4.44×10^{-15} . No additional forward pass was performed by this saved-record audit. Original data, source snapshots, settings, failures and results remain unchanged under `project1_toy/head_transfer_20261007`.

Stop this pair and grid. Do not change class IDs, raw-logit threshold, noise range or decoder after seeing no reversal. The result is stronger than the earlier unresolved channel pilot as evidence of a held-out sharing benefit, yet it is not hidden-feature monosemanticity, critical branch validation, arbitrary-load theory or a full-model robustness result.




Figure 8.10: Fixed class-evidence reconstruction comparison. Negative risk differences favor sharing. Pointwise conditional intervals remain below zero; this is a one-pair operational result, not a real-model phase boundary.




Figure 8.11: Saved decoder-policy contrast. The .20 interval includes zero. A positive contrast favors mono relative to frozen evaluation, but does not imply that mono's total reconstruction risk is lower.

Chapter 9

Project 2: the correct meaning of the covariance recursion

Learning objective. Prove the fixed-map identity, identify its spectral regimes, and avoid mistaking it for a theorem about retraining.

9.1 A fixed linear map

Suppose a population is repeatedly transformed by $x_{g+1} = Gx_g$, with the same symmetric G at every generation. If $C_g = \text{Cov}(x_g)$, then

$$C_{g+1} = GC_gG, \quad C_g = G^g C_0 G^g. \quad (9.1)$$

The proof is simply $\text{Cov}(Ax) = A \text{Cov}(x) A^\top$, followed by induction. Means evolve separately as $\mu_g = G^g \mu_0$.

Diagonalize $G = U \Lambda U^\top$ and let $\tilde{C}_0 = U^\top C_0 U$. Then

$$(U^\top C_g U)_{kl} = (\lambda_k \lambda_l)^g (\tilde{C}_0)_{kl}. \quad (9.2)$$

A direction with initially positive variance has immediate loss if $\lambda_k = 0$, gradual decay if $0 < \lambda_k < 1$, preservation if $\lambda_k = 1$, and growth if $\lambda_k > 1$. Positive semidefiniteness does not imply $\lambda_k \leq 1$.

If C_0 is positive definite, total variance tends to zero exactly when all eigenvalues are below one. If C_0 is singular, only directions on its support matter; a growing eigenvalue with no initial variance does not contribute. Even this elementary statement needs the support qualification.

9.2 Why rank deficiency does not imply gradual decay

A rank- m projector has eigenvalues one and zero. Applying it once discards its nullspace; applying it again changes nothing. Thus a bottleneck can cause a one-time loss without progressive generational collapse. A compression argument alone cannot establish a rate of repeated degradation.

Moreover, eigenvectors are generally mixtures of named concepts. A rare concept indexed by i does not automatically have an eigenvalue “ λ_i ” that determines its fate. Track its projections onto the eigenspaces, or directly track its decoded probability.

9.3 The same margin can mean shrinkage, stability or growth

Consider

$$G_c = \frac{c}{2} \begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix}, \quad c > 0. \quad (9.3)$$

Both features have $M_i = 1/\sqrt{2}$ for every c . The eigenvalues are c and zero. With $C_0 = I$, for $g \geq 1$ the trace is c^{2g} . For $c = 1/2$ it decays, for $c = 1$ it is constant after the first step, and for $c = 2$ it explodes.

Therefore the same normalized leakage score cannot determine recursive stability. The two questions can share a matrix without being controlled by the same scalar. Counting eigenvalues exactly equal to one is also not a robust monosemanticity measure: a tiny global rescaling changes the count while leaving the directions and M_i unchanged.

9.4 Adaptive maps and injected variation

If $x_{g+1} = A_g x_g + \epsilon_g$, with independent zero-mean ϵ_g of covariance Q_g , then

$$C_{g+1} = A_g C_g A_g^\top + Q_g. \quad (9.4)$$

Writing $P_{g:k} = A_{g-1} \cdots A_k$ and $P_{k:k} = I$ gives

$$C_g = P_{g:0} C_0 P_{g:0}^\top + \sum_{k=0}^{g-1} P_{g:k+1} Q_k P_{g:k+1}^\top. \quad (9.5)$$

Changing maps need not commute, so one cannot raise a single initial Gram matrix to a power. For nonnormal general maps, per-step eigenvalues alone do not control transient amplification; products and singular values matter.

9.5 Finite-sample shrinkage is estimator-specific

For N independent samples from a distribution with covariance C , the maximum-likelihood-style covariance estimator with denominator N and an estimated sample mean satisfies

$$\mathbb{E}[\hat{C} \mid C] = \frac{N-1}{N} C. \quad (9.6)$$

This follows from subtracting the random sample mean. The unbiased estimator with denominator $N-1$ has expectation C . These facts do not mean that every generative training loop has a universal shrinkage factor $1 - c/N$.

If a particular simplified recursion is $\mathbb{E}[C_{g+1} \mid C_g] = \kappa G C_g G$, then its expected solution is $\mathbb{E}C_g = \kappa^g G^g C_0 G^g$ for fixed G . Setting $G = I$ gives $\kappa^g C_0$, not $\kappa^{2g} C_0$. Almost-sure behavior and expected behavior may differ; neither should be substituted for the other.

Status of this result. The covariance identities are valid and useful baselines. They describe the specified linear processes. They do not prove that diffusion self-training obeys a fixed Gram-matrix map, or that monosemanticity must decay.

Project 2: an exact nonlinear generation operator

Learning objective. Use the same feature decoder from Project 1 to define and solve a genuine repeated-generation toy process.

10.1 Define the transition before discussing collapse

For the detector of Chapter 4, let

$$K_g(b' | b) = \mathbb{P}_\xi\{\text{decode}(W_g b + \xi) = b'\}. \quad (10.1)$$

This is a Markov kernel: every column is nonnegative and sums to one. If P_g is the distribution of current features, the synthetic distribution is

$$Q_{g+1} = K_g P_g. \quad (10.2)$$

With real-data replacement fraction $\alpha \in [0, 1]$, define

$$P_{g+1} = \alpha P_0 + (1 - \alpha) K_g P_g. \quad (10.3)$$

This is a completely specified nonlinear-decoder self-consumption model. At fixed $W_g = W$, it repeatedly passes data through the same channel. It becomes a retraining model only when a learning rule for W_g is supplied, as in Chapter 12.

10.2 The pair kernel in closed form

Use the code $h = a(b_1 - b_2) + \xi$ from Chapter 5 and state order 00, 01, 10, 11. For a source state b , set $d = a(b_1 - b_2)$. Then

$$K(01 | b) = \Phi\left(\frac{-a/2 - d}{\sigma}\right), \quad (10.4)$$

$$K(10 | b) = \Phi\left(\frac{d - a/2}{\sigma}\right), \quad (10.5)$$

$$K(00 | b) = 1 - K(01 | b) - K(10 | b), \quad (10.6)$$

$$K(11 | b) = 0. \quad (10.7)$$

The last equality is architectural: the two threshold events cannot occur together. Project 1's one-generation error and Project 2's redistribution are now computed from exactly the same conditional events.

At zero noise, the states 00, 01, 10 reproduce themselves, while 11 maps to 00. If the initial inputs are independent, the joint probability $p_1 p_2$ is lost immediately. Subsequent iterations cause no additional change. This is a clear feature-combination loss, but it is not progressive decay of all features and does not require changing G .

10.3 An explicit order parameter for a defined loss

For any pair (i, j) with nonzero initial co-occurrence, define joint retention

$$r_{ij,g} = \frac{\mathbb{P}_g(b_i = b_j = 1)}{\mathbb{P}_0(b_i = b_j = 1)}. \quad (10.8)$$

For the antipodal pair with $\alpha = 0$, $r_{12,g} = 0$ for all $g \geq 1$. With mixing, $\mathbb{P}_{g+1}(11) = \alpha \mathbb{P}_0(11)$, so $r_{12,g} = \alpha$ for $g \geq 1$. This is an exact indicator of a particular lost structure, not a universal definition of monosemanticity.

Marginal retention $\mathbb{P}_g(b_i = 1)/\mathbb{P}_0(b_i = 1)$ measures something else. Noise can increase a marginal through false positives while destroying correct combinations. A useful collapse evaluation needs both marginal and joint fidelity, plus semantic correctness.

10.4 What the eigenvalues mean here

At fixed K , $P_g = K^g P_0$ without mixing. A column-stochastic matrix always has eigenvalue one, representing a stationary distribution. Eigenvalues inside the unit disk govern how distributions approach stationary behavior under suitable irreducibility and aperiodicity conditions; reducible chains can have several stationary classes.

Stationarity is not the same as fidelity to P_0 . A process can converge stably to the wrong distribution. Thus a “stability eigenvalue” alone does not measure useful concept survival. Track the stationary bias and the transient rate separately.

For $\alpha > 0$, the affine fixed point is

$$P_* = \alpha \{I - (1 - \alpha)K\}^{-1} P_0 = \alpha \sum_{k=0}^{\infty} (1 - \alpha)^k K^k P_0. \quad (10.9)$$

The inverse exists because $(1 - \alpha)K$ has spectral radius below one. This mixture over channel depths makes the effect of real-data replenishment transparent.




Figure 10.1: Exact pair-kernel calculations with $a = 1$, $\sigma = 0.35$ and initial activation probability 0.2. Marginal frequencies can rise even though the joint state 11 is eliminated from generated outputs. Twenty-percent real-data mixing preserves some joint mass through replenishment.

10.5 A monosemantic comparator

A code retaining only feature 1 with zero output for feature 2 also has a kernel, but every output with $b'_2 = 1$ has zero probability. It drops all of feature 2, not only its simultaneous occurrence with feature 1. This is the generative analogue of the Project 1 capacity tradeoff: a more isolated code can be accurate on its retained concept while representing less of the full distribution.

The comparison should therefore ask how much meaningful information survives at fixed capacity, rather than declaring any noninterfering representation automatically superior.

Before moving on. Compute the zero-noise pair distribution after one and two generations. Explain why covariance alone would not be the clearest account of the lost joint concept.

Chapter 11

The mathematical link between the two projects

Learning objective. Prove how one-generation feature errors bound accumulated distribution drift, and identify what the bound cannot say.

11.1 Use the generation mechanism as a coupling

Under the kernel of Chapter 10, sample $b \sim P_g$ and then $b' \sim K_g(\cdot \mid b)$. This constructs both the current and generated distributions on one probability space. Let

$$e_{i,g} = \mathbb{P}(b'_i \neq b_i), \quad R_g = \sum_i I_i e_{i,g}. \quad (11.1)$$

These are exactly the one-step error quantities analyzed in Project 1, evaluated under the current distribution and geometry.

Theorem 11.1 (Feature-error control of distribution drift). *For any such binary-feature generation kernel,*

$$|\mathbb{P}(b'_i = 1) - \mathbb{P}(b_i = 1)| \leq e_{i,g}, \quad (11.2)$$

$$\text{TV}(K_g P_g, P_g) \leq \mathbb{P}(b' \neq b) \leq \sum_i e_{i,g}. \quad (11.3)$$

Consequently, without real-data mixing,

$$\text{TV}(P_t, P_0) \leq \min \left(1, \sum_{g=0}^{t-1} \sum_i e_{i,g} \right). \quad (11.4)$$

Proof. For a binary coordinate, $|\mathbb{E}(b'_i - b_i)| \leq \mathbb{E}|b'_i - b_i| = e_{i,g}$. For any event A , the indicators of $b \in A$ and $b' \in A$ can differ only if $b \neq b'$, so the difference in their probabilities is at most $\mathbb{P}(b' \neq b)$. Taking the supremum over A gives the coupling bound for total variation. A union bound over coordinates gives the next inequality. Finally use the triangle inequality across generations, and the fact that total variation is at most one. $\square$

11.2 Importance-weighted version

Define the weighted Hamming cost $d_I(b, b') = \sum_i I_i |b_i - b'_i|$. The optimal-transport distance using this cost is no larger than the cost of our particular coupling:

$$W_{d_I}(P_g, K_g P_g) \leq R_g. \quad (11.5)$$

If all $I_i \geq I_{\min} > 0$, then $\sum_i e_{i,g} \leq R_g / I_{\min}$ also converts the weighted risk into a total-variation bound. This may be loose when importance weights span many orders of magnitude. Reporting per-feature and joint errors remains more informative than a single weighted average.

11.3 A replenishment bound

Suppose $\text{TV}(K_g P_g, P_g) \leq \varepsilon_g$, and use Equation (10.3). Write $D_g = \text{TV}(P_g, P_0)$. Convexity and the triangle inequality give

$$D_{g+1} \leq (1 - \alpha)(D_g + \varepsilon_g). \quad (11.6)$$

Starting at $D_0 = 0$ and assuming $\varepsilon_g \leq \varepsilon$ yields, for $\alpha > 0$,

$$D_t \leq \min \left(1, \frac{1 - \alpha}{\alpha} \{1 - (1 - \alpha)^t\} \varepsilon \right). \quad (11.7)$$

This derivation allows K_g to depend on the current distribution; it does not require the same kernel every round. It is a bound on drift under a controlled per-round error, not proof that a small real-data fraction fixes any possible learning algorithm.

11.4 When the bound is informative

If errors are rare and concentrated in a small set of features, the result links a measurable one-step mechanism to a finite-horizon prediction. It can also be vacuous: summing many errors or many generations can exceed one. A stronger paper would exploit transition structure, cancellation, feature dependencies or rare-state hazards to obtain sharper predictions than the union bound.

The implication runs one way. A small coupling error guarantees small drift. A large error does not force drift: a process can swap states while preserving their overall distribution. Conversely, nearly unchanged marginal probabilities can hide lost co-occurrences. These are testable limitations that help choose order parameters.

11.5 Where monosemanticity enters the actual research claim

Geometry influences the one-step error through the exact risk or interference bound. The generation kernel propagates those errors. The candidate scientific chain is therefore

$$\text{feature organization} \longrightarrow \text{specific decoding errors} \longrightarrow \text{recursive distribution change}. \quad (11.8)$$

The theorem establishes the middle-to-last link for the specified toy mechanism. The first link requires either a geometry-controlled theorem or an intervention. The real-model experiment must test the corresponding chain rather than merely correlate an SAE statistic with generation number.

This is a mathematically coherent way to combine the projects. The coupling inequalities themselves are standard probability tools; the publishable contribution would need to be the nontrivial geometry-dependent prediction, its sharpness and its empirical support.

Before moving on. Construct a two-state process that flips both states with equal probability mass. Its per-sample error is one but its distribution does not change. Explain why this does not contradict the theorem.

What changes when the model is genuinely retrained

Learning objective. Write a learning operator and derive its local response, rather than assuming that geometry stays fixed.

12.1 Training is part of the dynamics

Let θ collect encoder, decoder and threshold parameters. For a distribution P on toy states, define a regularized population objective

$$\mathcal{L}(\theta; P) = \sum_b P(b) \ell(\theta, b) + \lambda \Omega(\theta), \quad \theta^*(P) \in \arg \min_{\theta} \mathcal{L}(\theta; P). \quad (12.1)$$

An explicit tie-breaking or optimization rule is needed if there are several solutions. A reconstruction or smooth detection surrogate can be used for training; if it differs from the evaluation risk, state that difference.

The population training-and-generation map is

$$F(P) = \alpha P_0 + (1 - \alpha) K(\theta^*(P)) P. \quad (12.2)$$

This includes refitting the model each generation. A finite-sample version also samples a dataset, estimates parameters and may warm-start from the previous generation. Those choices affect the map and should be logged.

12.2 Deriving the response of the learned parameters

Assume a differentiable local optimum satisfying $\nabla_{\theta} \mathcal{L}(\theta^*; P) = 0$, with nonsingular Hessian $H = \nabla_{\theta}^2 \mathcal{L}(\theta^*; P)$. For a distribution perturbation δP with $\sum_b \delta P(b) = 0$, differentiate the stationarity condition:

$$H \delta \theta + \sum_b \delta P(b) \nabla_{\theta} \ell(\theta^*, b) = 0. \quad (12.3)$$

Thus

$$D_P \theta^*[\delta P] = -H^{-1} \sum_b \delta P(b) \nabla_{\theta} \ell(\theta^*, b). \quad (12.4)$$

This is an implicit-function calculation. It is local and assumes a smooth, identifiable branch of solutions. Symmetries, singular Hessians, ReLU boundaries and optimizer changes can invalidate its use unless handled by gauge fixing, regularization, smooth surrogates or a more general analysis.

12.3 The local generation Jacobian

Differentiate Equation (12.2):

$$DF(P)[\delta P] = (1 - \alpha) [K\delta P + \{D_\theta K[D_P\theta^*[\delta P]]\}P]. \quad (12.5)$$

The first term propagates a changed input distribution through the current decoder. The second term captures how retraining changes the decoder itself. A fixed- G or fixed- K argument omits this second mechanism.

At a smooth fixed point, eigenvalues of this Jacobian on the probability-simplex tangent space provide a local stability criterion: spectral radius below one gives local asymptotic stability under the usual finite-dimensional differentiability conditions. A stable fixed point can still be far from P_0 , so quantify its fidelity separately.

12.4 A tractable staged experiment

Start with a small finite family of candidate geometries, such as isolated versus paired codes with stated norm constraints. Evaluate the training objective exactly under P_g , select the minimizer by a deterministic rule, and generate using its exact kernel. This yields a genuinely adaptive piecewise-defined recursion that can be enumerated and checked.

Then move to learned continuous dictionaries. At each generation, retrain on sampled data and compare observed one-step errors with the prediction conditional on the measured G_g . If the geometry changes, report it. The goal is not to force a fixed-matrix prediction onto a changing model.

12.5 Order parameters worth testing

Use a small set of quantities with distinct meanings: per-feature false negatives and false positives; reference-distribution feature retention; joint-state retention; normalized row leakage; decoder reconstruction fidelity; and the local transition or learning Jacobian when it is estimable. Candidate predictors should be evaluated on later generations or held-out regimes, not selected after seeing all collapse outcomes.

A basis-invariant covariance rank measures occupied linear directions. It does not by itself identify named concepts. A neuron purity measure depends on a semantic reference. Neither should be relabeled as the other. The research opportunity is to show which geometric and semantic measurements add predictive information beyond frequency, norm, ordinary reconstruction error and training loss.

Status of this result. Equations (12.4)–(12.5) are conditional local derivations. Solving a global nonlinear learning recursion and proving that a particular semantic scalar forecasts real-model collapse remain research work.

Diffusion: the bridge that is valid, and the one to test

Learning objective. Understand the denoising equations and why a feature Gram matrix is not automatically a diffusion-generation operator.

Source map: DDPM, score-based generative modeling and the recursive-diffusion references in Appendix D. The Gaussian calculation below is derived independently..

13.1 A diffusion model in five equations

A forward diffusion process gradually adds noise to a clean example x_0 . Let $\beta_\tau \in (0, 1)$, $\alpha_\tau = 1 - \beta_\tau$ and $\bar{\alpha}_\tau = \prod_{k=1}^\tau \alpha_k$. Then

$$q(x_\tau | x_{\tau-1}) = \mathcal{N}(\sqrt{\alpha_\tau}x_{\tau-1}, \beta_\tau I). \quad (13.1)$$

Because a composition of affine Gaussian channels is Gaussian,

$$x_\tau = \sqrt{\bar{\alpha}_\tau}x_0 + \sqrt{1 - \bar{\alpha}_\tau}\epsilon, \quad \epsilon \sim \mathcal{N}(0, I). \quad (13.2)$$

A neural network is trained to predict the added noise, often with the simplified objective

$$\mathcal{L}_{\text{noise}} = \mathbb{E}_{x_0, \tau, \epsilon} \|\epsilon - \epsilon_\theta(x_\tau, \tau)\|_2^2. \quad (13.3)$$

The pointwise minimizer of squared error is a conditional expectation. The prediction is therefore related to the score of the noisy data distribution, not a literal reconstruction Gram matrix. One common reverse-process mean parameterization is

$$\mu_\theta(x_\tau, \tau) = \frac{1}{\sqrt{\alpha_\tau}} \left(x_\tau - \frac{\beta_\tau}{\sqrt{1 - \bar{\alpha}_\tau}} \epsilon_\theta(x_\tau, \tau) \right). \quad (13.4)$$

Sampling uses a sequence of reverse transitions, often with additional noise:

$$x_{\tau-1} = \mu_\theta(x_\tau, \tau) + s_\tau z, \quad z \sim \mathcal{N}(0, I), \quad (13.5)$$

with a specified variance or sampler convention. The sampler, schedule, truncation and network all affect the generated distribution.

13.2 Two clocks, two different processes

Diffusion time τ indexes noise removal within one generation of one sample. Training generation g indexes fit-sample-refit cycles across models. A denoising step can reduce uncertainty while a correctly sampled full generative process preserves diversity. Confusing these clocks makes “denoising contracts covariance, so diffusion must collapse” an invalid argument.

Write the full recursive process as

$$\theta_g = \text{Train}(D_g), \quad Q_g = \text{SampleLaw}(\theta_g), \quad D_{g+1} \sim \alpha P_0 + (1 - \alpha)Q_g. \quad (13.6)$$

Here data size, optimization budget, initialization, conditioning and sampling parameters are part of the operator. There is no justified substitution $Q_g = G_g P_g$ merely because an SAE dictionary has a Gram matrix G_g .

13.3 A Gaussian denoiser that can be solved completely

Let the generation- g data distribution be $X \sim \mathcal{N}(0, C_g)$ and observe $Y = X + \epsilon$ with $\epsilon \sim \mathcal{N}(0, \sigma^2 I)$. Gaussian conditioning gives

$$\mathbb{E}[X | Y] = A_g Y, \quad A_g = C_g(C_g + \sigma^2 I)^{-1}. \quad (13.7)$$

This is the optimal population denoiser refitted to the current covariance. If the next generation's data consist of these posterior means, then

$$C_{g+1} = A_g(C_g + \sigma^2 I)A_g^\top \quad (13.8)$$

$$= C_g(C_g + \sigma^2 I)^{-1}C_g. \quad (13.9)$$

The eigenvectors are unchanged and each positive eigenvalue evolves by

$$c_{g+1} = \frac{c_g^2}{c_g + \sigma^2}. \quad (13.10)$$

It decreases strictly when $\sigma^2 > 0$ and converges to zero. To prove the limit, monotonicity gives $c_g \downarrow L \geq 0$; continuity implies $L = L^2/(L + \sigma^2)$, hence $L = 0$. When $c_g \ll \sigma^2$, the next value is approximately c_g^2/σ^2 , explaining rapid late-stage shrinkage.

This is a genuine refitting recursion in an analytically tractable denoising model. It is not a full DDPM theorem and does not yet distinguish monosemantic from rotated Gaussian representations.

13.4 Posterior sampling changes the answer completely

The conditional covariance is

$$V_g = C_g - C_g(C_g + \sigma^2 I)^{-1}C_g. \quad (13.11)$$

If the next sample is drawn from the exact posterior rather than set equal to its mean, the law of total covariance gives

$$\text{Cov}(X') = \text{Cov}(\mathbb{E}[X | Y]) + \mathbb{E}[\text{Cov}(X | Y)] = C_g. \quad (13.12)$$

Thus the same denoiser accompanied by the correct sampling variation preserves the original distribution. Removing sampling variance creates the collapse in Equation (13.9). This is why a real diffusion sampler must be modeled or held fixed, not discarded from the argument.




Figure 13.1: Exact Gaussian population refitting with initial variance 1 and observation-noise variance 0.2. Reusing posterior means contracts variance; exact posterior sampling preserves it. This is a didactic denoising model, not a trained image-diffusion experiment.

13.5 Where a representation-level contribution can enter

Gaussian isotropy cannot identify semantic axes: orthogonal rotations preserve the distribution. To study monosemanticity, use known sparse non-Gaussian factors or a controlled image generator with labeled concepts. Then test whether a feature’s interference and one-step error forecast later loss of its frequency or combinations.

On a real model, an SAE decoder dictionary D_g provides $D_g^\top D_g$, a measure of dictionary overlap. It is not the same as a tied encoder-decoder reconstruction map unless those constraints are imposed. A diffusion denoiser’s input-output sensitivity is a Jacobian, generally rectangular, input-dependent and noise-time-dependent. These are different objects and should retain different symbols.

A measured relation between dictionary overlap, causal feature interventions and future degradation would be evidence for the proposed mechanism. Calling the relation exact without measuring the missing decoder and sampler dynamics would not be justified.

13.6 Existing work the new claim must distinguish itself from

The MAD work studies self-consuming generative loops and the role of real-data supply. The accumulating-data work shows that replacement and accumulation can behave differently. A June 2026 paper by Khelifa, Turner and Venkataramanan already gives a spectral characterization of recursive diffusion collapse driven by reverse-process truncation, even under ideal score estimation. Therefore “the first spectral account of diffusion collapse” is not an available claim based on this reading.

The intended distinction is a tested mechanism involving sparse semantic features, interference, robust decoding and subsequent retention. Keep reverse truncation fixed and check whether it explains the apparent effect before attributing that effect to feature geometry.

Before moving on. Identify the covariance term that posterior sampling restores in the Gaussian example. Explain why posterior means alone lose it.

Experiments that test the derivations

Learning objective. Design experiments that can support or falsify the mechanism, with a realistic route from exact toys to a vision diffusion model.

14.1 Stage A: exact identities and finite-state calculations

The accompanying script already checks conditional moments from the anchor paper, pair risks by Monte Carlo, the pair transition matrix, the fixed-linear covariance identity, a same- M counterexample, the Gaussian recursion and the linear tight-frame comparison. These are unit-level mathematical validations, not the final research experiments.

Reproduce the exact pair boundary for several importance ratios and for both scale conventions. For each setting, compare the normal-CDF formula with independent random trials and binomial uncertainty. Include rare-but-important and common-but-low-importance combinations rather than varying all frequencies together.

For the transition model, enumerate the distribution rather than assuming it remains factorized. Track the entire four-state law, joint retention and marginal errors. Check the coupling inequalities directly. This stage should expose mistakes before any expensive training is run.

14.2 Stage B: learn the geometry in controlled toys

Train tied nonlinear autoencoders on planted features with a weighted reconstruction objective. Use a deliberately limited grid: a few loads (for example 1, 1.5, 2 and 4), several activation rates, two or three importance-decay schedules and multiple seeds. The values are starting design choices, not a mandatory exhaustive Cartesian product.

At each checkpoint, measure G , column norms, active-feature coverage, reconstruction error, statistical feature-detection error and an optional separately labeled adversarial metric. Predict risk using the measured geometry before inspecting held-out outcomes. This tests the conditional theory. Separately ask whether geometry itself is predictable from the training controls; that tests the stronger phase-diagram objective.

Necessary comparators include capacity-matched isolated features, a random dictionary, a learned unconstrained dictionary and a semantic or sparsity-promoting intervention. Control encoder norm or report matched-norm sensitivity analyses. Include a clean-performance comparison so a method is not credited merely for dropping difficult concepts.

14.3 Stage C: recursive training on the same toys

For each chosen starting model, generate a fixed-size synthetic dataset and train the next model under a fixed budget. Compare pure replacement with a specified real-data fraction. Keep a

real-only repeated-training control to estimate ordinary training variability. Distinguish warm starts from fresh initialization; choose one primary protocol and treat the other as an ablation.

The central prediction is prospective: use geometry and one-step errors at generation g to predict feature or joint losses at $g + 1$ and later. A model that only explains already-observed collapse is a description, not an onset predictor. Hold out seeds or regions of the sparsity/load grid when assessing prediction quality.

Analyze false positives and false negatives separately. For a rare feature, a small false-positive rate can dominate its marginal frequency. A seemingly preserved activation rate does not establish preserved semantic content.

14.4 Stage D: one controlled image generator

Before a general image dataset, use images generated from known factors: shape identity, location, color and a rare co-occurring attribute, with explicit train and evaluation distributions. This provides real pixel inputs and genuine convolutional or diffusion learning while retaining concept ground truth. It is a mechanistic testbed, not a benchmark contribution.

Train a small image diffusion model, then repeat the same controlled self-training protocol. Use the known rendering factors or a validated held-out evaluator to measure marginal and joint fidelity. Test feature robustness on fixed real reference images under controlled image corruptions or denoising conditions. Avoid defining a concept detector solely from the same model whose collapse is being measured.

14.5 Stage E: a small natural-image diffusion demonstration

Use one tractable vision setting, such as CIFAR-scale images, with a small established diffusion implementation. A pilot should measure wall-clock cost for one generation before committing to a long recursion. The primary comparison might use two representation/training conditions, two real-data fractions and three seeds over a short, informative horizon. The exact number of generations is determined by the pilot and the observed transition, not by a promise made before measuring compute.

This remains a real-model validation arm. If it fails to show the predicted mechanism, report that limitation and investigate the mismatch. It should not be reclassified as optional success evidence only after seeing unfavorable results.

14.6 How to measure geometry consistently

Fix the layer, diffusion-time sampling distribution and reference inputs. Fit a common diagnostic dictionary across checkpoints, or explicitly align independently learned dictionaries and report alignment uncertainty. Otherwise feature permutations, rescalings, splitting and merging can masquerade as semantic decay.

For a common dictionary, report reconstruction fidelity at every checkpoint: a fixed dictionary can fail because the basis changes, even if useful computation is preserved. With checkpoint-specific dictionaries, distinguish actual concept changes from changed coordinate systems. Use direct labeled-factor behavior as a reference wherever possible.

14.7 Interventions that support a mechanism

An intervention should change the hypothesized interference while monitoring clean quality, norms and coverage. Examples include a controlled coherence penalty in a toy model, a sparsity-promoting module, or a targeted direction manipulation with a random-direction control. A penalty that orthogonalizes everything may simply reduce stored information, so compare at matched capacity and inspect feature loss.

The strongest evidence is a linked change: the intervention alters measured geometry, changes the predicted one-step errors and changes later feature retention in the predicted direction. A correlation between low purity and late generations alone is much weaker.

14.8 Evaluation and uncertainty

Report training seeds as independent replications. For rare events, provide sample counts and uncertainty; zero observed events is not proof of zero probability. If a concept has true frequency p , a sample of size N contains none with probability $(1 - p)^N$, approximately e^{-Np} . To make this probability at most δ , one needs roughly $N \geq \log(1/\delta)/p$. This calculation belongs in the sampling plan.

For generative quality, combine a distributional image metric with concept-frequency and joint-frequency measures. Aggregate metrics can miss tail loss. Use held-out real reference data for evaluation, and never select thresholds or order parameters on the final test outcomes without reporting that selection.

14.9 A compact evidence table

Claim	Required evidence	A result that would weaken it
Geometry predicts one-step robustness	Analytic risk or bound agrees with held-out toy errors	Norm or feature frequency explains the effect equally well
Sparsity/load creates a tradeoff	Capacity-controlled risk crossing, plus learned-toy comparison	Crossing vanishes under energy matching
Errors predict recursive loss	Prospective prediction of later marginal or joint changes	Predictor adds nothing beyond current reconstruction error
Monosemanticity contributes causally	Controlled intervention changes the proposed chain	Apparent benefit is entirely due to dropping concepts
Mechanism transfers to diffusion	Same directional prediction and intervention in fixed-sampler experiments	Sampler truncation or evaluator drift explains the result

An eight-week program that preserves both projects

15.1 Status of this execution outline

This is the original September planning outline, retained as background rather than a renewed eight-week allowance. The October Project 1 results appear in Chapters 7–8. The current stop decision is to consolidate the established claims, compare their exact assumptions with the direct prior literature, and specify one bounded real-model transfer. Broad grids, additional toy families and recursive diffusion retraining are not automatically authorized by this historical outline. The full variable-load/importance boundary remains open.

Learning objective. Turn the derivations into deliverables without promising that a theorem outline already guarantees a conference paper.

15.2 One integrated paper structure

The intended paper can be organized around a precise question: how does sparse feature interference determine immediate decoding reliability and the retention of semantic structure across recursive training? Project 1 supplies the risk tradeoff; Project 2 supplies the propagation mechanism and predictive test. The shared toy model makes the connection explicit, while a diffusion demonstration tests how much of it survives outside the controlled setting.

The theoretical thread should remain active throughout: define assumptions, derive predictions, diagnose mismatches and refine the claim. Experimental work is equally part of the scientific argument. The document does not assign people to narrow permanent roles or prescribe authorship.

15.3 Weekly deliverables

Assuming a Monday 21 September 2026 start, eight full weeks end on 15 November; a tenth week ends on 29 November. If the actual kickoff differs, shift dates while preserving the ordering and compute gates.

Week	Dates	Concrete deliverable
1	Sep 21–27	Read the anchor paper using Volume I. Re-run the supplied checks. Fix task, norm budget, corruption location and two primary hypotheses. Run a diffusion cost pilot and a focused related-work check in parallel with mathematical work.

Week	Dates	Concrete deliverable
2	Sep 28–Oct 4	Reproduce pair risk and kernel. Train the first planted-feature model. Produce one robustness prediction plot and one recursive-distribution plot. Record how learned geometry differs from the solved pair.
3	Oct 5–11	Expand a small sparsity/load/importance grid. Validate conditional risk predictions on held-out seeds. Derive or tighten the part of the load-dependent theory supported by observed structure. Start the controlled-image diffusion loop.
4	Oct 12–18	Month-one checkpoint: a precise theorem statement with assumptions, verified toy evidence for both questions, and an operational real-model pipeline. Select the most informative regimes and eliminate redundant sweeps.
5	Oct 19–25	Test prospective collapse predictions and a geometry intervention. Run the primary natural-image diffusion conditions if the pilot supports the budget. Write the theory and methods sections using actual completed results.
6	Oct 26–Nov 1	Complete the primary recursive runs and matched controls. Investigate mismatches with the theory. Compare geometry predictors with frequency, norm, reconstruction-error and training-loss baselines.
7	Nov 2–8	Finish required seeds and decisive ablations. Assemble the complete paper narrative, plots and limitations. Reproduce the main results from clean configuration files.
8	Nov 9–15	Month-two checkpoint: integrated draft, checked proofs, final main experiments, uncertainty estimates and a reproducible artifact. Obtain technical review focused on the most vulnerable assumption.
9	Nov 16–22	If available: resolve reviewer-style objections and improve one transfer or intervention experiment. Avoid opening a second modality or unrelated task.
10	Nov 23–29	If available: finalize revisions, supplementary proofs and release materials; rerun only checks affected by changes.

15.4 What to reduce if the compute budget is tight

Reduce the number of architectures, modalities, redundant grid points and secondary metrics first. Keep the core mathematical question, both toy mechanisms, the decisive controls and one real-model test. A smaller, interpretable experiment is preferable to an underpowered collection of unrelated demonstrations.

If natural-image diffusion training is too expensive, a controlled-image diffusion model still tests a genuine nonlinear generative model with known concepts. Report the narrower generality honestly. Do not silently replace the real-model arm with another analytic plot and describe the original empirical objective as completed.

15.5 What would make the paper substantive

The standard identities in this volume are foundations. The eventual contribution needs a nontrivial result, such as a sharp learned-geometry tradeoff, a predictive feature-retention mechanism that survives controls, or a targeted intervention justified by theory and validated in a generative loop. The combined story should demonstrate why the two results illuminate each other.

An eight-week program can plausibly produce a focused submission if compute and participation support the plan. No document or toy calculation establishes main-track acceptance in advance. A spotlight is a distinction among accepted papers, not an easier fallback venue. The practical objective is a sound, novel and well-supported manuscript; venue decisions follow external review. The next step is to choose the primary corruption and resource convention, reproduce the analytic pair in a learned toy, and execute the diffusion pilot. This continues the same two-project program with explicit assumptions.

Appendix A

Formula ledger: retain, qualify or replace

Earlier statement	Correct status and replacement
$G = W^\top W$	Exact for the tied encoder geometry. An arbitrary SAE or diffusion model need not have this as its input-output map.
$M_i = G_{ii} / \ G_{i:}\ _2$	Exact alignment ratio for nonzero columns. Equals a radius only for a specified normalized linear decision problem.
$M_i < 1$ whenever an attacker can move the score $\sum \rho_{ij}^2 \approx pn/m$ inside the unrestricted attack margin	False. An isolated feature with $M_i = 1$ can still be perturbed.
$pn/m \leq \tau^{-2} - 1$ is a proved universal boundary	Not justified. Use state-averaged risk or a statistical interference variance.
Rare concepts have high M	Not established. Use the exact pair-risk equality or a stated sufficient error bound.
Monosemanticity is the limit $s \rightarrow 1$	Not implied. Frequency does not enter the fixed- G ratio directly.
$C_{g+1} = GC_g G$	Not an identity. Sparse activations can favor superposition; semantic organization and frequency are distinct.
Compression forces all non-retained modes to decay gradually	Exact for repeatedly applying a fixed linear map. Not a generic retraining law.
Eigenvalue-one fraction is a monosemanticity order parameter	False. A projector loses its nullspace in one step and is then stable. Eigenvalues above one can also occur.
κC_g recursion produces κ^{2g}	Unsupported. It is highly scale-sensitive and ignores semantic basis and initial support.
Two valid identities plus toy plots guarantee a paper	Algebraically inconsistent. The factor is κ^g for that covariance recursion.
Recursive synthetic training is automatically improvement	No. Novelty, strength of mechanism, empirical evidence and review determine the contribution.
	No. It is a recursive process; improvement is an outcome to measure.

Appendix B

Further derivations and worked exercises

B.1 Gaussian ReLU moments for continuous reconstruction

Let $Z = d + \sigma U$ with $U \sim \mathcal{N}(0, 1)$, and let ϕ, Φ be the standard normal density and CDF. Direct integration gives

$$A(d) = \mathbb{E} \text{ReLU}(Z) = \sigma \phi(d/\sigma) + d\Phi(d/\sigma), \quad (\text{B.1})$$

$$B(d) = \mathbb{E} \text{ReLU}(Z)^2 = (d^2 + \sigma^2)\Phi(d/\sigma) + d\sigma\phi(d/\sigma). \quad (\text{B.2})$$

For a target amplitude x , the conditional squared error is $x^2 - 2xA(d) + B(d)$. In the antipodal pair, use $d = x_1 - x_2$ for the first reconstruction and $-d$ for the second. Integrating this expression over the specified spike-and-uniform features yields an exact low-dimensional integral for the noisy continuous-amplitude risk. Evaluating an analytically specified integral is not symbolic regression.

To derive A , substitute $Z = d + \sigma U$ and integrate over $U > -d/\sigma$. Use $\int_a^\infty u\phi(u) du = \phi(a)$. For B , expand $(d + \sigma U)^2$ and use $\int_a^\infty u^2\phi(u) du = a\phi(a) + 1 - \Phi(a)$.

B.2 Mixture covariance includes a mean term

If real data have mean μ_0 and covariance C_0 , while generated data have mean μ_Q and covariance C_Q , a mixture with real fraction α has covariance

$$C_{\text{mix}} = \alpha C_0 + (1 - \alpha)C_Q + \alpha(1 - \alpha)(\mu_0 - \mu_Q)(\mu_0 - \mu_Q)^\top. \quad (\text{B.3})$$

The last term is between-component variation. Omitting it can make a drifting mean look like an unexplained covariance change. The finite-state kernel automatically accounts for this when moments are computed from the full law.

B.3 Exercise 1: an exact adversarial perturbation

For a correctly classified x and ℓ_2 budget, show that $\delta_* = -yf(x)\beta/\|\beta\|_2^2$ reaches the boundary. Its norm is $yf(x)/\|\beta\|_2$. This proves tightness of Equation (2.2). A slightly larger displacement crosses the boundary; equality only reaches it, and classification at equality depends on the tie convention.

B.4 Exercise 2: rare feature versus common nuisance

Take a fixed two-feature Gram matrix with unit diagonal and overlap ρ . Then $M_1 = 1/\sqrt{1 + \rho^2}$ independently of p_1, p_2 . But the statistical nuisance variance in the independent centered detector is $\rho^2 p_2(1 - p_2)$. Conditional robustness of feature 1 depends on its competitor's activity, not simply its own rarity. Its unconditional contribution to total risk also depends on its own false-positive and false-negative mixture.

B.5 Exercise 3: why a largest overlap and a sum differ

Compare one nuisance with overlap a and four nuisances with overlap $a/2$. Their total squared overlap is equal. Under isotropic unrestricted feature noise, the readout variance can agree. Under sparse nonnegative activations, their distributions of interference differ: one large jump is not the same random variable as a sum of four smaller jumps. The Bernstein bound records both V_i and B_i to reflect this distinction.

B.6 Exercise 4: solve the mixed fixed kernel

Starting from $P_{g+1} = \alpha P_0 + (1-\alpha)K P_g$, repeatedly substitute the previous step to obtain a finite geometric series plus $(1-\alpha)^g K^g P_0$. The final term vanishes for $\alpha > 0$ in the probability-vector norm, leaving the fixed-point series in Chapter 10. This is a statement about repeated use of a fixed channel; refitting generally changes K and prevents that same closed form.

B.7 Exercise 5: covariance does not identify concepts

Take a two-dimensional isotropic Gaussian and rotate it by 45 degrees. The covariance and distribution are unchanged, but coordinate axes are different. A covariance-based criterion alone cannot say which axes correspond to meaningful concepts. Non-Gaussian planted factors or external semantic measurements are necessary for that question.

B.8 Exercise 6: a sharply stated falsification

Suppose an SAE overlap score predicts late-generation image degradation until reconstruction error and feature frequency are included, after which it adds no held-out predictive information. This does not invalidate the exact toy algebra. It weakens the claim that the chosen geometry measure explains real-model collapse beyond simpler measurements. The paper should report and analyze that distinction rather than call the transfer confirmed.

Appendix C

Numerical verification record

All calculations here were performed on CPU with fixed seeds 20260921 for the main script and 20260922 for the additional checks. The scripts use NumPy, SciPy and exact SymPy algebra; figures use Matplotlib. They do not train a large neural model or establish publication novelty.

Check	Result
Anchor-paper conditional moments	1.5 million simulated examples; all conditional means and variances within 0.002 of the exact values.
Corrected label-noise J crossing	$\eta = 0.2427153958$ at $S = 0.2$.
Input-noise J crossing	$\lambda = 0.5479269441$ at $S = 0.2$.
Pair boundary, equal stored amplitude	$\sigma = 0.6583495016$ at $p_1 = p_2 = 0.2$, equal importance.
Pair boundary, equal encoder energy	$\sigma = 0.2799389129$ for the same probabilities and importance.
Equal-energy pair risk at $p = 0.2$, $\sigma = 0.3$	Exact 0.261396; independent simulation 0.261172 with 500,000 examples.
Pair risk at $(p_1, p_2, I_1, I_2, \sigma) = (.2, .2, 1, 1, .3)$	Exact 0.152641; simulation 0.151474 with 500,000 examples.
Pair risk at $(.1, .4, 1, .3, .7)$	Exact 0.252367; simulation 0.252943.
Pair risk at $(.6, .3, 1, 2, 1)$	Exact 0.900331; simulation 0.901558.
Pair transition kernel	Nonnegative entries; every column sums to one to numerical precision. Marginal-drift checks pass.
Fixed-map covariance identity	Iterative covariance updates agree with the matrix-power expression for random PSD initial covariance.
Same- M counterexample	$M = 1/\sqrt{2}$ for shrinking, stable and expanding maps; spectra and traces confirm all three regimes.
Linear tight-frame comparison	Distributed and mono risks agree: approximately 0.3333333333 in the tested setting.
General five-feature risk and bound	Eight random three-dimensional dictionaries: exact state-enumerated risks agree with independent simulations within 0.005; the stated error bound dominates the exact risks.
Relative J degradation identity	Exact symbolic simplification verifies the positive-factor expression in Volume I; no equation fitting is used.
Gaussian mean-only loop	Initial variance 1 reaches approximately 1.81×10^{-6} after ten refits at noise variance 0.2; posterior sampling preserves variance 1.

The main record is in `research_notes/verification_results.json`, generated by `research_notes/verify_math.py`. Additional checks are recorded in `research_notes/additional_verification.json` by `research_notes/verify_additional.py`. A tolerance assertion checks numerical implementation agreement; it is not a substitute for a proof. The proofs explain why the formulas hold, and the checks help catch algebra, coding and interpretation errors.

Appendix D

Primary references and claim boundaries

1. Zhang et al. *Beyond Interpretability: The Gains of Feature Monosemanticity on Model Robustness*. <https://arxiv.org/abs/2410.21331>. The source anchor; taught in Volume I. It does not analyze recursive diffusion collapse.
2. Elhage et al. *Toy Models of Superposition* (2022). <https://www.anthropic.com/research/toy-models-of-superposition>. The controlled sparse-feature representation framework.
3. Scherlis et al. *Polysemanticity and Capacity in Neural Networks* (2022). <https://arxiv.org/abs/2210.01892>. Its per-feature capacity, Equation (3), is exactly the squared geometric alignment used here; capacity allocation has prior analytic theory.
4. Shumailov et al. *The Curse of Recursion: Training on Generated Data Makes Models Forget* (2023 preprint). <https://arxiv.org/abs/2305.17493>. Prior model-collapse analysis; no universal finite-sample factor is assumed here.
5. Alemohammad et al. *Self-Consuming Generative Models Go MAD* (2023). <https://arxiv.org/abs/2307.01850>. Prior generative self-consumption work; real-data supply and generation choices are central protocol variables.
6. Gerstgrasser et al. *Is Model Collapse Inevitable? Breaking the Curse of Recursion by Accumulating Real and Synthetic Data* (2024). <https://arxiv.org/abs/2404.01413>. Distinguishes accumulation from replacement; supports treating the data policy as part of the problem.
7. Ho, Jain and Abbeel. *Denoising Diffusion Probabilistic Models* (2020). <https://arxiv.org/abs/2006.11239>. Source for the standard diffusion training and sampling framework.
8. Song et al. *Score-Based Generative Modeling through Stochastic Differential Equations* (2020). <https://arxiv.org/abs/2011.13456>. Score-based formulation of generative diffusion.
9. Khelifa, Turner and Venkataramanan. *Recursively Trained Diffusion Models: Limiting Collapse Distribution and Spectral Characterization* (11 June 2026). <https://arxiv.org/abs/2606.13796>. Existing spectral diffusion-collapse work; the abstract and stated contribution were checked for scope. A detailed theorem-by-theorem comparison is still required before claiming novelty against it.

The literature check for these notes is targeted, not an exhaustive novelty certification. The proposed contribution must be positioned by its exact assumptions, mechanism, prediction and empirical evidence. None of the phrases “first”, “strictly generalizes the source paper” or “guaranteed conference paper” follows from the calculations delivered here.

The completed work in this volume is the teaching and derivation foundation. The remaining scientific work is to characterize learned geometry, sharpen predictive claims where possible and validate the same mechanism in a real generative model within the agreed research program.